

CS5228 Knowledge Discovery and Data Mining

AY2024/25 Sem2 By Zhao Peiduo

Lecture 1

Common Data Mining Tasks Data mining encompasses various techniques to analyze and extract patterns from large datasets. Some of the most common data mining tasks include:

- **Association Rules:** This method analyzes *transactional data*, where a transaction is a data record consisting of a set of items from a fixed collection. The goal is to identify *association rules* that predict the occurrence of items based on the presence of other items in the dataset.
- **Clustering:** Clustering involves grouping data points based on a well-defined notion of similarity. The objective is to form *clusters*, ensuring that data points within the same cluster have high intra-cluster similarity while minimizing inter-cluster similarity with other clusters.
- **Classification:** This method uses datasets with multiple attributes to determine the *categorical value* of an attribute as a function of other attribute values. Popular classification techniques include K-Nearest Neighbor, Decision Trees, and Linear Classification.
- **Regression:** Similar to classification, regression also works with datasets having multiple attributes, but it predicts *numerical values* of an attribute as a function of other attributes. Common regression methods include K-Nearest Neighbor, Regression Trees, and Linear Regression.
- **Graph Mining:** This technique analyzes data represented as a graph, $G = (V, E)$, where V represents data points (vertices) and E represents relationships between them (edges). Typical patterns derived from graph mining include identifying communities of nodes and detecting important nodes within the network.
- **Recommender Systems:** Recommender systems work with *user-rated items* (such as movie ratings) to predict missing values and recommend items based on similarities. They exploit item features and user similarities to enhance recommendations.

Types of Attributes

- **Categorical (Qualitative):**
 - **Nominal:**
 - * Values are only labels.
 - * Operations: $=$, \neq
 - * Examples: sex (m/f), eye color, zip code.
 - **Ordinal:**
 - * Values are labels with a meaningful order.
 - * Operations: $=$, \neq , $<$, $>$
 - * Examples: street numbers, education level.
- **Numerical (Quantitative):**
 - **Interval:**
 - * Values are measurements with a meaningful distance.
 - * Operations: $=$, \neq , $<$, $>$, $+$, $-$
 - * Examples: body temperature in °C, calendar dates.
 - **Ratio:**
 - * Values are measurements with a meaningful ratio.
 - * Operations: $=$, \neq , $<$, $>$, $+$, $-$, \times , \div
 - * Examples: age, weight, income, blood pressure.

Types of Data

Data can be classified into three main types based on structure and organization:

- **(Well-)Structured Data:**
 - Highly organized: adheres to a predefined data model.
 - Each object has the same fixed set of attributes.
 - Easy to search, aggregate, manipulate, and analyze.
 - Examples: relational databases, spreadsheets.
- **Semi-Structured Data:**
 - No rigid data model: mix of structured and unstructured data.
 - Data exchange formats: XML, JSON, CSV.
 - Tagged unstructured data (e.g., photo with date/time, location, exposure, resolution, flash, etc.).
- **Unstructured Data:**
 - No fixed data model.
 - Requires more advanced data analysis techniques.
 - Examples: images, videos, audio, text, social media.

Data Quality

- **Noise:** Data can be defined as: true signal + **noise**. Sources of noise include:
 - Sensor readings from faulty devices (e.g., intrinsic noise or external influences).
 - Errors during data entry (by humans or machines).
 - Errors during data transmission.
 - Inconsistencies in data formats (e.g., ISO time vs. Unix time, DD/MM/YYYY vs. MM/DD/YYYY).
 - Inconsistencies in conventions (e.g., meters vs. miles, meters vs. centimeters).
- **Outliers:** An outlier is a data point with attribute values considerably different from other points. Outliers can be classified into:
 - **Outliers as noise:**
 - * They negatively interfere with data analysis.
 - * Removal of outliers or using robust methods is recommended.
 - **Outliers as targets:**
 - * The goal is to detect rare or anomalous events such as credit card fraud detection and intrusion detection in security systems.
- **Missing Values**
 - Common causes of missing values:
 - * Attribute values not collected (e.g., broken sensor, person refused to report age).
 - * Attributes not applicable in all cases (e.g., no income data for children).
 - Handling missing values:
 - * Remove data points with missing values.
 - * Remove attributes with missing values (if not essential).
 - * Try to fill in missing values (e.g., using average temperature from nearby sensors).
- **Duplicates**
 - Duplicates refer to data points representing the same object/entity.
 - * **Exact duplicates:** Data points have identical attribute values.
 - * **Near duplicates:** Data points slightly differ in their attribute values (e.g., same person with phone numbers in different formats).
 - Duplicate elimination:
 - * Relatively easy for exact duplicates.
 - * Challenging for near duplicates.

Exploratory Data Analysis (EDA)

Exploratory Data Analysis (EDA) is an essential step in data analysis to identify potential issues such as noise, outliers, missing values, and class distribution imbalances.

- **Identifying Noise**
 - Using histograms to inspect the distribution of data values.
- **Identifying Noise / Outliers**
 - Using box plots to inspect the distribution of attribute values.
 - * Make outliers explicit.
 - Using scatter plots to inspect correlations.
 - * Not always feasible in practice.
 - * Requires good understanding of data.
- **Handling Missing Values**
 - Example: Default value (0) if people did not disclose weight.
 - * Can negatively affect simple analysis such as calculating means/averages.
- **Distribution of Class Labels**
 - Classification tasks generally benefit from balanced datasets.
 - * Balanced = all classes are (almost) equally represented.
 - * Distribution of classes also affects the evaluation of found patterns.

Data Preprocessing

- **Main Purposes**
 - Improve data quality ("Garbage in, garbage out!").
 - Generate valid input for data mining algorithms.
 - Remove complexity from data to ease analysis.
- **Core Preprocessing Tasks**
 - Data cleaning
 - Data reduction
 - Data transformation
 - Data discretization
- **Data Cleaning**
 - Remove or fill missing values.
 - Identify and remove outliers (if outliers are not the goal of the analysis).
 - Identify and remove/merge duplicates.
 - Correct errors and inconsistencies (e.g., convert inches to centimeters).
 - *Non-trivial tasks that are typically very application-specific.*
- **Data Reduction**
 - **Reducing the number of data points**
 - * Sampling — selecting a subset of data points (random or stratified sampling).
 - * Used for preliminary analysis or large datasets.
 - **Reducing the number of attributes**
 - * Removing irrelevant attributes (e.g., IDs, sensitive attributes).
 - * Dimensionality reduction (PCA, LDA, t-SNE).
 - **Reducing the number of attribute values**
 - * Aggregation or generalization.
 - * Binning with smoothing.
- **Data Transformation**
 - Some data reduction techniques also transform data (e.g., dimensionality reduction, aggregation, binning).
 - Attribute construction:
 - * Add or replace attributes inferred from existing attributes.
 - * Example: weight, volume \rightarrow density.
 - Normalization:
 - * Scaling attribute values to a specified range (e.g., $[0,1]$).
 - * Standardization: scaling using mean and standard deviation.
- **Data Discretization**
 - Converting continuous attributes into ordinal attributes.
 - Some algorithms accept only categorical attributes.
 - Convert a regression task to a classification task.
- **One-Hot Encoding**
 - Converting categorical attributes into numerical attributes.
 - Transform categorical attributes into binary attributes (0/1).
 - Allows the application of numerical methods on categorical attributes.

Lecture 2

Goal of Clustering

- Separate **unlabeled** data into groups of **similar** objects/points
- Maximize **intra-cluster** similarity
- Minimize **inter-cluster** similarity

Meso-level perspective on data

- **micro:** individual data points
- **meso:** clusters
- **macro:** whole dataset

Ingredients for Clustering

- **Representation of objects, e.g.:**
 - (Multidimensional) point coordinates x, y
 - Sets A, B (e.g., items in a transaction)
 - Vectors u, v (e.g., TF-IDF)
- **Similarity Measure, e.g.:**
 - Euclidean Distance: $dist_{\text{euclidean}}(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$
 - Jaccard Similarity: $sim_{\text{jaccard}}(A, B) = \frac{|A \cap B|}{|A \cup B|}$
 - Cosine Similarity: $sim_{\text{cosine}}(u, v) = \frac{u \cdot v}{\|u\| \|v\|}$
- **Clustering Algorithm**
 - Process that determines if an object belongs to a cluster

Types of Clusters

- **Well-separated vs. Center-based**
 - Well-separated: Any object in a cluster is closer to every other object in the cluster than to any point outside the cluster.
 - Center-based: Any object in a cluster is closer to the "center" of the cluster than to the center of any other cluster. The cluster center is commonly called a **centroid**.
- **Contiguity-based vs. Density-based**
 - Contiguity-based: Two objects are in the same cluster if they are more similar than a specified threshold, ensuring each object is more similar to some object in the cluster than to any point in a different cluster.
 - Density-based: A cluster is defined as a dense region of objects surrounded by regions of lower density, which allows better handling of noise.
- **Partitional vs. Hierarchical**
 - Partitional: The data is divided into non-overlapping subsets (i.e., clusters), where each object belongs to exactly one cluster or no cluster at all.
 - Hierarchical: Clusters can be nested, and a point can belong to different clusters depending on the hierarchy level.
- **Exclusive vs. Non-exclusive / Overlapping**
 - Exclusive: Each object belongs to exactly one cluster.
 - Non-exclusive / Overlapping: An object can belong to more than one cluster at a time. Fuzzy clustering assigns each object to all clusters with a certain probability.
- **Complete vs. Partial**
 - Complete: Every object is assigned to at least one cluster, ensuring full coverage of data points.
 - Partial: An object might not belong to any cluster, allowing for the presence of noise and outliers.

K-Means

- **Basic Characteristics**
 - Clusters are centroid-based.
 - Clustering is partitional, exclusive, and complete.
- **Inputs (for d-dimensional Euclidean space)**
 - Data points: (x_1, x_2, \dots, x_N) , $x_i \in \mathbb{R}^d$
 - Number of clusters: $K \rightarrow C_1, C_2, \dots, C_K$ (cluster centers).
- **Optimization Objective**
 - Minimize Sum of Squared Error (SSE): $SSE = \sum_{i=1}^K \sum_{x \in C_i} \|x - c_i\|^2$
 - Finding the optimal solution is NP-hard: $O(N^{Kd+1})$.

K-Means — How to Define the Centroid of a Cluster?

Simple case in Euclidean space

The centroid is derived by minimizing SSE, which turns out to be the mean of all points in that cluster. Proof:

$$\begin{aligned} SSE &= \sum_{i=1}^K \sum_{x \in C_i} \|x - c_i\|^2 \\ \frac{\delta}{\delta c_k} SSE &= \frac{\delta}{\delta c_k} \sum_{i=1}^K \sum_{x \in C_i} \|x - c_i\|^2 \\ &= \sum_{x \in C_k} \frac{\delta}{\delta c_k} \|x - c_k\|^2 \\ &= \sum_{x \in C_k} 2(x - c_k) \end{aligned} \quad \begin{aligned} \sum_{x \in C_k} 2(x - c_k) &= 0 \\ \sum_{x \in C_k} x - \sum_{x \in C_k} c_k &= 0 \\ m_k c_k &= \sum_{x \in C_k} x \\ c_k &= \frac{1}{m_k} \sum_{x \in C_k} x \end{aligned}$$

K-Means — Basic Algorithm (Lloyd's Algorithm)

- **Initialization**
 - Select K points as initial centroids: C_1, C_2, \dots, C_K .
- **Repeat**
 - **Assignment:** Assign each point to the nearest cluster (i.e., centroid).
 - **Update:** Move each centroid to the average of its assigned points.
- **Stopping Criterion**
 - Repeat until no change in assignments.

Handling Empty Clusters

- **Artificially fill empty clusters after assignment step**
 - Replace empty cluster with the point that contributes most to SSE.
 - Replace empty cluster with a point from the cluster with the highest SSE.
- **Postprocessing**
 - Split "loose" clusters with very high SSE.
- **Modification of Lloyd's Algorithm**
 - K-Means variants aim to address the initial centroids issue.

K-Means — Limitations and Potential Workarounds

- **Susceptibility to "Natural" Clusters**
 - K-Means struggles with:
 - * Non-spherical clusters.
 - * Clusters of different sizes.
 - * Clusters of different densities.
 - **Potential Workaround: Choose a Larger Value for K**
 - * Split natural clusters into multiple "well-behaved" subclusters.
 - * Apply suitable postprocessing steps to merge subclusters.
- **Sensitivity to Initial Centroid Selection**
 - Different initializations of centroids may yield:
 - * Different clusterings with varying SSEs (leading to local optima).
 - * Empty clusters if a centroid is "blocked off" by other centroids.
 - **Potential Workaround: Improve Centroid Initialization**
 - * **Artificially Fill Empty Clusters After Assignment**
 - Replace empty cluster with the point that contributes most to SSE.
 - Replace empty cluster with a point from the cluster with the highest SSE.
 - * **Postprocessing**
 - Split "loose" clusters with very high SSE.
 - * **Modification of Lloyd's Algorithm (K-Means Variants)**
 - Use improved centroid initialization techniques.

K-Means Variants

- **K-Means++**
 - Only changes the initialization of centroids.
 - Goal: Spread out centroids for better performance with theoretical guarantees.
 - **Initialization Process:**
 - * Pick a random point as the first centroid C_1 .
 - * Repeat:
 - For each point x , calculate distance d_x to the nearest existing centroid.
 - Pick a random point for the next centroid with probability proportional to d_x^2 .
 - * Until K centroids have been picked.
- **X-Means**
 - Automatic method to choose K .
 - Iteratively applies K-Means with $K = 2$ to refine clustering.
 - **Example Scoring Functions:**
 - * Bayesian Information Criterion (BIC).
 - * Akaike Information Criterion (AIC).
 - * Minimum Description Length (MDL).
- **K-Medoids**
 - **Restriction:** Centroids are chosen from the data points.
 - * Does not require calculation of averages.
 - * Uses only a notion of distance or similarity.
 - * More robust to noise and outliers.
 - **Main Issue: Performance**
 - * More expensive update step.
 - * Swap medoid with each point in the cluster and calculate change in cost (e.g., SSE).
 - * Choose the point as the new medoid that minimizes cost after swapping.

DBSCAN (Density-Based Spatial Clustering of Applications with Noise)

Basic Characteristics

- **Clusters:** Density-based
- **Clustering:** Partitional, Exclusive, Partial

Inputs (for d-dimensional Euclidean space)

- $\mathbf{x} = (x_1, x_2, \dots, x_N)$, $x_i \in \mathbb{R}^d$
- ε (Epsilon): Radius defining a point's neighborhood
- *MinPts*: Minimum number of points in a neighborhood

$$density = \frac{mass}{volume} = \frac{MinPts}{\varepsilon}$$

Types of Points in DBSCAN

- **Core points**
 - Points with at least *MinPts* neighbors within radius ε
 - Form the **interior** of a cluster
- **Border points**
 - Non-core points with at least one core point in their neighborhood
 - Form the **border** of a cluster
- **Outliers / Noise**
 - All other points
 - Default node type

Algorithm 1 DBSCAN Algorithm

- 1: **Input:** Dataset D , radius ε , minimum points *MinPts*
- 2: **Output:** Clusters and outliers
- 3: **for** each point P in D **do**
- 4: **if** P is already visited **then**
- 5: Continue
- 6: **end if**
- 7: Mark P as visited
- 8: Retrieve all neighbors N of P within radius ε
- 9: **if** $|N| < MinPts$ **then**
- 10: Mark P as noise
- 11: **else**
- 12: Create a new cluster
- 13: Expand cluster by adding density-reachable points
- 14: **end if**
- 15: **end for**

Characteristics of DBSCAN

- **DBSCAN always converges**
 - Every data point is explored in either Phase 1 or Phase 2
 - A data point does not change its type (exception: noise \rightarrow border)
- **DBSCAN is not completely deterministic**
 - Phase 1 introduces randomness
 - Border points may be reachable from core points of different clusters
 - Noise and core points are deterministic

Hierarchical Clustering**Basic characteristics**

- Clusters: depends...
- Clustering: hierarchical, complete, exclusive (at each level!)

No parameterization (in principle)

- In practice, typically the number of clusters is specified (similar to K-means)
- Different choices of measures to calculate distances between clusters

Dendrograms

- A dendrogram is a visualization of hierarchical relationships
- Binary tree showing how clusters are hierarchically merged/split
- Each node represents a cluster
- Each leaf is a singleton cluster
- Height reflects distance between clusters

Hierarchical Clustering - Two Main Types

- **Agglomerative (bottom-up)**
 - Start with each point being its own cluster
 - Merge the closest pair of clusters at each step
 - Stop when only one cluster remains
 - Example: **AGNES (Agglomerative Nesting)**
- **Divisive (top-down)**
 - Start with all points in a single cluster
 - Recursively split clusters at each step
 - Stop when each cluster contains a single point
 - Example: **DIANA (Divisive Analysis)**

AGNES Algorithm

Algorithm 2 AGNES Algorithm

- 1: **Input:** Dataset D
 - 2: **Output:** Hierarchical clusters
 - 3: **for** each point P in D **do**
 - 4: Assign P to its own cluster
 - 5: **end for**
 - 6: **while** more than one cluster exists **do**
 - 7: Merge the two closest clusters into one
 - 8: **end while**
-

Single Linkage Clustering

- Distance between clusters = **minimum distance** between two points from each cluster:
 $d_{\text{single}}(C_i, C_j) = \min_{p \in C_i, q \in C_j} d(p, q)$
- **Strength:** Can handle non-globular shapes
- **Weakness:** Very susceptible to noise (*Chaining Effect*)

Complete Linkage Clustering

- Distance between clusters = **maximum distance** between two points from each cluster:
 $d_{\text{complete}}(C_i, C_j) = \max_{p \in C_i, q \in C_j} d(p, q)$
- **Strength:** Less susceptible to noise or outliers
- **Weakness:** Bias towards globular clusters, breaks large clusters

Average Linkage Clustering

- Distance between clusters = **average distance** between two points from each cluster:
 $d_{\text{average}}(C_i, C_j) = \text{avg}_{p \in C_i, q \in C_j} d(p, q)$

Linkage Alternatives

- **Centroid Linkage**
 - Distance between clusters = distance between centroids:
 $d_{\text{centroid}}(C_i, C_j) = d(m_i, m_j)$
- **Ward Linkage**
 - Variance-based merging criterion
 $d_{\text{ward}}(C_i, C_j) = \sum_{k \in C_i \cup C_j} \|x_k - m_{ij}\|^2 - \sum_{k \in C_i} \|x_k - m_i\|^2 - \sum_{k \in C_j} \|x_k - m_j\|^2$

Complexity Analysis

- **Space Complexity:** $O(N^2)$ (storing distance matrix)
- **Time Complexity:**
 - Baseline: $O(N^3)$ - (N-1) steps, each step $O(N^2)$
 - Using heap/priority queue: $O(N^2 \log N)$
 - Single Linkage special optimization: $O(N^2)$

DIANA Algorithm

- **Top-Down Hierarchical Clustering**
 - Start with all points in a single cluster
 - Recursively split clusters until each point is its own cluster
- **Challenge:** 2^n ways to split a cluster with n points
 - Heuristics required to restrict the search space
 - Generally slower and less common than AGNES
- **Cases where DIANA is preferable:**
 - When no complete clustering is needed (early stopping)
 - When splitting can use global knowledge