

# CS5228 Knowledge Discovery and Data Mining

AY2024/25 Sem2 By Zhao Peiduo

## Lecture 1

### Introduction

**Common Data Mining Tasks** Data mining encompasses various techniques to analyze and extract patterns from large datasets. Some of the most common data mining tasks include:

- **Association Rules:** This method analyzes *transactional data*, where a transaction is a data record consisting of a set of items from a fixed collection. The goal is to identify *association rules* that predict the occurrence of items based on the presence of other items in the dataset.
- **Clustering:** Clustering involves grouping data points based on a well-defined notion of similarity. The objective is to form *clusters*, ensuring that data points within the same cluster have high intra-cluster similarity while minimizing inter-cluster similarity with other clusters.
- **Classification:** This method uses datasets with multiple attributes to determine the *categorical value* of an attribute as a function of other attribute values. Popular classification techniques include K-Nearest Neighbor, Decision Trees, and Linear Classification.
- **Regression:** Similar to classification, regression also works with datasets having multiple attributes, but it predicts *numerical values* of an attribute as a function of other attributes. Common regression methods include K-Nearest Neighbor, Regression Trees, and Linear Regression.
- **Graph Mining:** This technique analyzes data represented as a graph,  $G = (V, E)$ , where  $V$  represents data points (vertices) and  $E$  represents relationships between them (edges). Typical patterns derived from graph mining include identifying communities of nodes and detecting important nodes within the network.
- **Recommender Systems:** Recommender systems work with *user-rated items* (such as movie ratings) to predict missing values and recommend items based on similarities. They exploit item features and user similarities to enhance recommendations.

#### Types of Attributes

- **Categorical (Qualitative):**
  - **Nominal:**
    - \* Values are only labels.
    - \* Operations:  $=$ ,  $\neq$
    - \* Examples: sex (m/f), eye color, zip code.
  - **Ordinal:**
    - \* Values are labels with a meaningful order.
    - \* Operations:  $=$ ,  $\neq$ ,  $<$ ,  $>$
    - \* Examples: street numbers, education level.
- **Numerical (Quantitative):**
  - **Interval:**
    - \* Values are measurements with a meaningful distance.
    - \* Operations:  $=$ ,  $\neq$ ,  $<$ ,  $>$ ,  $+$ ,  $-$
    - \* Examples: body temperature in °C, calendar dates.
  - **Ratio:**
    - \* Values are measurements with a meaningful ratio.
    - \* Operations:  $=$ ,  $\neq$ ,  $<$ ,  $>$ ,  $+$ ,  $-$ ,  $\times$ ,  $\div$
    - \* Examples: age, weight, income, blood pressure.

#### Types of Data

Data can be classified into three main types based on structure and organization:

- **(Well-)Structured Data:**
  - Highly organized: adheres to a predefined data model.
  - Each object has the same fixed set of attributes.
  - Easy to search, aggregate, manipulate, and analyze.
  - Examples: relational databases, spreadsheets.
- **Semi-Structured Data:**
  - No rigid data model: mix of structured and unstructured data.
  - Data exchange formats: XML, JSON, CSV.
  - Tagged unstructured data (e.g., photo with date/time, location, exposure, resolution, flash, etc.).
- **Unstructured Data:**
  - No fixed data model.
  - Requires more advanced data analysis techniques.
  - Examples: images, videos, audio, text, social media.

#### Data Quality

- **Noise:** Data can be defined as: true signal + **noise**. Sources of noise include:
  - Sensor readings from faulty devices (e.g., intrinsic noise or external influences).
  - Errors during data entry (by humans or machines).
  - Errors during data transmission.
  - Inconsistencies in data formats (e.g., ISO time vs. Unix time, DD/MM/YYYY vs. MM/DD/YYYY).
  - Inconsistencies in conventions (e.g., meters vs. miles, meters vs. centimeters).
- **Outliers:** An outlier is a data point with attribute values considerably different from other points. Outliers can be classified into:
  - **Outliers as noise:**
    - \* They negatively interfere with data analysis.
    - \* Removal of outliers or using robust methods is recommended.
  - **Outliers as targets:**
    - \* The goal is to detect rare or anomalous events such as credit card fraud detection and intrusion detection in security systems.
- **Missing Values**
  - Common causes of missing values:
    - \* Attribute values not collected (e.g., broken sensor, person refused to report age).
    - \* Attributes not applicable in all cases (e.g., no income data for children).
  - Handling missing values:
    - \* Remove data points with missing values.
    - \* Remove attributes with missing values (if not essential).
    - \* Try to fill in missing values (e.g., using average temperature from nearby sensors).
- **Duplicates**
  - Duplicates refer to data points representing the same object/entity.
    - \* **Exact duplicates:** Data points have identical attribute values.
    - \* **Near duplicates:** Data points slightly differ in their attribute values (e.g., same person with phone numbers in different formats).
  - Duplicate elimination:
    - \* Relatively easy for exact duplicates.
    - \* Challenging for near duplicates.

#### Exploratory Data Analysis (EDA)

Exploratory Data Analysis (EDA) is an essential step in data analysis to identify potential issues such as noise, outliers, missing values, and class distribution imbalances.

- **Identifying Noise**
  - Using histograms to inspect the distribution of data values.
- **Identifying Noise / Outliers**
  - Using box plots to inspect the distribution of attribute values.
    - \* Make outliers explicit.
  - Using scatter plots to inspect correlations.
    - \* Not always feasible in practice.
    - \* Requires good understanding of data.
- **Handling Missing Values**
  - Example: Default value (0) if people did not disclose weight.
    - \* Can negatively affect simple analysis such as calculating means/averages.
- **Distribution of Class Labels**
  - Classification tasks generally benefit from balanced datasets.
    - \* Balanced = all classes are (almost) equally represented.
  - Distribution of classes also affects the evaluation of found patterns.

#### Data Preprocessing

- **Main Purposes**
  - Improve data quality ("Garbage in, garbage out!").
  - Generate valid input for data mining algorithms.
  - Remove complexity from data to ease analysis.
- **Core Preprocessing Tasks**
  - Data cleaning
  - Data reduction
  - Data transformation
  - Data discretization
- **Data Cleaning**
  - Remove or fill missing values.
  - Identify and remove outliers (if outliers are not the goal of the analysis).
  - Identify and remove/merge duplicates.
  - Correct errors and inconsistencies (e.g., convert inches to centimeters).
  - *Non-trivial tasks that are typically very application-specific.*
- **Data Reduction**
  - **Reducing the number of data points**
    - \* Sampling — selecting a subset of data points (random or stratified sampling).
    - \* Used for preliminary analysis or large datasets.
  - **Reducing the number of attributes**
    - \* Removing irrelevant attributes (e.g., IDs, sensitive attributes).
    - \* Dimensionality reduction (PCA, LDA, t-SNE).
  - **Reducing the number of attribute values**
    - \* Aggregation or generalization.
    - \* Binning with smoothing.
- **Data Transformation**
  - Some data reduction techniques also transform data (e.g., dimensionality reduction, aggregation, binning).
  - Attribute construction:
    - \* Add or replace attributes inferred from existing attributes.
    - \* Example: weight, volume  $\rightarrow$  density.
  - Normalization:
    - \* Scaling attribute values to a specified range (e.g., [0,1]).
    - \* Standardization: scaling using mean and standard deviation.
- **Data Discretization**
  - Converting continuous attributes into ordinal attributes.
  - Some algorithms accept only categorical attributes.
  - Convert a regression task to a classification task.
- **One-Hot Encoding**
  - Converting categorical attributes into numerical attributes.
  - Transform categorical attributes into binary attributes (0/1).
  - Allows the application of numerical methods on categorical attributes.

## Lecture 2

### Clustering

#### Goal of Clustering

- Separate **unlabeled** data into groups of **similar** objects/points
- Maximize **intra-cluster** similarity
- Minimize **inter-cluster** similarity

#### Meso-level perspective on data

- **micro:** individual data points
- **meso:** clusters
- **macro:** whole dataset

#### Ingredients for Clustering

- **Representation of objects, e.g.:**
  - (Multidimensional) point coordinates  $x, y$
  - Sets  $A, B$  (e.g., items in a transaction)
  - Vectors  $u, v$  (e.g., TF-IDF)
- **Similarity Measure, e.g.:**
  - Euclidean Distance:  $dist_{\text{euclidean}}(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$
  - Jaccard Similarity:  $sim_{\text{jaccard}}(A, B) = \frac{|A \cap B|}{|A \cup B|}$
  - Cosine Similarity:  $sim_{\text{cosine}}(u, v) = \frac{u \cdot v}{\|u\| \|v\|}$
- **Clustering Algorithm**
  - Process that determines if an object belongs to a cluster

#### Types of Clusters

- **Well-separated vs. Center-based**
  - Well-separated: Any object in a cluster is closer to every other object in the cluster than to any point outside the cluster.
  - Center-based: Any object in a cluster is closer to the "center" of the cluster than to the center of any other cluster. The cluster center is commonly called a **centroid**.
- **Contiguity-based vs. Density-based**
  - Contiguity-based: Two objects are in the same cluster if they are more similar than a specified threshold, ensuring each object is more similar to some object in the cluster than to any point in a different cluster.
  - Density-based: A cluster is defined as a dense region of objects surrounded by regions of lower density, which allows better handling of noise.
- **Partitional vs. Hierarchical**
  - Partitional: The data is divided into non-overlapping subsets (i.e., clusters), where each object belongs to exactly one cluster or no cluster at all.
  - Hierarchical: Clusters can be nested, and a point can belong to different clusters depending on the hierarchy level.
- **Exclusive vs. Non-exclusive / Overlapping**
  - Exclusive: Each object belongs to exactly one cluster.
  - Non-exclusive / Overlapping: An object can belong to more than one cluster at a time. Fuzzy clustering assigns each object to all clusters with a certain probability.
- **Complete vs. Partial**
  - Complete: Every object is assigned to at least one cluster, ensuring full coverage of data points.
  - Partial: An object might not belong to any cluster, allowing for the presence of noise and outliers.

## K-Means

- **Basic Characteristics**
  - Clusters are centroid-based.
  - Clustering is partitional, exclusive, and complete.
- **Inputs (for d-dimensional Euclidean space)**
  - Data points:  $(x_1, x_2, \dots, x_N)$ ,  $x_i \in \mathbb{R}^d$
  - Number of clusters:  $K \rightarrow C_1, C_2, \dots, C_K$  (cluster centers).
- **Optimization Objective**
  - Minimize Sum of Squared Error (SSE):  $SSE = \sum_{i=1}^K \sum_{x \in C_i} \|x - c_i\|^2$
  - Finding the optimal solution is NP-hard:  $O(N^{Kd+1})$ .

#### K-Means — How to Define the Centroid of a Cluster?

##### Simple case in Euclidean space

The centroid is derived by minimizing SSE, which turns out to be the mean of all points in that cluster. Proof:

$$\begin{aligned} SSE &= \sum_{i=1}^K \sum_{x \in C_i} \|x - c_i\|^2 \\ \frac{\delta}{\delta c_k} SSE &= \frac{\delta}{\delta c_k} \sum_{i=1}^K \sum_{x \in C_i} \|x - c_i\|^2 \\ &= \sum_{x \in C_k} \frac{\delta}{\delta c_k} \|x - c_k\|^2 \\ &= \sum_{x \in C_k} 2(x - c_k) \end{aligned}$$
$$\begin{aligned} \sum_{x \in C_k} 2(x - c_k) &= 0 \\ \sum_{x \in C_k} x - \sum_{x \in C_k} c_k &= 0 \\ m_k c_k &= \sum_{x \in C_k} x \\ c_k &= \frac{1}{m_k} \sum_{x \in C_k} x \end{aligned}$$

#### K-Means — Basic Algorithm (Lloyd's Algorithm)

- **Initialization**
  - Select  $K$  points as initial centroids:  $C_1, C_2, \dots, C_K$ .
- **Repeat**
  - **Assignment:** Assign each point to the nearest cluster (i.e., centroid).
  - **Update:** Move each centroid to the average of its assigned points.
- **Stopping Criterion**
  - Repeat until no change in assignments.

#### Handling Empty Clusters

- **Artificially fill empty clusters after assignment step**
  - Replace empty cluster with the point that contributes most to SSE.
  - Replace empty cluster with a point from the cluster with the highest SSE.

- **Postprocessing**
  - Split "loose" clusters with very high SSE.
- **Modification of Lloyd's Algorithm**
  - K-Means variants aim to address the initial centroids issue.

#### K-Means — Limitations and Potential Workarounds

- **Susceptibility to "Natural" Clusters**
  - K-Means struggles with:
    - \* Non-spherical clusters.
    - \* Clusters of different sizes.
    - \* Clusters of different densities.
  - **Potential Workaround: Choose a Larger Value for K**
    - \* Split natural clusters into multiple "well-behaved" subclusters.
    - \* Apply suitable postprocessing steps to merge subclusters.
- **Sensitivity to Initial Centroid Selection**
  - Different initializations of centroids may yield:
    - \* Different clusterings with varying SSEs (leading to local optima).
    - \* Empty clusters if a centroid is "blocked off" by other centroids.
  - **Potential Workaround: Improve Centroid Initialization**
    - \* **Artificially Fill Empty Clusters After Assignment**
      - Replace empty cluster with the point that contributes most to SSE.
      - Replace empty cluster with a point from the cluster with the highest SSE.
    - \* **Postprocessing**
      - Split "loose" clusters with very high SSE.
    - \* **Modification of Lloyd's Algorithm (K-Means Variants)**
      - Use improved centroid initialization techniques.

#### K-Means Variants

- **K-Means++**
  - Only changes the initialization of centroids.
  - Goal: Spread out centroids for better performance with theoretical guarantees.
  - **Initialization Process:**
    - \* Pick a random point as the first centroid  $C_1$ .
    - \* Repeat:
      - For each point  $x$ , calculate distance  $d_x$  to the nearest existing centroid.
      - Pick a random point for the next centroid with probability proportional to  $d_x^2$ .
    - \* Until  $K$  centroids have been picked.
- **X-Means**
  - Automatic method to choose  $K$ .
  - Iteratively applies K-Means with  $K = 2$  to refine clustering.
  - **Example Scoring Functions:**
    - \* Bayesian Information Criterion (BIC).
    - \* Akaike Information Criterion (AIC).
    - \* Minimum Description Length (MDL).
- **K-Medoids**
  - **Restriction:** Centroids are chosen from the data points.
    - \* Does not require calculation of averages.
    - \* Uses only a notion of distance or similarity.
    - \* More robust to noise and outliers.
  - **Main Issue: Performance**
    - \* More expensive update step.
    - \* Swap medoid with each point in the cluster and calculate change in cost (e.g., SSE).
    - \* Choose the point as the new medoid that minimizes cost after swapping.

## DBSCAN (Density-Based Spatial Clustering of Applications with Noise)

#### Basic Characteristics

- **Clusters:** Density-based
- **Clustering:** Partitional, Exclusive, Partial

#### Inputs (for d-dimensional Euclidean space)

- $\mathbf{x} = (x_1, x_2, \dots, x_N)$ ,  $x_i \in \mathbb{R}^d$
- $\epsilon$  (Epsilon): Radius defining a point's neighborhood
- *MinPts*: Minimum number of points in a neighborhood

$$density = \frac{mass}{volume} = \frac{MinPts}{\epsilon^d}$$

#### Types of Points in DBSCAN

- **Core points**
  - Points with at least *MinPts* neighbors within radius  $\epsilon$
  - Form the **interior** of a cluster
- **Border points**
  - Non-core points with at least one core point in their neighborhood
  - Form the **border** of a cluster
- **Outliers / Noise**
  - All other points
  - Default node type

## Algorithm 1 DBSCAN Algorithm

```
1: Input: Dataset  $D$ , radius  $\epsilon$ , minimum points  $MinPts$ 
2: Output: Clusters and outliers
3: for each point  $P$  in  $D$  do
4:   if  $P$  is already visited then
5:     Continue
6:   end if
7:   Mark  $P$  as visited
8:   Retrieve all neighbors  $N$  of  $P$  within radius  $\epsilon$ 
9:   if  $|N| < MinPts$  then
10:    Mark  $P$  as noise
11:  else
12:    Create a new cluster
13:    Expand cluster by adding density-reachable points
14:  end if
15: end for
```

#### Characteristics of DBSCAN

- **DBSCAN always converges**
  - Every data point is explored in either Phase 1 or Phase 2
  - A data point does not change its type (exception: noise  $\rightarrow$  border)
- **DBSCAN is not completely deterministic**
  - Phase 1 introduces randomness
  - Border points may be reachable from core points of different clusters
  - Noise and core points are deterministic

#### DBSCAN — Limitations

- DBSCAN cannot handle different densities.
- DBSCAN is generally very sensitive to parameters. Choosing  $\epsilon$  and *MinPts* requires a good understanding of data and context.

#### DBSCAN — How to Choose Parameter Values?

- Informed by results of EDA, e.g.:
  - Distribution of all pairwise distances.
  - First insights into suitable values for  $\epsilon$ .
- Density of data points has intuitive semantic meaning, e.g.:
  - Geographic distance between bars in a city.
  - Task: Find areas (clusters) with more than 10 bars within 500m.

## Hierarchical Clustering

### Basic characteristics

- Clusters: depends...
- Clustering: hierarchical, complete, exclusive (at each level!)

### No parameterization (in principle)

- In practice, typically the number of clusters is specified (similar to K-means)
- Different choices of measures to calculate distances between clusters

### Dendrograms

- A dendrogram is a visualization of hierarchical relationships
- Binary tree showing how clusters are hierarchically merged/split
- Each node represents a cluster
- Each leaf is a singleton cluster
- Height reflects distance between clusters

### Two Main Types of Hierarchical Clustering

- **Agglomerative (bottom-up)**
  - Start with each point being its own cluster
  - Merge the closest pair of clusters at each step
  - Stop when only one cluster remains
  - Example: **AGNES (Agglomerative Nesting)**
- **Divisive (top-down)**
  - Start with all points in a single cluster
  - Recursively split clusters at each step
  - Stop when each cluster contains a single point
  - Example: **DIANA (Divisive Analysis)**

## AGNES Algorithm

### Algorithm 2 AGNES Algorithm

```

1: Input: Dataset  $D$ 
2: Output: Hierarchical clusters
3: for each point  $P$  in  $D$  do
4:   Assign  $P$  to its own cluster
5: end for
6: while more than one cluster exists do
7:   Merge the two closest clusters into one
8: end while

```

### AGNES Algorithm Linkages

- **Single Linkage Clustering**
  - Distance between clusters = **minimum distance** between two points from each cluster:  
 $d_{\text{single}}(C_i, C_j) = \min_{p \in C_i, q \in C_j} d(p, q)$
  - **Strength:** Can handle non-globular shapes.
  - **Weakness:** Very susceptible to noise (*Chaining Effect*).
- **Complete Linkage Clustering**
  - Distance between clusters = **maximum distance** between two points from each cluster:  
 $d_{\text{complete}}(C_i, C_j) = \max_{p \in C_i, q \in C_j} d(p, q)$
  - **Strength:** Less susceptible to noise or outliers.
  - **Weakness:** Bias towards globular clusters, breaks large clusters.
- **Average Linkage Clustering**
  - Distance between clusters = **average distance** between two points from each cluster:  
 $d_{\text{average}}(C_i, C_j) = \text{avg}_{p \in C_i, q \in C_j} d(p, q)$
- **Centroid Linkage**
  - Distance between clusters = distance between centroids:  
 $d_{\text{centroid}}(C_i, C_j) = d(m_i, m_j)$
- **Ward Linkage**
  - Variance-based merging criterion:

$$d_{\text{ward}}(C_i, C_j) = \sum_{k \in C_i \cup C_j} \|x_k - m_{ij}\|^2 - \sum_{k \in C_i} \|x_k - m_i\|^2 - \sum_{k \in C_j} \|x_k - m_j\|^2$$

### AGNES Complexity Analysis

- **Space Complexity:**  $O(N^2)$  (storing distance matrix)
- **Time Complexity:**
  - Baseline:  $O(N^3)$  - (N-1) steps, each step  $O(N^2)$
  - Using heap/priority queue:  $O(N^2 \log N)$
  - Single Linkage special optimization:  $O(N^2)$

## DIANA Algorithm

- **Top-Down Hierarchical Clustering**
  - Start with all points in a single cluster
  - Recursively split clusters until each point is its own cluster
- **Challenge:**  $2^n$  ways to split a cluster with  $n$  points
  - Heuristics required to restrict the search space
  - Generally slower and less common than AGNES
- **Cases where DIANA is preferable:**
  - When no complete clustering is needed (early stopping)
  - When splitting can use global knowledge

## Cluster Evaluation

### Cluster Evaluation Problems

- Challenges in Visual Inspection: Just eyeballing the clustering is rarely possible.
- Algorithmic Bias: Clustering algorithms will always find some clusters.

### Purpose of Cluster Evaluation

- Comparing the results of different clustering algorithms.
- Comparing the results of a clustering algorithm with different parameters.
- Minimizing the effects of noise on the clustering.

### Two Main Approaches of Clustering Evaluation

- **External quality measures:** Evaluate clustering against a ground truth (if available).
- **Internal quality measures:** Evaluate clustering from the data itself.

### External Quality Measures

- **Ground Truth: Labeled Data**
  - Labels indicate that two points "belong together."
  - If clustering reflects this, it is considered good clustering.
- **Cluster Purity**
  - Measures how pure clusters are concerning the most common label.
  - Formula:  $P = \frac{1}{N} \sum_{c \in C} \max_{l \in L} |c \cap l|$
  - Example Calculation:  $P = \frac{1}{8}(3 + 4) = 0.875$
  - **Limitation:** Purity does not penalize having many clusters.
  - $P = 1$  is easy to achieve if each cluster contains only a single point.
- **Information Retrieval Metrics**
  - **True Positives (TP):** Same cluster, same label.
  - **True Negatives (TN):** Different clusters, different labels.
  - **False Positives (FP):** Same cluster, different labels.
  - **False Negatives (FN):** Different clusters, same label.
- **Rand Index (RI)**
  - Reflects clustering accuracy.
  - Formula:  $RI = \frac{TP+TN}{TP+FP+FN+TN}$
  - Example:  $RI_{\text{example}} = 0.82$
- **Precision, Recall, and F1-Score**
  - Precision:  $P = \frac{TP}{TP+FP}$
  - Recall:  $R = \frac{TP}{TP+FN}$
  - F1-Score:  $F1 = \frac{2 \cdot P \cdot R}{P+R}$
  - Example values:  $P_{\text{example}} = 0.75$ ,  $R_{\text{example}} = 0.69$ ,  $F1_{\text{example}} = 0.72$

### Internal Quality Measures

- **Sum of Squared Errors (SSE)**
  - SSE is often used to determine the number of clusters.
  - Formula:  $SSE = \sum_{i=1}^k \sum_{x \in C_i} \|x - \mu_i\|^2$
  - **Limitations:**
    - \* Does not penalize having a large number of clusters.
    - \* SSE decreases as the number of clusters increases.
    - \* Applicable beyond K-Means but less intuitive for non-globular clusters.
  - SSE can also be used for complex data but inherently favors globular clusters.
- **Silhouette Coefficient**
  - A good clustering has:
    - \* High inter-cluster distances.
    - \* Low intra-cluster distances.
  - Defined for each data point  $x$ :
    - \* **Cohesion**  $a(x)$ : average distance to points in the same cluster.
    - \* **Separation**  $b(x)$ : minimum average distance to points in a different cluster.
    - \* **Silhouette Score:**

$$s(x) = \frac{b(x) - a(x)}{\max\{a(x), b(x)\}}, \quad \text{if } |C_x| > 1$$
  - \* If  $|C_x| = 1$ , then  $s(x) = 0$ .
  - The overall **\*\*Silhouette Coefficient\*\*** for the clustering:

$$SC = \frac{1}{N} \sum_{i=1}^N s(x_i)$$

- Interpretation:  $-1 \leq s(x) \leq 1$  where **\*\*negative values indicate poor clustering\*\*** and **\*\*values closer to 1 indicate well-separated clusters\*\***.

### Pragmatic Considerations

- The choice of the "best" clustering is often pragmatic:
  - Fixed number of clusters (problematic for DBSCAN).
  - Parameters defined by tasks (e.g., "areas with more than 5 McDonalds within 500m").
  - Maximum, minimum, or average size of clusters.
  - Focus on individual clusters instead of the whole clustering (e.g., largest/smallest cluster).
  - Setting  $k$  too high and merging later if needed.

## Association Rule Mining

### Basic Setup

- **Input database:**
  - Set of **transactions**
  - Transaction = set of **items**
- **Output: Association Rules**
  - Rules predicting the occurrence of some items based on the occurrence of other items.

Example Association Rules

- Antecedent → Consequent
- {item2, item3} → {item5}
- {item1} → {item3}

Problem Statement

- Association rules are not "hard"rules
  - Example: {cereal} → {milk} does not mean customers always buy milk with cereal.
  - Each possible combination (e.g., {yogurt, bread} → {milk}) is a potential rule.
- Given  $d$  unique items, the number of possible rules is:  $3^d - 2^{d+1} + 1$
- Example: If  $d = 6$ , then there are 602 possible rules.

Applications of Association Rule Mining

- Market Basket Analysis: Understanding customer shopping behavior.
- Medical Data Analysis: Diagnosis Support Systems and ADR discovery (adverse drug reaction)
- Census Data Analysis: Getting insights into a population.
- Behavior Data Analysis: User preferences & linkings.

Association Rule Definitions

- Itemset: A subset of items.
- K-itemset: An itemset containing  $k$  items.
- Support Count (SC): Number of transactions containing an itemset.
- Support (S): Fraction of transactions containing an itemset.
- Frequent Itemset: An itemset with a support greater or equal to a minimum threshold  $minsup$ .
- Association Rule: An implication expression  $X \rightarrow Y$ , where  $X$  and  $Y$  are itemsets.
- Support of an Association Rule: Fraction of transactions containing all items in an association rule  $X \rightarrow Y$ :

$$S(X \rightarrow Y) = \frac{SC(X \cup Y)}{N} = S(X \cup Y)$$

- Confidence of an Association Rule: Probability of  $Y$  given  $X$ :

$$C(X \rightarrow Y) = \frac{S(X \rightarrow Y)}{S(X)} = \frac{S(X \cup Y)}{S(X)}$$

Support-Confidence Combinations

- Low Support, Low Confidence: Items in  $(X \cup Y)$  do not frequently appear together. Even when items in  $X$  appear together, they often do so without items in  $Y$ .
- High Support, Low Confidence: Items in  $(X \cup Y)$  frequently appear together, but if items in  $X$  appear together, they often do so without items in  $Y$ .
- Low Support, High Confidence: Items in  $(X \cup Y)$  do not frequently appear together, but when items in  $X$  appear together, they often do so with items in  $Y$ .
- High Support, High Confidence: Items in  $(X \cup Y)$  frequently appear together, and if items in  $X$  appear together, they often do so with items in  $Y$ .

Brute Force Algorithm

- Finding Association Rules: Given a set of transactions, find all association rules  $X \rightarrow Y$  with:
  - Support  $S(X \rightarrow Y) \geq minsup$
  - Confidence  $C(X \rightarrow Y) \geq minconf$
- Brute Force Algorithm:
  - List all possible association rules  $X \rightarrow Y$ .
  - Calculate support  $S(X \rightarrow Y)$  and confidence  $C(X \rightarrow Y)$  for each rule.
  - Drop rules where  $S(X \rightarrow Y) < minsup$  and  $C(X \rightarrow Y) < minconf$ .

Algorithm 3 Brute Force Association Rule Mining

- Input: Transaction database  $D$ , minimum support  $minsup$ , minimum confidence  $minconf$
- Output: Set of association rules  $R$
- Initialize an empty set  $F$  for frequent itemsets
- Initialize an empty set  $R$  for valid rules

▷ Frequent Itemset Generation

- ```
5: for each transaction  $t$  in  $D$  do
6:   for  $k = 1$  to  $|t|$  do
7:     Generate all possible  $k$ -itemsets
8:     for each itemset  $I$  in  $k$ -itemsets do
9:       Count occurrences of  $I$ 
10:    end for
11:  end for
12: end for
13: Retain only itemsets with support  $\geq minsup$ 
```

▷ Rule Generation

- ```
14: for each frequent itemset do
15:   Generate all possible rules  $X \rightarrow Y$ 
16:   Compute  $S(X \rightarrow Y)$  and  $C(X \rightarrow Y)$ 
17:   if  $S(X \rightarrow Y) \geq minsup$  and  $C(X \rightarrow Y) \geq minconf$  then
18:     Add  $X \rightarrow Y$  to  $R$ 
19:   end if
20: end for
21: return  $R$ 
```

Computational Complexity:

- Given  $d$  unique items, the number of possible rules is:

$$3^d - 2^{d+1} + 1 \in O(3^d)$$

- Let  $w$  be the maximum number of items in a transaction within the database:

$$O(N \cdot (3^w - 2^{w+1} + 1))$$

- where typically  $w \ll d$ .
- Frequent itemset generation complexity:

$$O(N \cdot 2^w)$$

Observation: Decoupling Support and Confidence:

- A rule  $X \rightarrow Y$  has sufficient support only if  $X \cup Y$  is a frequent itemset.
- No need to calculate confidence of rules where  $X \cup Y$  is not a frequent itemset.

Apriori Algorithm

- Apriori Principle (Anti-Monotonicity):
  - If an itemset  $X$  is "not frequent", then any superset  $Y \supset X$  is "also not frequent".
  - If an itemset  $Y$  is "frequent", then all its subsets  $X \subset Y$  must also be "frequent".
- Apriori Algorithm:
  - Iteratively generate candidate frequent itemsets.
  - Use previously found frequent  $(k - 1)$ -itemsets to "prune" candidate  $k$ -itemsets.
  - Filter frequent itemsets based on the minimum support threshold.

Algorithm 4 Apriori Algorithm for Frequent Itemset Mining

- Input: Transaction database  $D$ , minimum support  $minsup$
- Output: Frequent itemsets  $F$
- Initialize  $F_1$  with frequent 1-itemsets

▷ Frequent Itemset Generation

- ```
4: for  $k = 2$  to max itemset length do
5:   Generate candidate  $k$ -itemsets  $L_k$  from  $F_{k-1}$ 
6:   Prune infrequent itemsets from  $L_k$ 
7:   for each transaction  $t$  in  $D$  do
8:     for each candidate itemset  $I$  in  $L_k$  do
9:       Count occurrences of  $I$ 
10:    end for
11:  end for
12:  Retain only itemsets with support  $\geq minsup$  as  $F_k$ 
13:  if  $F_k$  is empty then
14:    break
15:  end if
16: end for
```

▷ Rule Generation

- ```
17: for each frequent itemset do
18:   Generate all possible rules  $X \rightarrow Y$ 
19:   Compute  $S(X \rightarrow Y)$  and  $C(X \rightarrow Y)$ 
20:   if  $C(X \rightarrow Y) \geq minconf$  then
21:     Add  $X \rightarrow Y$  to  $R$ 
22:   end if
23: end for
24: return  $F$ 
```

- Rule Generation:
  - Generate association rules  $X \rightarrow Y$  from frequent itemsets.
  - Compute confidence:

$$C(X \rightarrow Y) = \frac{S(X \cup Y)}{S(X)}$$

- Retain rules where  $C(X \rightarrow Y) \geq minconf$ .
- Computational Complexity:
  - Generates at most  $2^d - 1$  itemsets for  $d$  unique items.
  - Reduces complexity by "pruning" infrequent itemsets early.
  - Worst-case complexity:

$$O(N \cdot 2^w)$$

where  $N$  is the number of transactions and  $w$  is the largest transaction size.

- Advantages:
  - Reduces the number of candidate itemsets compared to brute-force.
  - Efficient for dense datasets with many frequent itemsets.
- Limitations:
  - Computationally expensive when itemsets are very large.
  - Many database scans required, increasing execution time.

## Classification & Regression

### Classification AND Regression:

- Core tasks of supervised machine learning.
- Training:** Find patterns in a dataset between the values of **dependent variable(s)** given the values of **independent variable(s) / features**
- Prediction:** Use patterns to assign values to dependent variables for new/unseen data.

### Classification VS. Regression:

- Classification:** Dependent variable is categorical.
- Regression:** Dependent variable is continuous.

### Application Examples

- Real Estate:** Prediction of flat prices (Regression task).
- Health Analytics:** Prediction of heart disease (Classification task).
- Text Analytics:** Sentiment Analysis / Opinion Mining (over text).
- Self-Driving Vehicles:**
  - Input:** Image & sensor data.
  - Regression tasks, e.g.:**
    - Acceleration.
    - Steering angle.
    - Event prediction (%).
  - Classification tasks, e.g.:**
    - Obstacle detection.
    - Street sign identification.

## Supervised Learning

- Basic Setup:**
  - Given: Dataset  $D$  of pairs  $\{(x, y)\}$ 
    - $x$  — features (independent variables).
    - $y$  — label (dependent variable).
  - Split dataset  $D$  into:
    - $D_{train}$  — data for training the model.
    - $D_{test}$  — data for evaluating the model.
    - Important:  $D_{train} \cap D_{test} = \emptyset$  (test data is not used during training).
- Model:** Parameterized function  $h(x, \Theta) = y$ .
  - Maps input space (features) to output space (labels).
  - $\Theta$  — trainable/learnable parameters of the model.
- Model Selection:**
  - Selection of a "family" of functions  $h(x, \Theta) = y$ .
  - Examples:
    - K-Nearest Neighbor.
    - Decision Trees.
    - Linear Models.
    - Neural Networks.
- Training / Learning:**
  - Process of systematically finding the best values for  $\Theta$ .
  - $\Theta_{best} \iff$  best mapping between features and labels.
- Evaluation Process:**
  - Train model on  $(X_{train}, y_{train})$ .
  - Test model on unseen  $(X_{test})$ .
  - Generate predictions  $\hat{y}_{test}$ .
  - Compare  $\hat{y}_{test}$  with ground truth  $y_{test}$ .

## Classification and Regression Evaluation

- Regression Evaluation**
  - Direct comparison of numerical values.
  - Common metric: Root Mean Squared Error (RMSE):

$$RMSE = \sqrt{\frac{\sum_{i=1}^N (\hat{y}_i - y_i)^2}{N}}$$

- Classification Evaluation:** Comparison using a confusion matrix:
  - True Positives (TP): Correctly predicted positives.
  - True Negatives (TN): Correctly predicted negatives.
  - False Positives (FP): Incorrectly predicted positives.
  - False Negatives (FN): Incorrectly predicted negatives.

### Popular Classification Metrics

$$Accuracy = \frac{TP + TN}{TP + FP + TN + FN}$$

$$Sensitivity/Recall = \frac{TP}{TP + FN}$$

$$Specificity = \frac{TN}{TN + FP}$$

$$Precision = \frac{TP}{TP + FP}, \quad Recall = \frac{TP}{TP + FN}$$

$$F1 = 2 \cdot \frac{Precision \cdot Recall}{Precision + Recall}$$

### Challenges in Evaluation

- Imbalanced datasets:** Standard metrics may not be reliable
- False Positives (FP) vs False Negatives (FN):** Different costs of errors
  - Example: **Heart disease prediction** (Recall > Precision)
  - Example: **News classification** (Precision > Recall)

### Numerical Class Scores

- Many models output class scores instead of hard labels
- Thresholding:** Convert scores to binary labels (e.g., threshold = 0.5)

### Receiver Operating Characteristic (ROC) Curve

- Plots True Positive Rate (Sensitivity) vs. False Positive Rate (1-Specificity)
- Area Under ROC Curve (AUROC) quantifies classifier performance
  - AUROC = 0.5: Random classifier
  - AUROC = 1.0: Perfect classifier

### One-vs-Rest: Macro vs. Micro Averaging

- Micro-Averaging:** Averages over all individual predictions (favors larger classes)
- Macro-Averaging:** Averages over metrics from each class equally

## K-Nearest Neighbor Algorithm (KNN)

### Intuition behind KNN:

- Label of an unseen data point  $x$  derives from the labels of the  $k$ -nearest neighbors of  $x$ .
- Similar data points  $\rightarrow$  similar labels.
- Required:** Notion of similarity/distance.

### KNN for Classification

- Training:** Remember training data.
- Prediction for unseen point  $x_i$** 
  - Calculate distances to all training data points  $x_j$ , e.g.:

$$d(p, q) = \sqrt{\sum_i^d (q_i - p_i)^2} \quad (1)$$

- Get the  $k$ -nearest neighbors.
- Label of  $x_i$  = **most frequent label among all  $k$ -nearest neighbors.**

### KNN for Classification — Distance Metrics

- Examples for different distance metrics**
  - Euclidean distance:  $d(p, q) = \sqrt{\sum_{i=1}^d (q_i - p_i)^2}$
  - Manhattan distance:  $d(p, q) = \sum_{i=1}^d |q_i - p_i|$
  - Chebyshev distance:  $d(p, q) = \max_i |q_i - p_i|$
- Other metrics**
  - Cosine similarity:  $\text{sim}(p, q) = \frac{\sum_{i=1}^d p_i q_i}{\|p\| \|q\|}$
  - Jaccard similarity:  $\text{sim}(A, B) = \frac{|A \cap B|}{|A \cup B|}$

### Common Outcomes

- Different  $k$  values yield different results:
  - Very small  $k$ :** Sensitive to noise and outliers.
  - Large  $k$ :** Insufficient capacity to properly separate.
  - Very large  $k$ :** Most frequent class in training data dominates.

### KNN for Regression

- Training:** Remember training data.
- Prediction for unseen point  $x_i$** 
  - Calculate distances to all training data points  $x_j$ , e.g.:

$$d(p, q) = \sqrt{\sum_i^d (q_i - p_i)^2} \quad (2)$$

- Get the  $k$ -nearest neighbors.
- Label of  $x_i$  = **mean of scores of all  $k$ -nearest neighbors.**

## K-Fold Cross Validation & Model Evaluation

### Extended Training Setup

- Building a good classifier or regressor**
  - Find the best data preprocessing steps.
  - Find the best model (model selection).
  - Find the best hyperparameter values.
  - Iterative evaluation required.**
- Important: Not allowed to use test data for this!**
  - Test data is meant to be **unseen**.
  - Using test data during optimization makes it no longer unseen.

### Training & Evaluation Using Validation Data

- Common data split to ensure generalizability**
  - Use **test data** only at the end for final evaluation.
  - Use **validation data** for model selection and hyperparameter tuning.

### K-Fold Cross Validation

- Core idea**
  - Split training data into  **$k$  blocks** of equal size (e.g.,  $k = 10$ ).
  - Use ( **$k-1$** ) **blocks** for training and the remaining block for evaluation.
  - Repeat for  **$k$  rounds** with different permutations.

- Advantages**
  - Averaged results** are more reliable.
  - Variance between results** is a useful indicator.

### Information Leakage through Normalization

- Important guidelines**
  - Do not normalize before splitting** into training & test data!
  - Normalize training & test data, but only based on training data!**

## Decision Tree

### Decision Tree — Idea

- Represent mapping between features and label/value as a flowchart-like structure.

### Components

- (Inner) node — test on a single feature.
- Branch — outcome of a test; corresponds to a feature value or range of values.
- Leaf — label (classification) or real value (regression).

### Same Dataset, Different Decision Tree

- In general, there are multiple trees that match a dataset.

### Decision Tree — Diversity

- Different branching factors.
- Different depth:
  - Leaves can have more than one label or real value.
  - Required based on dataset or choice (pruning).
  - Final output: majority label (classification) or mean of values (regression).

### Building a Decision Tree

- Notations:
  - $D_t$  — set of records that reach node  $t$ .
  - $D_0$  — set of all records at root node.
- General procedure:
  - If  $|D_t| = 1$  or all records in  $D_t$  have the same class or value  $\Rightarrow t$  is a leaf node.
  - Otherwise, choose a test (feature + conditions) to split  $D_t$  into smaller subsets (subtrees).
  - Recursively apply the procedure to each subtree.

### How to Split? — Nominal Attributes

- Binary split: Partition all  $d$  values into two subsets.  $\frac{2^d - 2}{2}$  possible splits.
- Multiway split: Each value yields a subtree. Arbitrary splits into  $2 \leq s \leq d$  subtrees are possible, but the number of possible splits grows exponentially.

### How to Split? — Ordinal Attributes

- Binary split: Partition all  $d$  values into two subsets. Partitions must preserve the natural order of values.  $d - 1$  possible splits.
- Multiway split: Each value yields a subtree.

### Finding the Best Splits — Approaches

- Global Optimum:
  - Best splits = splits that result in a Decision Tree with the highest accuracy.
  - Problem: Finding the optimal tree is NP-complete  $\Rightarrow$  not practical for large datasets.
- Greedy Approach:
  - Faster heuristics but no guarantees for optimal results.
  - Pick the split that minimizes the impurity of subtrees (w.r.t. class labels).

### Finding the Best Splits — General Procedure

- Calculate impurity  $I(t)$  of node  $t$  before splitting.
- For each possible split, calculate impurity of split  $I_{split}$ .
- Select split with lowest impurity  $I_{split}$ .
- Perform split if  $I_{split} < I(t)$  (not necessarily always the case).

### Impurity of a Node (Classification)

- Gini Index:

$$Gini(t) = 1 - \sum_{c \in C} P(c|t)^2$$

- Entropy:

$$Entropy(t) = - \sum_{c \in C} P(c|t) \log P(c|t)$$

- $P(c|t)$  = relative frequency of class  $c$  in node  $t$ .

### Impurity of a Split (Classification)

- Assume node  $t$  is split into  $k$  children:
  - $n_i$  — number of records at  $i$ -th child.
  - $n$  — number of records at node  $t$ .
  - $I(i)$  — impurity of node (e.g., Gini, Entropy).
- Impurity of split:

$$I_{split} = \sum_i \frac{n_i}{n} I(i)$$

- Information Gain:

$$IG = I(t) - I_{split}$$

- Choose split that minimizes  $I_{split}$  = split that maximizes  $IG$ .
- Required condition:  $IG > 0$ .

### Impurity of a Split (Regression)

- Residual Sum of Squares (RSS):

$$RSS_{split} = \sum_{k=1}^K \sum_{i \in R_k} (y_i - \mu_{R_k})^2$$

### Decision Trees — Pros & Cons

- **Pros:**
  - Inexpensive to train and test.
  - Easy to interpret (if tree is not too large).
  - Can handle categorical and numerical data.
- **Cons:**
  - Sensitive to small changes in the training data.
  - Greedy approach does not guarantee the optimal tree.
  - Each decision involves only a single feature.
  - Does not take interactions between features into account.

### Decision Trees — Overfitting

- Decision Tree algorithm can always split the training data perfectly.
- Solution: Limit size/height of Decision Tree  $\Rightarrow$  Pruning.

### Pre-Pruning

- **Maximum Depth**
  - Define the maximum depth/height of the tree.
  - Stop splitting if the maximum depth is reached.
- **Minimum Sample Count**
  - Define the minimum number of samples each node must have.
  - Stop splitting if a node has fewer than the minimum number of samples.
  - Example: Minimum sample count = 16.

## Bagging (Bootstrap Aggregation)

### Bagging — Basic Idea (not limited to Decision Trees)

- Train many models (classifiers/regressors) on different training data.
- Combine predictions of each model for the final prediction.
- Increases accuracy and lowers variance.

### Bootstrap Sampling

- Take repeated samples from a single training dataset  $D$ .
- Bootstrap sample  $D_i$  sampled from  $D$ , uniformly and with replacement ( $|D_i| = |D|$ ).
- Train a model over each bootstrap dataset  $D_i$ .

### Limitations and Consequences

- Strong predictors in  $D$  will likely be in most bootstrap samples  $D_i$ .
- Most bagged trees will use strong predictors on top.
- Most bagged trees will look very similar.
- Predictions of bagged trees will be highly correlated.
- **Only limited reduction in variance!**

## Random Forest

### Random Forest = Bootstrap Sampling (Bagging) + Feature Sampling

- Create bootstrap samples  $D_i$  like for bagging.
- Feature sampling: For each  $D_i$ , consider only a random subset of features of size  $m$ .
- Typically,  $m \approx \sqrt{d}$ .

### Effects of Feature Sampling

- Strong predictors in  $D$  are often absent in  $D_i$ .
- Resulting trees often look very different.
- Predictions of trees are much less correlated.
- **Higher reduction in variance + typically higher accuracy.**

### Pros and Cons of Random Forests (Compared to Decision Trees)

- **Pros:**
  - High accuracy—fairly close to state of the art.
  - Sampling and training independent across  $D_i \rightarrow$  parallelizable!
  - Not much tuning required.
- **Cons:**
  - Less interpretable.
  - Slower training and prediction.

## Boosting

### Bagging vs. Boosting

- **Bagging:** Trees are trained independently (can be done in parallel).
- **Boosting:** Trees are trained in sequence; the accuracy of the last tree affects the training of the next tree.
- **Bagging:** All trees have the same amount of say in the final prediction.
- **Boosting:** Trees have different amounts of say in the final prediction (depending on their individual accuracy).

### Boosting and Weak Learners

- **Strong Learners:** Perform as best as possible on a given classification or regression task.
- **Weak Learners:** Perform slightly better than guessing.
  - A common weak learner: **Decision Stump** (decision tree of height 1, i.e., only one split).
  - Very simple model  $\rightarrow$  very fast training.
- **Boosting:** Combine many weak classifiers into a single strong learner.
  - Basic idea: subsequent models try to improve the errors of previous models.

## AdaBoost (Adaptive Boosting)

- Applicable to many classification/regression algorithms to improve performance.
- Very commonly combined with Decision Trees.

### Basic Training Algorithm for AdaBoost

- Train a Weak Learner over  $D_i$  (e.g., Decision Stump).
- Identify all misclassified samples.
- Calculate the error rate of the learner to quantify its amount of say.
- Sample  $D_{i+1}$  such that misclassified samples are more likely to be picked than correctly classified samples.
- Repeat...

## Gradient Boosted Trees

- **Gradient Boosting**
  - Mainly applied to regression algorithms to improve performance.
  - Very commonly combined with Decision Trees (for regression).
- **Basic training algorithm**
  - Start with an initial prediction (e.g., mean over all values).
  - Calculate residuals = error between true value and current prediction.
  - Train a Decision Stump to predict residuals.
  - Update predictions based on predicted residuals.
  - Repeat...

### Gradient Boosting

- Mainly applied to regression algorithms to improve performance.
- Very commonly combined with Decision Trees (for regression).

### Basic Training Algorithm for Gradient Boosting

- Start with an initial prediction (e.g., mean over all values).
- Calculate residuals = error between true value and current prediction.
- Train Decision Stump to predict residuals.
- Update predictions based on predicted residuals.
- Repeat...

### Boosting Methods — Pros & Cons (Compared to Decision Trees)

- **Pros**
  - High accuracy — often state of the art.
- **Cons**
  - Less interpretable (arguably even less compared to Random Forests).
  - Slower training and prediction  $\rightarrow$  sequential processing  $\rightarrow$  not parallelizable.

## Algorithm 5 AdaBoost Training Algorithm

```
1: Initialization: Dataset  $D$ ,  $|D| = N$ , with initial sample weights  $w_i = \frac{1}{N}$ 
2: for  $m = 1$  to  $M$  do
3:   Generate  $D_m$  by sampling from  $D$  w.r.t. sampling weights  $w$ 
4:   Train Decision Stump  $h_m$  over  $D_m$ 
5:   Apply  $h_m$  to all samples in  $D$  and identify misclassified samples
6:   Calculate total error:  $\epsilon_m = \sum_i^N w_i \cdot \delta(h_m(x_i) \neq y_i)$ 
7:   Calculate amount of say:  $\alpha_m = \frac{1}{2} \ln \frac{1-\epsilon_m}{\epsilon_m}$ 
8:   Update sample weights:
```

$$w_i = w_i \cdot \begin{cases} e^{\alpha_m}, & \text{if } x_i \text{ was misclassified} \\ e^{-\alpha_m}, & \text{if } x_i \text{ was correctly classified} \end{cases}$$

$$w_i = \frac{w_i}{\sum_i^N w_i}$$

9: **end for**

## Algorithm 6 Gradient Boosting Training Algorithm

```
1: Initialization: Dataset  $D$ ,  $f_0(x_i) = \text{mean}(y)$ ,  $\eta = 0.1$ 
2: for  $m = 1$  to  $M$  do
3:   Calculate residuals:  $r_{i,m} = y_i - f_{m-1}(x_i)$ 
4:   Train Decision Stump  $h_m$  over  $D$  with  $r_{i,m}$  as targets
5:   Predicted residuals  $\hat{r}_m(x_i)$  for all training samples
6:   Calculate new predictions:  $f_m(x_i) = f_{m-1} + \eta \cdot h_m(x_i)$ 
7: end for
```

## Algorithm 7 Gradient Boosting Training

```
1: Initialization: Dataset  $D$ ,  $f_0(x_i) = \text{mean}(y)$ ,  $\eta = 0.1$ 
2: for  $m = 1$  to  $M$  do
3:   Calculate residuals:  $r_{i,m} = y_i - f_{m-1}(x_i)$ 
4:   Train Decision Stump  $h_m$  over  $D$  with  $r_{i,m}$  as targets
5:   Predict residuals  $\hat{r}_m(x_i)$  for all training samples
6:   Calculate new predictions:  $f_m(x_i) = f_{m-1} + \eta \cdot h_m(x_i)$ 
7: end for
8: Output:  $M$  Decision Stumps  $h_1, h_2, \dots, h_M$ 
```

## Lecture 7

### Linear Models

#### Basic setup

- Dataset of  $n$  samples:  $\{(x_i, y_i)\}_{i=1}^n$
- Input data with  $d$  features:  $x_i = (x_{i1}, x_{i2}, \dots, x_{id})$

#### Assumption

- There exists a linear relationship between  $x_i$  and the dependent variable  $y_i$ :

$$\hat{y}_i = h_\theta(x_i) = f(\theta_0 x_{i0} + \theta_1 x_{i1} + \theta_2 x_{i2} + \dots + \theta_d x_{id})$$

$$\theta = \{\theta_0, \theta_1, \theta_2, \dots, \theta_d\}, \quad \theta_i \in \mathbb{R}$$

#### Vector Representation

- Introduce constant feature  $x_{i0}$

$$h_\theta(x_i) = f(\theta_0 x_{i0} + \theta_1 x_{i1} + \theta_2 x_{i2} + \dots + \theta_d x_{id})$$

- Represent  $x_i$  with new constant feature:

$$x_i = (1, x_{i1}, x_{i2}, \dots, x_{id})$$

- Rewrite linear relationship using vectors:

$$h_\theta(x_i) = f(\theta^T x_i)$$

#### Linear Regression

- Regression**  $\Rightarrow$  Real-valued predictions
- Function  $f$  is the identity function  $f(x) = x$

$$\hat{y}_i = h_\theta(x_i) = f(\theta^T x_i) = \theta^T x_i$$

$$\hat{y} = X\theta$$

$$\hat{y} = \begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \vdots \\ \hat{y}_n \end{bmatrix}, \quad X = \begin{bmatrix} 1 & x_{11} & x_{12} & \dots & x_{1d} \\ 1 & x_{21} & x_{22} & \dots & x_{2d} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n1} & x_{n2} & \dots & x_{nd} \end{bmatrix}, \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_d \end{bmatrix}$$

#### Linear Regression — Loss Function

- Loss function** (also: cost function, error function)
  - Quantifies how good or bad a given set of values for  $\theta$  is.
  - Measures the difference between predictions  $\hat{y}$  and true values  $y$ .
- Mean Squared Error (MSE)**

$$L = \frac{1}{n} \sum_{i=1}^n e_i^2 = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

#### Find Minimum of $L$ Analytically

- Minimum of loss function  $L \Rightarrow$  Calculus to the rescue!
  - Partial derivatives w.r.t. to all  $\theta_i$  are 0:

$$\frac{\partial L}{\partial \theta_0} = 0, \quad \frac{\partial L}{\partial \theta_1} = 0, \quad \dots, \quad \frac{\partial L}{\partial \theta_d} = 0$$

—  $d + 1$  equations with  $d + 1$  unknowns

#### Rewrite loss function $L$

- Vector representation makes it easier to handle:

$$L = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2 = \frac{1}{n} \|X\theta - y\|^2$$

- Derive  $L$  w.r.t.  $\theta$ :

$$\frac{\partial L}{\partial \theta} = \frac{2}{n} X^T (X\theta - y)$$

- Set  $\frac{\partial L}{\partial \theta} = 0$ :

$$\frac{2}{n} X^T (X\theta - y) = \vec{0}$$

#### Linear Regression — Normal Equation

- Solve for  $\theta$ :

$$X^T X \theta = X^T y$$

$$(X^T X)^{-1} X^T X \theta = (X^T X)^{-1} X^T y$$

$$\theta = (X^T X)^{-1} X^T y$$

$$\theta = X^\dagger y, \quad \text{with } X^\dagger = (X^T X)^{-1} X^T$$

#### Pseudo Inverse $X^\dagger$

- $X^\dagger = (X^T X)^{-1} X^T$
- Performance analysis:
  - Most expensive operation: calculating the inverse of  $(X^T X)^{-1}$
  - Calculation of inverse depends on number of features  $d$ , not on number of data samples  $n$
  - Complexity of calculating inverse of a  $d \times d$  matrix:  $\mathcal{O}(d^3)$

#### Analytical Solving Algorithm

- Construct matrix  $X$  and vector  $y$  from data
- Calculate pseudo inverse  $X^\dagger = (X^T X)^{-1} X^T$
- Return  $\theta = X^\dagger y$

#### Gradient Descent

- Core idea**
  - Start with a random setting of  $\theta$
  - Adjust  $\theta$  iteratively to minimize  $L$

#### Gradient

- Vector of partial derivatives of a multivariable function (e.g.,  $\theta_0, \theta_1, \dots, \theta_d$ )
- Partial derivative: slope w.r.t. to a single variable given a current set of values for all  $\theta_0, \theta_1, \dots, \theta_d$
- Points in the direction of the steepest ascent

$$\nabla_\theta L = \frac{\partial L}{\partial \theta} = \begin{bmatrix} \frac{\partial L}{\partial \theta_0} \\ \frac{\partial L}{\partial \theta_1} \\ \frac{\partial L}{\partial \theta_2} \\ \vdots \\ \frac{\partial L}{\partial \theta_d} \end{bmatrix}$$

- Gradient Descent Algorithm
- Important concept: learning rate
    - Scaling factor for gradient (typical range: 0.01 - 0.0001)

Algorithm 8 Gradient Descent Algorithm

- 1: **Input:** Data  $(X, y)$ , loss function  $L$ , learning rate  $\eta$   
2: **Initialization:** Set  $\theta$  to random values  
3: **while** not converged **do**  
4:     Compute gradient:

$$\nabla_{\theta} L = \frac{\partial L}{\partial \theta}$$

- 5:     Update parameters:

$$\theta \leftarrow \theta - \eta \cdot \nabla_{\theta} L$$

- 6: **end while**  
7: **Output:** Optimal  $\theta$

*In practice: stop loop when  $\theta$  converges*

Gradient Descent — Variations

- (Basic) Gradient Descent
  - Calculate gradient and update  $\theta$  for the whole dataset
- Stochastic Gradient Descent (SGD)
  - Calculate gradient and update  $\theta$  for each data sample
- Mini-batch Gradient Descent
  - Calculate gradient and update  $\theta$  for batches of samples
  - e.g., batch = 64 data samples
  - In practice, often referred to as SGD

Normal Equation vs. Gradient Descent

- Gradient Descent
  - Works well even if  $d$  is large
  - Works even if  $X^T X$  is non-invertible
  - Iterative process; may not find optimal solution in practice
  - Learning rate is a critical hyperparameter
- Normal Equation
  - Finds optimal solutions
  - Non-iterative; no need for a learning rate
  - Calculation of  $(X^T X)^{-1}$  in  $\mathcal{O}(d^3)$

Polynomial Linear Regression

- Linear Regression  $\nrightarrow$  line / plane / hyperplane
- Polynomial Linear Regression
  - Allows capturing nonlinear relationships between  $X$  and  $y$
  - Polynomial regression model for 1 input feature

$$\hat{y}_i = \theta_0 1 + \theta_1 x_i + \theta_2 x_i^2 + \dots + \theta_p x_i^p$$

Matrix representation (again, 1 input feature!)

$$X^{(1)} = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix} \quad X^{(2)} = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & \vdots & \vdots \\ 1 & x_n & x_n^2 \end{bmatrix} \quad X^{(3)} = \begin{bmatrix} 1 & x_1 & x_1^2 & x_1^3 \\ 1 & x_2 & x_2^2 & x_2^3 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & x_n & x_n^2 & x_n^3 \end{bmatrix}$$

Polynomial Linear Regression — Overfitting

- Increasing degree of polynomial  $p$ 
  - More capacity to capture nonlinear relationships
  - Much higher sensitivity to noise and outliers
- Countermeasure: Regularization
  - Extend loss function to "punish" large values of  $\theta$

$$L = \frac{1}{n} \|X\theta - y\|^2 + \lambda \frac{1}{n} \|\theta\|_2^2$$

$$\|\theta\|_2^2 = \sum_{i=1}^d \theta_i^2$$

*Note: excludes  $\theta_0$ !*

Polynomial Linear Regression — Minimizing Loss  $L$

- Normal Equation

$$\theta = (X^T X + \lambda \begin{bmatrix} 0 & & \\ & 1 & \\ & & \ddots \\ & & & 1 \end{bmatrix})^{-1} X^T y$$

- Gradient Descent

$$\nabla_{\theta} L = \frac{2}{n} X^T (X\theta - y) + \lambda \frac{2}{n} \theta$$

Polynomial Linear Regression — More than 1 Feature

- Number of terms in multivariate polynomial given  $p, d$

$$\theta_i, \quad 0 \leq i \leq M, \quad \text{with } M = \binom{p+d}{p}$$

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}$$

- Practical considerations
  - Limited to small number of features and small polynomial degrees
  - In principle, terms can be dropped (e.g., all interaction terms)

Linear Regression — Interpretation of  $\theta_i$

- Change of value of feature  $i$  by 1 unit  $\Rightarrow$  change of output value by  $\theta_i$
- Assumption: all other feature values remain the same

Data Normalization — Yes or No? (standardization, min-max scaling)

Data normalization does not affect model performance (assuming basic Linear Regression without regularization)

In favor of "No"

- Preserves unit of feature  $i \Rightarrow$  direct interpretation of  $\theta_i$
- Better for comparing  $\theta_i$  for the same features across different datasets

In favor of "Yes"

- When using regularization
- When using Polynomial Linear Regression
- Better for comparing  $\theta_i$  within a model

$$(\theta_i > \theta_j \Rightarrow \text{feature } i \text{ more important than feature } j)$$

Logistic Regression

- Logistic Regression  $\Rightarrow$  Real-valued predictions interpreted as probability
- Function  $f$  is the standard **Logistic Function** (Sigmoid function)

$$f(x) = \frac{L}{1 + e^{-k(x-x_0)}}$$

For  $L = 1, k = 1, x_0 = 0$ :

$$f(x) = \frac{1}{1 + e^{-x}}$$

Logistic Regression: Probabilistic Interpretation

- $\hat{y}$  interpreted as a probability

$$\hat{y} = h_{\theta}(x) = f(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}, \quad \hat{y} \in [0, 1]$$

- $\hat{y} = h_{\theta}(x)$  is the estimated probability that  $y_i = 1$  given  $x$  and  $\theta$

$$\hat{y} = P(y = 1|x, \theta)$$

Given only discrete 2 outcomes:

$$P(y = 1|x, \theta) + P(y = 0|x, \theta) = 1$$

$$\hat{y} = 1 - P(y = 0|x, \theta)$$

$$\hat{y} = P(y = 1|x, \theta) = 1 - P(y = 0|x, \theta)$$

- $P(y|x)$  is a Bernoulli distribution

$$P(y|x) = \begin{cases} \hat{y}, & y = 1 \\ 1 - \hat{y}, & y = 0 \end{cases}$$

$$P(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

Logistic Regression: Loss Function

Goal: Maximize probability of true  $y$  label given training sample  $x$

- Find  $\theta$  that maximizes

$$P(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

$$\log P(y|x) = \log \left[ \hat{y}^y (1 - \hat{y})^{1-y} \right] = y \log \hat{y} + (1 - y) \log(1 - \hat{y})$$

- Find  $\theta$  that minimizes

$$L = -P(y|x) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$



- Cross-Entropy Loss**  
**Logistic Regression: Loss Function**
- Loss for all training samples

$$L = -\frac{1}{n} \sum_{i=1}^n [y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)]$$

$$= -\frac{1}{n} \sum_{i=1}^n [y_i \log h_{\theta}(x_i) + (1 - y_i) \log(1 - h_{\theta}(x_i))]$$

$$= -\frac{1}{n} \sum_{i=1}^n \left[ y_i \log \frac{1}{1 + e^{-\theta^T x_i}} + (1 - y_i) \log \left( 1 - \frac{1}{1 + e^{-\theta^T x_i}} \right) \right]$$

**Gradient of the Loss Function**

- After lots of tedious math...

$$\frac{\partial L}{\partial \theta_j} = \frac{1}{n} \sum_{i=1}^n [h_{\theta}(x_i) - y_i] x_{i,j}$$

$$\nabla_{\theta} L = \frac{1}{n} X^T (h_{\theta}(X) - y)$$

- Problem:

$$\frac{1}{n} X^T (h_{\theta}(X) - y) \neq 0$$

has no closed-form solution for  $\theta$

⇒ **Gradient Descent!**

**Polynomial Logistic Regression**

- Analogous to Polynomial Linear Regression
  - Allows to capture nonlinear relationships between  $X$  and  $y$
  - Polynomial Logistic Regression model for 1 input feature

$$\hat{y}_i = \frac{1}{1 + e^{-\theta^T x_i}}, \quad \text{with } \theta^T x_i = \theta_0 + \theta_1 x_i + \theta_2 x_i^2 + \dots + \theta_p x_i^p$$

**Identical practical considerations**

- Solve using Gradient Descent as usual (optionally with regularization to avoid overfitting)
- Limited to small number of features and small polynomial degrees

$$\nabla_{\theta} L = \frac{1}{n} X^T (h_{\theta}(X) - y) + \frac{\lambda}{n} \theta$$

**Lecture 8**

## Recommender System

**Motivation**

- In the online world: information and item overload
  - Too many items: products, songs, movies, news articles, restaurants, etc.
  - More choices require better filters ⇒ recommendation engines

**Manual Recommendations — Pros & Cons**

- Pros**
  - Semantically rich (ratings, plain text, images, videos, etc.)
  - Explainability / Interpretability
- Cons**
  - Manual effort — What is the incentive for writing a review?
  - Lack of personalization

**Simple Aggregations**

- Rank items based on aggregated
- Pros**
  - Relatively easy to compute (typically weighted aggregated based on different factors)
  - Typically good/safe recommendations (particularly for new/unknown users)
- Cons**
  - Requires sufficient number of ratings per item
  - High risk of popularity bias; lack of diversity ("rich get richer" effects, "few get richer" effects)
  - Lack of personalization

**Personalized Recommendations**

- Users have different preferences that define the relevance of items
  - Preferences = interests, likings, needs, wants, desires, etc.
  - Relevant items = items that match users' preferences best

**Basic Setup**

- Set of users  $U = \{u_1, u_2, \dots, u_n\}$
- Set of items  $V = \{v_1, v_2, \dots, v_m\}$
- Rating matrix  $R$  with  $|U|$  rows and  $|V|$  columns
- Matrix element  $R_{u,v}$ :  $u$ 's rating of  $v$  (e.g., 1-5 stars, binary 0/1)

**Personalized Recommendations — Core Tasks**

- Collect ratings  $R_{uv}$
- Infer missing values  $\hat{R}_{uv}$ 
  - In practice, mainly interested in high values
  - Algorithmic component of recommender systems
  - Wide range of existing approaches
- Evaluation: How good are the recommendations?
  - Compare  $\hat{R}_{uv}$  with true  $R_{uv}$  from test set

**Collecting Ratings**

- Explicit**
  - Ask/invite/encourage users to rate items
  - Pay users to rate items (e.g., crowdsourcing)
- Implicit — derive ratings from users' behavior, e.g.:**
  - Product bought
  - Video watched
  - Article read
  - Link clicked
  - ...
- Key challenge:** Rating matrix  $R$  is in practice very sparse!

**Evaluation**

- Split  $R$  into training and test set
- Performance metrics:
  - Root Mean Squared Error (for numerical ratings)
  - Precision, Recall, F1 score, etc. (TP, TN, FP, FN for binary ratings or binary recommendation after converting numerical ratings)
  - Precision@k, Recall@k (precision and recall w.r.t. the top-k highest predicted ratings)
  - Compare rankings induced by  $\hat{R}_{uv}$  and  $R_{uv}$  with  $(u, v) \in S$

**Recommendations Using Association Rules**

- User preferences & likings**
  - Items: movies, songs, books, etc.
  - Transaction: viewing/listening/reading history
- Interesting rules (movies)**
  - Viewer who watched movies  $\{a, b\}$  also watched movies  $\{x, y\}$
  - Example:  $\{Jaws\} \rightarrow \{It\}$
- Limitations**
  - Basic AR algorithm ignores ratings
  - Popularity bias: user with very unique tastes likely to get subpar recommendations

**Recommendations Using Clustering**

- Approach**
  - Cluster movies based on "useful" features (genre, director, writer, length, ...)
  - Recommend movies from clusters with movies a user has rated highly
- Limitations**
  - Find good features in practice is very difficult
  - Unsystematic: no well-defined process to pick recommendations

**Recommendations Using Regression (or Classification)**

- Example approach: Linear Regression**
  - Independent variable: movie features
  - Dependent variable: user rating
- Build a linear regression model for each user
- Limitations**
  - Requires good features for each item
  - Cold-start problem: requires a lot of user ratings to build a good model

**Content-Based Recommender System**

**Intuition**

- Recommend item  $v$  to user  $u$  that are similar to  $v$  and  $u$  has rated highly.
- Examples:
  - Movies of the same genre.
  - Songs from the same artist.
  - Articles about the same topic.
  - Products with similar features.

**Basic Requirement: Item Profiles = Feature Vector for Each Item**

- Movie: genre, director, writer, cast, length, year, ...
- Product: type, brand, price, weight, color, ...
- Article: set of (important) words / tf-idf vector / ...

**Simple Approach — Pairwise Item Similarity**

- Pairwise item similarity  $sim(x, y)$ 
  - $x, y$  — feature vectors of movies.
  - Common metric: cosine similarity:

$$sim(x, y) = \cos(\theta) = \frac{x \cdot y}{\|x\| \|y\|}$$

**Limitation: Requires Reference Item**

- Movie(s) the user was most recently watching.
- Movie(s) the user has rated the highest.
- Movie(s) the user is currently browsing.

**User-Item Similarities**

**Needed: User Profiles = Feature Vector for Each User**

- Requirement: same shape as item profiles to calculate similarities.
- Approach: user profile = "some aggregate" over item profiles rated by the user.

**User-Item Similarities — Binary Utility Matrix**

- $R_{uv} \in \{0, 1\}$  — for example,  $R_{uv} = 1$  if:
  - User  $u$  bought movie  $v$ .
  - User  $u$  watched movie  $v$ .
- Implicit rating that  $u$  likes  $v$  (no explicit ratings available here; but also no implicit dislikes!).

**Simple Average**

- Computes average across binary ratings:

$$\text{User profile} = \frac{1}{|M|} \sum_{v \in M} \text{feature vector of } v$$

where  $M$  is the set of movies the user has rated.

User-Item Similarities — Real-Valued Utility Matrix

- $R_{uv} \in \mathbb{R}$  — for example,  $R_{uv} \in \{1.0, 1.5, 2.0, \dots, 5.0\}$  star rating.
- Explicit rating of user  $u$  for movie  $v$ .
- Important: semantic interpretation — ratings express both likes and dislikes.
- Use rating as weights for features for a weighted aggregation:

$$w_f = \frac{\sum_{v \in M} R_{uv} \cdot f_v}{\sum_{v \in M} R_{uv}}$$

where  $w_f$  is the weight of feature  $f$ , and  $f_v$  is the feature value for item  $v$ .

Intuition

- The user likes romantic and animated movies.
- The user dislikes fantasy and adventure movies.

Step 1: Normalize Ratings

- Subtract average user rating from each movie rating:

$$R'_{uv} = R_{uv} - \frac{1}{|M|} \sum_{v \in M} R_{uv}$$

- Converts ratings into positive (liked) and negative (disliked) scale.
- Distinguishes "generous" users (mostly rate highly and 3.0 is a low rating) from more "grumpy" users (mostly rate low and 3.0 is a high rating).

Step 2: Calculate Weighted Features for User Profile

- The weights are the normalized weights:

$$w_f = \frac{\sum_{v \in M} R'_{uv} \cdot f_v}{|M|}$$

User-Item Similarities

Pairwise Item Similarity  $\text{sim}(u, v)$

- $u$  — user profile;  $v$  — item profile.
- Suitable metric: cosine similarity:

$$\text{sim}(u, v) = \frac{u \cdot v}{\|u\| \|v\|}$$

Recommend items  $v_i$  to user  $u$  with max. similarities  $\text{sim}(u, v_i)$

Practical Considerations

- Top  $k$  most similar items always the same  $\Rightarrow$  add some randomization for diversity.
- Top  $k$  most similar items might include items the user has already rated  $\Rightarrow$  remove those items.
- More sophisticated ways to aggregate item profiles to user profiles conceivable (e.g., ignore underrepresented features).

Content-Based Recommender System — Pros & Cons

Pros

- Recommendations for user  $u$  do not depend on other users.
- Recommendations can also include new or unpopular items.
- Good explainability (features that had most effect on the high similarity).

Cons

- Cold-start problem: How to build a profile for new users?
  - Naive approach: recommend generally popular items to new users.
- Finding good features (and values!) for items is a non-trivial task.
- Overspecialization: By default, no recommendations outside a user's profile.
  - In practice: add some randomization into the recommendation process.

Collaborative Filtering

- **Idea: Utilize the opinions of others**
  - Recommend items that other users with similar tastes/preferences/needs have liked
  - Does not require item or user-specific features
- **Two perspectives**
  - User-based — two users are similar if they rated the same items similarly
  - Item-based — two items are similar if they are equally rated by users

User-Based Collaborative Filtering — Calculating Similarities

- **Represent all users by their rating vectors**
  - Rows of rating matrix:

$$r_A = (2, 4, 5, 0, 1)^T$$

$$r_B = (1, 0, 4, 0, 2)^T$$

- **Using Cosine Similarity**
  - Example calculations:

$$\text{sim}(r_A, r_B) = 0.77$$

- **Problem with this approach**
  - Missing values (0) are treated as negative
  - All ratings are positive values
  - No explicit notion of dissimilarity (only less or more similar)

User-Based Collaborative Filtering — Normalizing Ratings

- **Idea: Normalize rating vectors**
  - Mean-centering — subtract row mean from each rating vector
  - Missing values (0) now represent the average rating
  - Bad ratings (i.e., below average) now represented by negative values
- **Cosine similarity between mean-centered vectors**
  - Uses the Pearson Correlation Coefficient
  - Example:

$$\text{sim}(r_A, r_B) = 0.78$$

$$\text{sim}(r_D, r_B) = -0.65$$

User-Based Collaborative Filtering — Predicting Ratings

- Estimated rating  $\hat{R}_{uv}$  is the weighted average of ratings from similar users

$$\hat{R}_{uv} = \frac{\sum_{w \in N} \text{sim}(u, w) \cdot R_{wv}}{\sum_{w \in N} \text{sim}(u, w)}$$

- $N$  is the set of  $k$  users most similar to  $u$  who have already rated item  $v$

Item-Based Collaborative Filtering

- Analog to user-based approach
  - Find the most similar items
  - Two items are similar if their ratings across all users are similar
- **Predicting Ratings**

$$\hat{R}_{uv} = \frac{\sum_{i \in M} \text{sim}(i, v) \cdot R_{ui}}{\sum_{i \in M} \text{sim}(i, v)}$$

- $M$  is the set of  $k$  items most similar to  $v$  that have already been rated by  $u$

Collaborative Filtering — User-Based vs. Item-Based

- In theory, user-based and item-based are dual approaches
- In practice, item-based typically outperforms user-based
  - Items are "simpler" than users
  - Items can be more easily described
  - Users can have very varied tastes
- Item-item similarity is typically more meaningful.

Model-Based Collaborative Filtering

- **Latent Factor Models**
  - Latent representation:  $k$ -dimensional vector for each user  $u$  and item  $v$
  - Learn latent representations from the data
  - Estimate unknown ratings:

$$\hat{R}_{uv} = w_u^T h_v$$

- **Matrix Factorization Approach**
  - Put all user vectors into a matrix  $W$
  - Put all item vectors into a matrix  $H$
  - Find  $W, H$  such that:

$$R \approx WH$$

Finding Matrices  $W, H$

- **Minimize loss function**

$$L = \sum_{R_{uv} > 0} e_{uv} = \sum_{R_{uv} > 0} (R_{uv} - \hat{R}_{uv})^2 = \sum_{R_{uv} > 0} (R_{uv} - w_u^T h_v)^2$$

- **With regularization**

$$L = \sum_{R_{uv}} (R_{uv} - w_u^T h_v)^2 + \lambda (\|w_u\|^2 + \|h_v\|^2)$$

- **Using Gradient Descent**
  - Compute gradients:

$$\frac{\partial e_{uv}}{\partial w_u} = -2(R_{uv} - w_u^T h_v) h_v + 2\lambda w_u$$

$$\frac{\partial e_{uv}}{\partial h_v} = -2(R_{uv} - w_u^T h_v) w_u + 2\lambda h_v$$

- Update rules:

$$w_u \leftarrow w_u - \eta \frac{\partial e_{uv}}{\partial w_u}$$

$$h_v \leftarrow h_v - \eta \frac{\partial e_{uv}}{\partial h_v}$$

Collaborative Filtering — Pros & Cons

- **Pros**
  - No need to find and create good features (e.g., genres for movies)
  - Intuitive approach
- **Cons**
  - Similarity calculations rely on sufficient number of ratings
  - Cold-start problem in case of new users or items
  - Popularity bias: users with unique tastes likely receive subpar recommendations
  - Naive implementation is expensive: finding  $k$  most similar users/items has complexity  $O(|R|)$

**Graph:** Formalism for representing *relationships* between items

**A graph is a tuple**  $G = (V, E)$

- Set of vertices (or nodes)  $V = \{v_1, v_2, \dots, v_n\}$
- Set of edges  $E = \{e_1, e_2, \dots, e_m\}$ , where an edge is a pair of vertices:  $e_i = (v_j, v_k)$

**Example:**

$$V = \{A, B, C, D\} \quad E = \{(A, B), (A, C), (C, D), (B, A), (C, B)\}$$

**Types of Graphs**

- **Directed vs Undirected:**
  - *Undirected:* Edges have no direction (e.g., Facebook friendships)
  - *Directed:* Edges have a direction (e.g., Twitter followers)
- **Weighted vs Unweighted:**
  - *Unweighted:* All edges are equal (e.g., MRT connectivity)
  - *Weighted:* Edges have weights (e.g., travel time, number of co-authored papers)
- **Cyclic vs Acyclic:**
  - *Cyclic:* Graph has at least one cycle
  - *Acyclic:* No cycles exist
- **Simple vs Multigraph:**
  - *Simple Graph:* At most one edge between a pair of nodes
  - *Multigraph:* Multiple edges allowed between same pair of nodes
- **Sparse vs Dense:**
  - *Sparse:* Relatively few edges
  - *Dense:* Many edges, close to complete graph
- **Connected Graphs:**
  - *Strongly Connected:* Directed graph where a path exists between all pairs of nodes
  - *Weakly Connected:* Underlying undirected graph is connected
  - *Disconnected:* Not all nodes are connected

**Adjacency Matrix:** A matrix  $A$  used to represent a graph. Entry  $A_{ij}$  is nonzero if there is an edge from node  $i$  to node  $j$ .

**Example (Unweighted Undirected):**

$$\text{Graph: } V = \{A, B, C, D\} \quad A = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

**Example (Weighted Directed):**

$$A = \begin{bmatrix} 0 & 2 & 3 & 0 \\ 3 & 0 & 0 & 0 \\ 0 & 5 & 0 & 10 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

**Community Detection** No formal definition of a community. From "Networks: An Introduction" (Mark Newman): "*Loosely stated, [community detection] is the problem of finding the natural divisions of a network into groups of vertices such that there are many edges within groups and few edges between groups. What exactly we mean by "many" or "few," however, is debatable.*"

**Wide range of application**

- Identifying groups in social networks
- Recommendation systems
- Market segmentation
- Outlier/anomaly detection

**Community Detection — Modularity**

Modularity  $Q \in [-\frac{1}{2}, 1]$  of an undirected graph  $G$  with adjacency matrix  $A$ . Measures the relative density of edges inside communities with respect to edges outside communities. Optimizing modularity is NP-hard  $\rightarrow$  practical algorithms based on heuristics.

$$Q = \frac{1}{2m} \sum_{vw} \left[ A[v, w] - \frac{k_v k_w}{2m} \right] \delta(c_v, c_w)$$

where:

- $A[v, w]$ : weight of edge between nodes  $v$  and  $w$
- $k_i$ : sum of weights of edges attached to node  $i$
- $c_i$ : community of node  $i$
- $\delta(c_v, c_w) = \begin{cases} 1 & c_v = c_w \\ 0 & \text{otherwise} \end{cases}$

**Community Detection — Louvain Algorithm**

## Algorithm 9 Louvain Algorithm

- 1: Initialization: Each node is a community
- 2: **repeat**
- 3:   Phase 1: Modularity Optimization
- 4:   **for** each node  $v$  **do**
- 5:     Check if moving  $v$  to an adjacent community improves modularity
- 6:     Move  $v$  to community that maximizes modularity
- 7:   **end for**
- 8:   Phase 2: Graph Aggregation
- 9:   Represent each community as a new node
- 10:   Update weights between new nodes
- 11: **until** no further change

**Louvain Algorithm — Remarks**

**Heuristic**

- Optimizes modularity locally on all nodes
- No guarantees for optimal modularity globally (in practice often superior to other methods)

**Performance optimization**

- Phase 1 requires computing change in modularity  $\Delta Q$
- $\Delta Q$  can be computed based on local changes in community assignments (no need to recompute modularity after each change)

**Community Detection — Girvan-Newman Algorithm**

**Divisive hierarchical approach**

- Start with whole graph representing a community
- Iteratively remove edges until the community is split into two sub-communities (can recurse)

**Criteria for removing edges: Edge Betweenness Centrality**

$$c_B(e) = \sum_{v, w \in V} \frac{\sigma(v, w|e)}{\sigma(v, w)}$$

where:

- $\sigma(v, w)$ : number of shortest paths from  $v$  to  $w$
- $\sigma(v, w|e)$ : number of those paths that pass through edge  $e$

## Algorithm 10 Girvan-Newman Algorithm

- 1: **repeat**
- 2:   Calculate  $c_B(e)$  for all  $e \in E$
- 3:   Remove edge with max  $c_B(e)$
- 4: **until** graph is split into two components

**Recursive step**

- Apply algorithm to each new component
- Stop if a component contains only a single node (or based on user-specified stopping)

**Girvan-Newman Algorithm — Remarks**

**Complexity Analysis**

- Core concept: Edge Betweenness Centrality
- Requires solving the All-Pairs Shortest Path (APSP) problem
- Time complexity depends on graph type (directed/undirected, cyclic/acyclic, weighted/unweighted, etc.)

**Min-Cut Problem**

- Given a graph  $G$ , cut  $G$  into 2 components such that the number of edges between both components is minimal.
- Example:  $|\text{Min-Cut}| = 2$  in the illustrated graph.
- Fundamental problem in graph theory  $\Rightarrow$  many existing algorithms (focus depends on directed vs. undirected, etc.).

**Karger's Algorithm**

- **Randomized** method to find Min-Cut.
- Applicable to undirected graphs with positive weights (includes unweighted graphs).

## Algorithm 11 Karger's Min-Cut Algorithm

**while**  $|V| > 2$  **do**

    Randomly pick a remaining edge  $e = (v, u)$

    Merge/contract  $v$  and  $u$  into a new node

        - Update edges to neighbors of  $v$  and  $u$

        - Remove self-loops

**end while**

**Return** edges between the final 2 nodes as Min-Cut

*Intuition:* Edges that are in the Min-Cut have a lower probability of being picked. *Runtime:*  $O(|V|^2)$  (basic), but further optimizations exist.

**Analysis**

- What is the probability that the algorithm finds the *correct* Min-Cut?
- For an undirected graph  $G = (V, E)$  with  $n = |V|$  and  $m = |E|$ :

$$\text{Average degree} = \frac{1}{n} \sum_{v \in V} \text{degree}(v) = \frac{2m}{n}$$

$$|\text{Min-Cut}| \leq \frac{2m}{n} \Rightarrow P(\text{edge is in Min-Cut}) \leq \frac{2}{n}$$

- Let  $P(\text{success}) = P(\text{final cut is Min-Cut})$ :

$$P(\text{success}) \geq \left(1 - \frac{2}{n}\right) \left(1 - \frac{2}{n-1}\right) \dots \left(1 - \frac{2}{3}\right) = \frac{2}{n(n-1)} = \left(\binom{n}{2}\right)^{-1}$$

Repeated Runs and Total Runtime

- If run  $k$  times and smallest cut chosen, what is  $P(\text{failure})$ ?

$$P(\text{failure}) = \left[ 1 - \left( \binom{n}{2} \right)^{-1} \right]^k$$

With  $k = \binom{n}{2} \ln n \Rightarrow P(\text{failure}) \leq \left( \frac{1}{e} \right)^{\ln n} = \frac{1}{n}$

- Total runtime  $O(n^2 \log n) \times O(n^2) = O(n^4 \log n)$

Remarks

- In general, a graph has multiple possible Min-Cuts.
- Choice is application-specific:
  - Favor Min-Cuts where the two components are of similar size.
  - Ignore Min-Cuts where a component size is below a threshold.

**Centrality Measures:** Quantify the **importance** of a node given its topological position in a graph. Centrality is usually assigned to nodes but can also be extended to edges (e.g., Edge Betweenness Centrality). Different measures capture different "flavours" of importance.

**What makes a node important?** That depends on the specific notion of importance the centrality measure encodes.

Applications:

- Identify influential social network users
- Identify superspreaders of diseases
- Identify key nodes in infrastructure networks

**Degree Centrality:** A local measure; counts how many edges are directly incident to a node.

Undirected Graph:

$$c_d(v_i) = \sum_{v_j \in V} A[i, j]$$

Directed Graph:

$$c_{d\_in}(v_i) = \sum_{v_j \in V} A[j, i], \quad c_{d\_out}(v_i) = \sum_{v_j \in V} A[i, j]$$

**Note:** For unweighted graphs,  $A[i, j] = 1$  for an existing edge. Thus, the sum equals the degree.

Pros:

- Simple and efficient to compute
- Often sufficient for basic applications

Cons:

- Only considers immediate neighbors
- All edges treated equally
- Vulnerable to manipulation (e.g., spam links or fake accounts)

Examples of manipulation:

- Fake websites linking to boost ranking
- Fake social media followers
- Fake reviews on e-commerce platforms

**Eigenvector Centrality:** A recursive measure—nodes are important if they connect to other important nodes. Computed using the principal eigenvector of the adjacency matrix  $A$ :

$$c_{ev}(v_i) = \frac{1}{\lambda} \sum_{v_j \in V} A[i, j] \cdot c_{ev}(v_j), \quad \lambda c_{ev} = A c_{ev}$$

Power Iteration Algorithm (to compute largest eigenvector):

Algorithm 12 Power Iteration Method

1: **Input:** Matrix  $M$ , error threshold  $\epsilon$ , max iterations  $T$

2: Initialize:  $t = 0, x_0 = [1/|V|, 1/|V|, \dots]$

3: **repeat**

4:      $t \leftarrow t + 1$

5:      $x_t \leftarrow Mx_{t-1}$

6:      $x_t \leftarrow x_t / \|x_t\|$

7:      $\delta \leftarrow \|x_t - x_{t-1}\|$

8: **until**  $\delta < \epsilon$  or  $t > T$

9: **Return**  $x_t$

Normalize

Compute change

**PageRank:** Eigenvector-like measure adapted for directed graphs (e.g., web graphs). Uses a *Random Surfer* model:

$$c_{pr}(v_i) = \alpha M c_{pr}(v_i) + (1 - \alpha) E, \quad E = \left[ \frac{1}{|V|}, \dots \right]^T$$

Where  $M$  is the transition matrix of the graph.  $M$  is column-stochastic so  $\lambda = 1$ .

Remarks on Eigenvector-Based Measures:

- Recursive and intuitive definition
- Extended models: HITS, SALSA, Katz, personalized PageRank
- Costlier to compute but highly parallelizable

**Closeness Centrality:** A node is central if it is "close" to all others (short average distance).

$$c_{cl}(v) = \frac{N}{\sum_{w \in V} d(v, w)}$$

Where  $d(v, w)$  is the shortest path length. Works on directed/undirected graphs.

**Betweenness Centrality:** Measures how often a node appears on shortest paths between other pairs:

$$c_b(v) = \sum_{s, t \in V; s \neq v \neq t} \frac{\sigma_{st}(v)}{\sigma_{st}}$$

Where  $\sigma_{st}(v)$  is the number of shortest paths from  $s$  to  $t$  that pass through  $v$ .

Remarks on Closeness & Betweenness:

- Both rely on **shortest path distances**
- Require solving the **All-Pairs Shortest Path (APSP)** problem
- Complexities vary by graph type (e.g., cyclic vs acyclic, with or without negative edges)

Summary of Popular Centrality Measures:

- **Local:** Degree, InDegree, OutDegree
- **Eigenvector-based:** Eigenvector Centrality, PageRank
- **Distance-based:** Closeness, Betweenness