

DSA5103 Optimization Problem for Data Modelling

AY2024/25 Sem2 By Zhao Peiduo

Lecture 1

Nonlinear Programming A general **nonlinear programming problem (NLP)** is to minimize/maximize a function $f(x)$, subject to equality constraints $g_i(x) = 0$, $i \in [m]$, and inequality constraints $h_j(x) \leq 0$, $j \in [p]$. Here, f , g_i , and h_j are functions of the variable $x = (x_1, \dots, x_n)^T \in \mathbb{R}^n$. The term definitions are as follows:

- f : **Objective function**
- $g_i(x) = 0$: **Equality constraints**
- $h_j(x) \leq 0$: **Inequality constraints**

It suffices to discuss minimization problems since minimizing $f(x)$ is equivalent to maximizing $-f(x)$.

Feasible Set

$$S = \{x \in \mathbb{R}^n \mid g_1(x) = 0, \dots, g_m(x) = 0, h_1(x) \leq 0, \dots, h_p(x) \leq 0\}.$$

A point in the feasible set is a **feasible solution** or **feasible point** where all constraints are satisfied; otherwise, it is an **infeasible solution** or **infeasible point**. When there is no constraint, $S = \mathbb{R}^n$, we say the NLP is **unconstrained**.

Local and Global Minimizer Let S be the feasible set. Define $B_\epsilon(y) = \{x \in \mathbb{R}^n \mid \|x - y\| < \epsilon\}$ to be the open ball with center y and radius ϵ . Here, $\|x\| = \sqrt{x_1^2 + \dots + x_n^2}$.

1. A point $x^* \in S$ is said to be a **local minimizer** of f if there exists $\epsilon > 0$ such that

$$f(x^*) \leq f(x) \quad \forall x \in S \cap B_\epsilon(x^*).$$

2. A point $x^* \in S$ is said to be a **global minimizer** of f if

$$f(x^*) \leq f(x) \quad \forall x \in S.$$

Interior point Let $S \subseteq \mathbb{R}^n$ be a nonempty set. An point $x \in S$ is called an **interior point** of S if

$$\exists \epsilon > 0 \quad s.t. \quad B_\epsilon(x) \subseteq S.$$

Gradient Vector Let $S \subseteq \mathbb{R}^n$ be a nonempty set. Suppose $f : S \rightarrow \mathbb{R}$, and x is an interior point of S such that f is differentiable at x . Then the **gradient vector** of f at x is

$$\nabla f(x) = \begin{bmatrix} \frac{\partial f}{\partial x_1}(x) \\ \vdots \\ \frac{\partial f}{\partial x_n}(x) \end{bmatrix}.$$

Hessian Matrix Let $S \subseteq \mathbb{R}^n$ be a nonempty set. Suppose $f : S \rightarrow \mathbb{R}$, and x is an interior point of S such that f has second-order partial derivatives at x . Then the **Hessian** of f at x is the $n \times n$ matrix:

$$H_f(x) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2}(x) & \frac{\partial^2 f}{\partial x_1 \partial x_2}(x) & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n}(x) \\ \frac{\partial^2 f}{\partial x_2 \partial x_1}(x) & \frac{\partial^2 f}{\partial x_2^2}(x) & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n}(x) \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1}(x) & \frac{\partial^2 f}{\partial x_n \partial x_2}(x) & \cdots & \frac{\partial^2 f}{\partial x_n^2}(x) \end{bmatrix}.$$

- The ij -entry of $H_f(x)$ is $\frac{\partial^2 f}{\partial x_i \partial x_j}(x)$.
- In general, $H_f(x)$ is not symmetric. However, if f has continuous second-order derivatives, then the Hessian matrix is symmetric since ∂x_i and ∂x_j are interchangeable.

Positive (Semi)Definite Let A be a real $n \times n$ matrix.

1. A is said to be **positive semidefinite** if $x^T A x \geq 0$, $\forall x \in \mathbb{R}^n$.
2. A is said to be **positive definite** if $x^T A x > 0$, $\forall x \neq 0$.
3. A is said to be **negative semidefinite** if $-A$ is positive (semi)definite.
4. A is said to be **negative definite** if $-A$ is positive definite.
5. A is said to be **indefinite** if A is neither positive nor negative semidefinite.

Eigenvalue Test Theorem Let A be a real symmetric $n \times n$ matrix.

1. A is **positive semidefinite** iff every eigenvalue of A is nonnegative.
2. A is **positive definite** iff every eigenvalue of A is positive.
3. A is **negative semidefinite** iff every eigenvalue of A is nonpositive.
4. A is **negative definite** iff every eigenvalue of A is negative.
5. A is **indefinite** iff it has both a positive eigenvalue and a negative eigenvalue.

Proof for: A is positive semidefinite iff every eigenvalue of A is nonnegative

(Forward) Suppose A is positive semidefinite, show that its eigenvalues are nonnegative. By definition, a Hermitian matrix A is positive semidefinite if for all nonzero vectors $x \in \mathbb{C}^n$:

$$x^* A x \geq 0$$

Let λ be an eigenvalue of A with corresponding eigenvector x such tha $Ax = \lambda x$. Taking the inner product of both sides with x , we obtain:

$$x^* A x = x^* (\lambda x) = \lambda (x^* x)$$

Since $x^* x$ (the squared norm of x) is always positive for nonzero x , the above equation implies $\lambda \geq 0$.

(Backward) Since A is Hermitian, it has an orthonormal basis of eigenvectors $\{q_1, q_2, \dots, q_n\}$ with corresponding real eigenvalues $\lambda_1, \lambda_2, \dots, \lambda_n$.

For any vector x , we can express it in terms of the eigenvectors as:

$$x = \sum_{i=1}^n c_i q_i$$

for some scalars c_i , and compute the quadratic form:

$$x^* A x = \left(\sum_{i=1}^n c_i^* q_i^* \right) A \left(\sum_{j=1}^n c_j q_j \right)$$

Expanding the expression using the orthonormality of the eigenvectors:

$$x^* A x = \sum_{i=1}^n \lambda_i |c_i|^2$$

Since we are given that all eigenvalues $\lambda_i \geq 0$, and the squared magnitudes $|c_i|^2$ are nonnegative, it follows that:

$$x^* A x \geq 0 \quad \forall x \neq 0$$

Thus, A is positive semidefinite.

Necessary and Sufficient Conditions

Assume that $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is nonlinear and differentiable. A point x^* is called a **stationary point** of f if $\nabla f(x^*) = 0$.

Necessary condition: Confine our search for global minimizers within the set of stationary points

If x^* is a local minimizer of f , then

1. x^* is a stationary point, i.e., $\nabla f(x^*) = 0$
2. The Hessian $H_f(x^*)$ is positive semidefinite

Sufficient condition: Verify that a point is indeed a local minimizer

If the following conditions hold, then x^* is a local minimizer of f .

1. x^* is a stationary point, i.e., $\nabla f(x^*) = 0$
2. The Hessian $H_f(x^*)$ is positive definite,

Convex set A set $D \subseteq \mathbb{R}^n$ is said to be a **convex** set if for any two points x and y in D , the line segment joining x and y also lies in D . That is,

$$x, y \in D \Rightarrow \lambda x + (1 - \lambda)y \in D \quad \forall \lambda \in [0, 1].$$

Strictly convex function

Let $D \subseteq \mathbb{R}^n$ be a convex set. Consider a function $f : D \rightarrow \mathbb{R}$.

1. The function f is said to be **convex** if $f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y)$, $\forall x, y \in D$, $\lambda \in [0, 1]$.
2. The function f is said to be **strictly convex** if $f(\lambda x + (1 - \lambda)y) < \lambda f(x) + (1 - \lambda)f(y)$. for all distinct $x, y \in D$, $\lambda \in (0, 1)$.

For a convex f It holds that

1. any local minimizer is a global minimizer.
2. if f is strictly convex, then the global minimizer is unique.

Test for convexity of a differentiable function

Suppose that f has continuous second partial derivatives on an open convex set D in \mathbb{R}^n .

1. The function f is convex on D iff the Hessian matrix $H_f(x)$ is positive semidefinite at each $x \in D$.
2. If $H_f(x)$ is positive definite at each $x \in D$, then f is strictly convex on D .
3. If $H_f(\hat{x})$ is indefinite at some point $\hat{x} \in D$, then f is not a convex nor a concave function on D .

Eigenvalue Decomposition: The eigenvalue decomposition of $A \in \mathbb{S}^n$ is given by:

$$A = Q \Lambda Q^T = [Q_{\cdot 1} \quad \cdots \quad Q_{\cdot n}] \begin{bmatrix} \lambda_1 & & \\ & \ddots & \\ & & \lambda_n \end{bmatrix} [Q_{\cdot 1} \quad \cdots \quad Q_{\cdot n}]^T$$

where Q is an orthogonal matrix whose **columns** are eigenvectors of A , Λ is a diagonal matrix with eigenvalues of A on the diagonal.

Change of bases using eigenvectors Denote the i th column of orthogonal matrix Q as q_i . Change the bases to $\{q_1, q_2\}$. With new bases,

- For any vector x , $x = Q(Q^T x)$, so its representation becomes

$$\tilde{x} = Q^T x = \begin{bmatrix} \tilde{x}_1 \\ \tilde{x}_2 \end{bmatrix}$$

- Since $y = Ax = Q \Sigma Q^T x$, the representation of y is

$$\tilde{y} = \Sigma \tilde{x} = \begin{bmatrix} \lambda_1 \tilde{x}_1 \\ \lambda_2 \tilde{x}_2 \end{bmatrix}$$

Hence, the linear transformation results in a scaling of λ along the eigenvector associated with λ .

Statistical Properties Let $x_1, \dots, x_n \in \mathbb{R}^p$ be n observations of a random variable x .

- Mean vector: $\mu = \bar{x} = \frac{1}{n} \sum_{i=1}^n x_i \in \mathbb{R}^p$
- (Sample/Empirical) Covariance matrix: $\Sigma = \frac{1}{n-1} \sum_{i=1}^n (x_i - \mu)(x_i - \mu)^T \in \mathbb{R}^{p \times p}$ (Covariance matrices are symmetric and positive semidefinite)
- Standard deviation (for $p = 1$): $\sigma = \sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2}$

PCA

- PCA is often used to **reduce the dimensionality** of large data sets while preserving as much information as possible.
- PCA allows us to identify the **principal directions in which the data varies**.

Let $x_1, \dots, x_n \in \mathbb{R}^p$ be n observations of a random variable x and

$$X = \begin{bmatrix} x_1^T \\ \vdots \\ x_n^T \end{bmatrix}.$$

The mean vectors of x_i and $Q^T x_i$ (for $i = 1, \dots, n$) are, respectively,

$$\mu = \frac{1}{n} \sum_{i=1}^n x_i, \quad \hat{\mu} = \frac{1}{n} \sum_{i=1}^n Q^T x_i = Q^T \mu.$$

Consequently, the associated covariance matrices are, respectively,

$$\Sigma = \frac{1}{n-1} \sum_{i=1}^n (x_i - \mu)(x_i - \mu)^T,$$

$$\hat{\Sigma} = \frac{1}{n-1} \sum_{i=1}^n (Q^T x_i - Q^T \mu)(Q^T x_i - Q^T \mu)^T = Q^T \Sigma Q.$$

Optimization problem of PCA

$$\max_{Q \in \mathbb{R}^{p \times k}, Q^T Q = I} \text{trace}(Q^T \Sigma Q).$$

Let the eigenvalue decomposition of Σ be

$$\Sigma = \begin{bmatrix} q_1 & \cdots & q_p \end{bmatrix} \begin{bmatrix} \lambda_1 & & \\ & \ddots & \\ & & \lambda_p \end{bmatrix} \begin{bmatrix} q_1 & \cdots & q_p \end{bmatrix}^T,$$

where

$$\lambda_1 \geq \cdots \geq \lambda_p \geq 0.$$

Then

$$Q = \begin{bmatrix} q_1 & \cdots & q_k \end{bmatrix}.$$

Standard PCA workflow

1. Make sure the data X are rows = observations and columns = variables.
2. Standardize the columns of X .
3. Run `[Q, X_new, d, tsquared, explained] = pca(X)`.
4. Using the variance% in "explained", choose k (usually 1, 2, or 3) components for visual analysis.
 - For example, if $d = (1.9087, 0.0913)$, explained = (95.4, 4.6), one may choose $k = 1$ as the first principal component carries 95.4% of the information.
 - For example, if $d = (2.9108, 0.9212, 0.1474, 0.0206)$, explained = (72.8, 23.0, 3.7, 0.5), one may choose $k = 2$ as the first two principal components carry 95.8% of the information.
5. Plot $X_{\text{new}}(:, 1), \dots, X_{\text{new}}(:, k)$ on a k -dimensional plot.

Lecture 2

Gradient Descent Method Given $x_0 \in \mathbb{R}^n$, for $k = 0, 1, 2, \dots$ do:

$$\begin{aligned} r_k &= Ax_k - b, \\ \alpha_k &= \frac{(r_k, r_k)}{(Ar_k, r_k)}, \\ x_{k+1} &= x_k - \alpha_k r_k. \end{aligned}$$

Gradient Descent Method Example: $Ax = b$ where A is Symmetric Positive Definite
Let

$$f(x) = \|x - x_\star\|_A^2 = (A(x - x_\star), (x - x_\star)) = (x - x_\star)^T A (x - x_\star),$$

where x_\star is the solution of

$$Ax = b.$$

It is obvious that

$$f(x) = 0 \quad \text{if and only if} \quad x = x_\star.$$

Denote

$$x = x_0 + \delta_0.$$

Then,

$$\begin{aligned} f(x) &= f(x_0) + (A\delta_0, \delta_0) + 2\delta_0^T (Ax_0 - b) \\ &= f(x_0) + \delta_0^T A\delta_0 + 2\delta_0^T r_0, \end{aligned}$$

where

$$r_0 = Ax_0 - b.$$

It is clear that

$$f(x) \leq f(x_0)$$

only if

$$\delta_0^T r_0 \leq 0,$$

in particular,

$$-r_0 = b - Ax_0$$

is the negative of the gradient direction $-\nabla f$ at the point x_0 .

The negative of the gradient direction is locally the direction that yields the fastest rate of decrease for f . Hence, we can choose

$$\delta_0 = -\alpha_0 r_0,$$

so that

$$\begin{aligned} f(x) &= f(x_0) + \alpha_0^2 (Ar_0, r_0) - 2\alpha_0 r_0^T r_0 \\ &= f(x_0) + \alpha_0^2 r_0^T Ar_0 - 2\alpha_0 r_0^T r_0 \leq f(x_0), \end{aligned}$$

provided

$$\alpha_0 \geq 0.$$

It is obvious, we have

$$f(x) \leq f(x_0), \quad \forall 0 \leq \alpha \leq \frac{2(r_0, r_0)}{(Ar_0, r_0)}.$$

The optimal α shall satisfy

$$f(x) = \min_{\alpha_0 \in \mathbb{R}} f(x_0) + \alpha_0^2 (Ar_0, r_0) - 2\alpha_0 r_0^T r_0,$$

i.e.,

$$\alpha_0 = \frac{(r_0, r_0)}{(Ar_0, r_0)} \geq 0.$$

Therefore, we conclude

$$\text{If } x = x_0 - \alpha_0 r_0, \quad \text{then } f(x) \leq f(x_0).$$

Kantorovich Inequality Let B be any Symmetric Positive Definite real matrix and λ_{\max} and λ_{\min} its largest and smallest eigenvalues. Then,

$$\frac{(Bx, x)(B^{-1}x, x)}{(x, x)^2} \leq \frac{(\lambda_{\max} + \lambda_{\min})^2}{4\lambda_{\max}\lambda_{\min}}, \quad \forall x \neq 0.$$

Kantorovich Inequality Proof

Clearly, it is equivalent to show that the result is true for any unit vector x . Since B is symmetric, we have

$$B = Q^T D Q,$$

where Q is orthogonal and

$$D = \begin{bmatrix} \lambda_1 & & \\ & \ddots & \\ & & \lambda_n \end{bmatrix},$$

and

$$\lambda_{\max} = \lambda_1 \geq \cdots \geq \lambda_n = \lambda_{\min} > 0.$$

We have

$$(Bx, x)(B^{-1}x, x) = (DQx, Qx)(D^{-1}Qx, Qx).$$

Setting

$$y = Qx = [y_1 \quad \cdots \quad y_n]^T, \quad \beta_i = y_i^2.$$

Note that $\sum_{i=1}^n \beta_i = 1$, and

$$\lambda = (Dy, y) = \sum_{i=1}^n \beta_i \lambda_i$$

is a convex combination of the eigenvalues λ_i , $i = 1, \dots, n$, and furthermore, the following relation holds,

$$(Bx, x)(B^{-1}x, x) = \lambda\psi(y),$$

with

$$\psi(y) = (D^{-1}y, y) = \sum_{i=1}^n \beta_i \frac{1}{\lambda_i}.$$

Noting that

$$\psi(y) \leq \frac{1}{\lambda_1} + \frac{1}{\lambda_n} - \frac{\lambda}{\lambda_1 \lambda_n}, \quad (\text{since } \sum_{i=1}^n \beta_i = 1, \text{ proved later})$$

therefore,

$$(Bx, x)(B^{-1}x, x) = \lambda\psi(y) \leq \lambda \left(\frac{1}{\lambda_1} + \frac{1}{\lambda_n} - \frac{\lambda}{\lambda_1 \lambda_n} \right).$$

The maximum of the right-hand side is reached for

$$\lambda = \frac{\lambda_1 + \lambda_n}{2}$$

yielding

$$\begin{aligned} (Bx, x)(B^{-1}x, x) = \lambda\psi(y) &\leq \lambda \left(\frac{1}{\lambda_1} + \frac{1}{\lambda_n} - \frac{\lambda}{\lambda_1 \lambda_n} \right) \\ &\leq \frac{\lambda_1 + \lambda_n}{4} \left(\frac{1}{\lambda_1} + \frac{1}{\lambda_n} \right) \end{aligned}$$

Proof for $\psi(y) \leq \frac{1}{\lambda_1} + \frac{1}{\lambda_n} - \frac{\lambda}{\lambda_1 \lambda_n}$

Since

$$0 < \lambda_n \leq \cdots \leq \lambda_i \leq \cdots \leq \lambda_1, \quad i = 1, \dots, n,$$

we have for any $i = 1, \dots, n$ that

$$\lambda_1 \geq \lambda_i > 0, \quad \lambda_i - \lambda_n \geq 0, \quad i = 1, \dots, n,$$

which gives

$$\lambda_1(\lambda_i - \lambda_n) \geq \lambda_i(\lambda_1 - \lambda_n),$$

i.e.,

$$\lambda_1 \lambda_n \leq \lambda_i(\lambda_1 + \lambda_n - \lambda_i),$$

and

$$\frac{1}{\lambda_i} \leq \frac{\lambda_1 + \lambda_n - \lambda_i}{\lambda_1 \lambda_n}.$$

Note that

$$\beta_i \geq 0, \quad \sum_{i=1}^n \beta_i = 1,$$

we get

$$\beta_i \frac{1}{\lambda_i} \leq \beta_i \frac{\lambda_1 + \lambda_n - \lambda_i}{\lambda_1 \lambda_n},$$

and so,

$$\begin{aligned} \sum_{i=1}^n \beta_i \frac{1}{\lambda_i} &\leq \sum_{i=1}^n \beta_i \frac{\lambda_1 + \lambda_n - \lambda_i}{\lambda_1 \lambda_n} \\ &= \frac{1}{\lambda_1} + \frac{1}{\lambda_n} - \frac{\sum_{i=1}^n \beta_i \lambda_i}{\lambda_1 \lambda_n}. \end{aligned}$$

This lemma helps to establish the following result regarding the convergence rate of the method.

Theorem Let A be a Symmetric Positive Definite matrix. Then, the A -norms of the error vectors

$$d_k = x_* - x_k = -A^{-1}r_k$$

generated by the Gradient Descent Algorithm satisfy the relation

$$\|d_{k+1}\|_A \leq \frac{\lambda_{\max} - \lambda_{\min}}{\lambda_{\max} + \lambda_{\min}} \|d_k\|_A,$$

and so,

$$\lim_{k \rightarrow \infty} \|d_k\|_A = 0,$$

which gives

$$\lim_{k \rightarrow \infty} d_k = 0,$$

i.e., the algorithm converges for any initial guess x_0 .

Proof First, we have

$$\|d_k\|_A^2 = (Ad_k, d_k) = (-r_k, d_k) = (r_k, A^{-1}r_k).$$

Then we have

$$\|d_{k+1}\|_A^2 = (Ad_{k+1}, d_{k+1}) = (-r_{k+1}, d_{k+1})$$

and by simple substitution,

$$\begin{aligned} d_{k+1} &= d_k + \alpha_k r_k, \\ \|d_{k+1}\|_A^2 &= (-r_{k+1}, d_k + \alpha_k r_k), \\ &= (-r_{k+1}, d_k) - \alpha(r_{k+1}, r_k), \\ &= (-r_{k+1}, d_k), \end{aligned}$$

since

$$(r_{k+1}, r_k) = 0.$$

Thus,

$$\begin{aligned} \|d_{k+1}\|_A^2 &= (-r_{k+1}, d_k), \\ &= (-r_k + \alpha_k A r_k, d_k), \\ &= (-r_k, d_k) + \alpha_k (A r_k, d_k), \\ &= (r_k, A^{-1}r_k) - \alpha_k (A r_k, A^{-1}r_k), \\ &= (r_k, A^{-1}r_k) - \frac{(r_k, r_k)^2}{(A r_k, r_k)}, \\ &= \|d_k\|_A^2 \left(1 - \frac{(r_k, r_k)}{(A r_k, r_k)} \times \frac{(r_k, r_k)}{(r_k, A^{-1}r_k)} \right). \end{aligned}$$

The result follows by applying the Kantorovich inequality.

Unconstrained problem

To minimize a **differentiable** function f

$$\min_{x \in \mathbb{R}^n} f(x)$$

Recall that a global minimizer is a local minimizer, and a local minimizer is a stationary point.

- We may try to find stationary points x , i.e., $\nabla f(x) = 0$ for solving an unconstrained problem.
- When it is difficult to solve $\nabla f(x) = 0$, we look for an approximate solution via iterative methods.

A general algorithmic framework

Choose $x^{(0)}$ and repeat

$$x^{(k+1)} = x^{(k)} + \alpha_k p^{(k)}, \quad k = 0, 1, 2, \dots$$

until some stopping criteria is satisfied.

- $x^{(0)}$ initial guess of the solution.
- $\alpha_k > 0$ is called the step length/step size/learning rate.
- $p^{(k)}$ is a search direction.

Descent Direction

The search direction $p^{(k)}$ should be a descent direction at $x^{(k)}$.

- We say $p^{(k)}$ is a descent direction at $x^{(k)}$ if

$$\nabla f(x^{(k)})^T p^{(k)} < 0$$

- The function value f can be reduced along this descent direction with “appropriate” step length

$$\exists \delta > 0 \quad \text{such that} \quad f(x^{(k)} + \alpha_k p^{(k)}) < f(x^{(k)}) \quad \forall \alpha_k \in (0, \delta)$$

Algorithm 1 Steepest Descent Method

1: **Initialization:** Choose initial point $x^{(0)}$, tolerance $\epsilon > 0$, set $k \leftarrow 0$.

2: **while** $\|\nabla f(x^{(k)})\| > \epsilon$ **do**

3: Find the step length α_k (e.g., by a certain line search rule).

4: Update the solution:

$$x^{(k+1)} = x^{(k)} - \alpha_k \nabla f(x^{(k)})$$

5: Increment $k \leftarrow k + 1$.

6: **end while**

7: **Output:** $x^{(k)}$ (approximate solution)

One may choose to use a constant step length (say $\alpha_k = 0.1$), or find it via line search rules:

- Exact line search
- Backtracking line search

Exact line search

Exact line search tries to find α_k by solving the one-dimensional problem:

$$\min_{\alpha > 0} \varphi(\alpha) := f(x^{(k)} + \alpha p^{(k)})$$

- In general, exact line search is the most difficult part of the steepest descent method.
- If f is a simple function, it may be possible to obtain an analytical solution for α_k by solving $\varphi'(\alpha) = 0$.

Contour plot A contour is a fixed height $f(x_1, x_2) = c$.

Algorithm 2 Steepest Descent Method with Exact Line Search

- 1: **Initialization:** Choose initial point $x^{(0)}$, tolerance $\epsilon > 0$, set $k \leftarrow 0$.
2: **while** $\|\nabla f(x^{(k)})\| > \epsilon$ **do**
3: Compute search direction: $p^{(k)} = -\nabla f(x^{(k)})$.
4: Find optimal step length:

$$\alpha_k = \arg \min_{\alpha > 0} f(x^{(k)} + \alpha p^{(k)})$$

- 5: Update the solution:

$$x^{(k+1)} = x^{(k)} + \alpha_k p^{(k)}$$

- 6: Increment $k \leftarrow k + 1$.

7: **end while**

8: **Output:** $x^{(k)}$ (approximate solution)

Properties of steepest descent method with exact line search

Let $\{x^{(k)}\}$ be the sequence generated by the steepest descent method with exact line search.

- **Monotonic decreasing property:**

$$f(x^{(k+1)}) < f(x^{(k)}) \quad \text{if } \nabla f(x^{(k)}) \neq 0.$$

- Suppose f is a *coercive* function with continuous first-order derivatives on \mathbb{R}^n . Then some subsequence of $\{x^{(k)}\}$ converges.

The limit of any convergent subsequence of $\{x^{(k)}\}$ is a stationary point of f .

Backtracking Line Search

Backtracking line search starts with a relatively large step length and iteratively shrinks it (i.e., "backtracking") until the Armijo condition holds.

Algorithm 3 Backtracking Line Search

- 1: Choose $\bar{\alpha} > 0$, $\rho \in (0, 1)$, $c_1 \in (0, 1)$; Set $\alpha \leftarrow \bar{\alpha}$.
2: **repeat**
3: Until

$$f(x^{(k)} + \alpha p^{(k)}) \leq f(x^{(k)}) + c_1 \alpha \nabla f(x^{(k)})^T p^{(k)}$$

▷ Armijo Condition

- 4: $\alpha \leftarrow \rho \alpha$

5: **until** Armijo condition holds

6: **return** $\alpha_k = \alpha$

Notes:

- $p^{(k)}$ is a descent direction:

$$\nabla f(x^{(k)})^T p^{(k)} < 0$$

- The Armijo condition:

$$f(x^{(k)} + \alpha p^{(k)}) \leq f(x^{(k)}) + c_1 \alpha \nabla f(x^{(k)})^T p^{(k)}$$

ensures a reasonable amount of decrease in the objective function.

- Example parameter choices:

$$\bar{\alpha} = 1, \quad \rho = 0.9, \quad c_1 = 10^{-4}$$

Algorithm 4 Steepest Descent Method with Backtracking Line Search

- 1: Choose $x^{(0)}$, $\epsilon > 0$, $\bar{\alpha} > 0$, $\rho \in (0, 1)$, $c_1 \in (0, 1)$; Set $k \leftarrow 0$.
2: **while** $\|\nabla f(x^{(k)})\| > \epsilon$ **do**
3: Compute search direction: $p^{(k)} = -\nabla f(x^{(k)})$.
4: Set $\alpha \leftarrow \bar{\alpha}$.
5: **repeat**
6: Until Armijo condition holds:

$$f(x^{(k)} + \alpha p^{(k)}) \leq f(x^{(k)}) + c_1 \alpha \nabla f(x^{(k)})^T p^{(k)}$$

- 7: $\alpha \leftarrow \rho \alpha$.

8: **until** Armijo condition holds

9: $\alpha_k \leftarrow \alpha$.

10: Update the solution:

$$x^{(k+1)} = x^{(k)} + \alpha_k p^{(k)}$$

11: Increment $k \leftarrow k + 1$.

12: **end while**

13: **return** $x^{(k)}$.

Steepest Descent Method for Multivariate Linear Regression

Algorithm 5 Steepest Descent for Multivariate Linear Regression

- 1: Choose $\beta_0^{(0)}, \beta^{(0)} = (\beta_1^{(0)}, \dots, \beta_p^{(0)})^T$ and $\epsilon > 0$; Set $k \leftarrow 0$.
2: **while** $\|\nabla L(\beta_0^{(k)}, \beta^{(k)})\| > \epsilon$ **do**
3: Determine step length α_k .
4: Update parameters:

$$\beta_0^{(k+1)} = \beta_0^{(k)} - \alpha_k \sum_{i=1}^n ((\beta^{(k)})^T x_i + \beta_0^{(k)} - y_i)$$

- 5: **for** $j = 1, 2, \dots, p$ **do**

$$\beta_j^{(k+1)} = \beta_j^{(k)} - \alpha_k \sum_{i=1}^n ((\beta^{(k)})^T x_i + \beta_0^{(k)} - y_i) x_{ij}$$

6: **end for**

7: Increment $k \leftarrow k + 1$.

8: **end while**

9: **return** $\beta_0^{(k)}, \beta^{(k)} = (\beta_1^{(k)}, \dots, \beta_p^{(k)})^T$.

Normal Equation

$$\min_{\beta_0, \beta_1, \dots, \beta_p} L(\beta_0, \beta_1, \dots, \beta_p) = \frac{1}{2} \sum_{i=1}^n (\beta^T x_i + \beta_0 - y_i)^2$$

$$\hat{X}^T \hat{X} \hat{\beta} = \hat{X}^T Y$$

How to solve

$$\hat{X}^T \hat{X} \hat{\beta} = \hat{X}^T Y$$

Case 1. When $\hat{X}^T \hat{X}$ is invertible, the normal equation implies that

$$\hat{\beta} = (\hat{X}^T \hat{X})^{-1} \hat{X}^T Y$$

is the **unique** solution of linear regression.

This often happens when we face an **over-determined system** — number of training examples n is much larger than number of features p .

We have many training samples to fit but do not have enough degree of freedom.

Case 2. When $\hat{X}^T \hat{X}$ is not invertible, the normal equation will have infinite number of solutions.

$\hat{X}^T \hat{X}$ is not invertible when we face an **under-determined problem** — $n < p$.

We have too many degrees of freedom and do not have enough training samples.

We can apply any method for solving a linear system (e.g., Gaussian elimination) to obtain a solution.

Lecture 3

Classification Binary classification:

- Email: spam/not spam
- Student: fail/pass

We usually assign:

$$\text{label} \begin{cases} 0, & \text{normal state/negative class, e.g., not spam} \\ 1, & \text{abnormal state/positive class, e.g., spam} \end{cases}$$

However, the label assignment can be arbitrary:

$$0 = \text{not spam}, 1 = \text{spam} \quad \text{or} \quad 0 = \text{spam}, 1 = \text{not spam}$$

Data: $x_i \in \mathbb{R}^p$, $y_i \in \{0, 1\}$, $i = 1, 2, \dots, n$.

Multi-class classification:

- Iris flower (3 species: Setosa, Versicolor, Virginica)
- Optical character recognition

Data: $x_i \in \mathbb{R}^p$, $y_i \in \{1, \dots, K\}$, $i = 1, 2, \dots, n$.

Linear Regression vs. Logistic Regression

Linear Regression

- Data $x_i, y_i \in \mathbb{R}$
- Fit: $f(x) = \beta^T x + \beta_0 = \hat{\beta}^T \hat{x}$, where $\hat{\beta} = [\beta_0; \beta]$, $\hat{x} = [1; x]$

Logistic Regression

- Data $x_i, y_i \in \{0, 1\}$
- Fit:

$$f(x) = g(\hat{\beta}^T \hat{x})$$

where

$$g(z) = \frac{1}{1 + e^{-z}} \quad (\text{logistic function})$$

so,

$$f(x) = g(\hat{\beta}^T \hat{x}) = \frac{1}{1 + e^{-(\beta^T x + \beta_0)}}$$

Logistic Regression Decision Rule

$$f(x) = g(\hat{\beta}^T \hat{x}), \quad g(z) = \frac{1}{1 + e^{-z}}$$

$$f(x) = p(y = 1|x; \hat{\beta})$$

Predict $y = 1$ (class 1) if:

$$f(x) \geq 0.5 \quad \text{i.e.,} \quad \hat{\beta}^T \hat{x} \geq 0$$

Predict $y = 0$ (class 0) if:

$$f(x) < 0.5 \quad \text{i.e.,} \quad \hat{\beta}^T \hat{x} < 0$$

Decision Boundary

The set of all $x \in \mathbb{R}^p$ such that:

$$\beta_0 + \beta^T x = 0$$

is called the decision boundary between classes 0 and 1.

The logistic regression has a linear decision boundary; it is:

- a point when $p = 1$
- a line when $p = 2$
- a plane when $p = 3$
- in general a $(p - 1)$ -dimensional subspace

Maximum Likelihood Estimation

- Data (x_i, y_i) , $i = 1, 2, \dots, n$, $x_i \in \mathbb{R}^p$, $y_i \in \{0, 1\}$.
- The likelihood of a single training example (x_i, y_i) is:

$$\begin{aligned} \text{probability}(x_i \in \text{class } y_i) &= \begin{cases} p(y_i = 1|x_i; \hat{\beta}) = f(x_i), & \text{if } y_i = 1 \\ p(y_i = 0|x_i; \hat{\beta}) = 1 - f(x_i), & \text{if } y_i = 0 \end{cases} \\ &= f(x_i)^{y_i} [1 - f(x_i)]^{1-y_i} \end{aligned}$$

- Assuming independence of training samples, the likelihood is:

$$\prod_{i=1}^n f(x_i)^{y_i} [1 - f(x_i)]^{1-y_i}$$

- Want to find $\hat{\beta}$ to maximize the log-likelihood:

$$\begin{aligned} L(\hat{\beta}) &= -\log \left(\prod_{i=1}^n f(x_i)^{y_i} [1 - f(x_i)]^{1-y_i} \right) \\ &= -\sum_{i=1}^n (y_i \log f(x_i) + (1 - y_i) \log(1 - f(x_i))) \end{aligned}$$

For a single training example (x_i, y_i) , the cost is:

$$\begin{aligned} &-y_i \log f(x_i) - (1 - y_i) \log(1 - f(x_i)) \\ &= \begin{cases} -\log f(x_i), & \text{if } y_i = 1 \\ -\log(1 - f(x_i)), & \text{if } y_i = 0 \end{cases} \end{aligned}$$

Simplifying the Cost Function

$$\begin{aligned} \log \left(\frac{f(x_i)}{1 - f(x_i)} \right) &= \log \left(\frac{\frac{1}{1 + e^{-\hat{\beta}^T \hat{x}_i}}}{1 - \frac{1}{1 + e^{-\hat{\beta}^T \hat{x}_i}}} \right) \\ &= \log(e^{\hat{\beta}^T \hat{x}_i}) = \hat{\beta}^T \hat{x}_i \\ \log(1 - f(x_i)) &= \log \left(1 - \frac{1}{1 + e^{-\hat{\beta}^T \hat{x}_i}} \right) \\ &= \log \left(\frac{1 + e^{\hat{\beta}^T \hat{x}_i} - 1}{1 + e^{\hat{\beta}^T \hat{x}_i}} \right) \\ &= -\log \left(1 + e^{\hat{\beta}^T \hat{x}_i} \right) \end{aligned}$$

Gradient of the cost function

- Cost function

$$L(\beta_0, \beta_1, \dots, \beta_p) = \sum_{i=1}^n \log(1 + e^{\beta_0 + \beta^T x_i}) - y_i(\beta_0 + \beta^T x_i)$$

$$\beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}$$

- Calculate

$$\begin{aligned} \frac{\partial}{\partial \beta_0} L &= \sum_{i=1}^n \left(\frac{1}{1 + e^{-(\beta_0 + \beta^T x_i)}} - y_i \right) = \sum_{i=1}^n (f(x_i) - y_i) \\ \frac{\partial}{\partial \beta_1} L &= \sum_{i=1}^n \left(\frac{1}{1 + e^{-(\beta_0 + \beta^T x_i)}} - y_i \right) x_{i1} = \sum_{i=1}^n (f(x_i) - y_i) x_{i1} \\ \frac{\partial}{\partial \beta_2} L &= \sum_{i=1}^n \left(\frac{1}{1 + e^{-(\beta_0 + \beta^T x_i)}} - y_i \right) x_{i2} = \sum_{i=1}^n (f(x_i) - y_i) x_{i2} \\ &\vdots \\ &\vdots \end{aligned}$$

$$\frac{\partial}{\partial \beta_p} L = \sum_{i=1}^n \left(\frac{1}{1 + e^{-(\beta_0 + \beta^T x_i)}} - y_i \right) x_{ip} = \sum_{i=1}^n (f(x_i) - y_i) x_{ip}$$

Solution may not exist

The solution (global minimizer) of the minimization problem

$$\min_{\beta_0, \beta_1, \dots, \beta_p} \sum_{i=1}^n \log(1 + e^{\beta_0 + \beta^T x_i}) - y_i(\beta_0 + \beta^T x_i)$$

may not exist. (Regularization will help solve this issue)

Example. $n = 1$, $x_1 = -1$, $y_1 = 0$. Then the cost function

$$L(\beta_0, \beta_1) = \log(1 + e^{\beta_0 - \beta_1})$$

We can see that $\min L = 0$. However, this value cannot be attained.

Multi-class classification: one-vs-rest

Idea: transfer multi-class classification to multiple binary classification problems

Data: $x_i \in \mathbb{R}^p$, $y_i \in \{1, \dots, K\}$, $i = 1, 2, \dots, n$.

For each $k \in \{1, 2, \dots, K\}$

- Construct a new label $\tilde{y}_i = 1$ if $y_i = k$ and $\tilde{y}_i = 0$ otherwise
- Learn a binary classifier f_k with data x_i, \tilde{y}_i

Multi-class classifier predicts class k where k achieves the maximal value

$$\max_{k \in \{1, 2, \dots, K\}} f_k(x)$$

Overfitting

- Underfitting:** a model is too simple and does not adequately capture the underlying structure of the data
- Overfitting:** a model is too complicated and contains more parameters than can be justified by the data; it does not generalize well from training data to test data
- Good fit:** a model adequately learns the training data and generalizes well to test data

Ridge regularization

In linear/logistic regression, over-fitting occurs frequently. Regularization will make the model simpler and works well for most of the regression/classification problems.

- Ridge regularization:

$$\lambda \|\beta\|^2 = \lambda \sum_{j=1}^p \beta_j^2$$

λ : regularization parameter, $\|\beta\|^2$: regularizer

- It is differentiable. It forces β_j 's to be small
- Extreme case: suppose λ is a huge number, it will push all β_j 's to be zero and the model will be naive

Ridge regularized problems

- Logistic regression + ridge regularization (Gradient methods can be used, a solution exists)

$$\min_{\beta_0, \beta_1, \dots, \beta_p} \sum_{i=1}^n \log(1 + e^{\beta_0 + \beta^T x_i}) - y_i(\beta_0 + \beta^T x_i) + \lambda \sum_{j=1}^p \beta_j^2$$

- Linear regression + ridge regularization (Apply either normal equation or gradient methods)

$$\min_{\beta_0, \beta_1, \dots, \beta_p} \frac{1}{2} \sum_{i=1}^n (\beta^T x_i + \beta_0 - y_i)^2 + \lambda \sum_{j=1}^p \beta_j^2$$

Lasso regularization

- Lasso (Least Absolute Shrinkage and Selection Operator) regularization:

$$\lambda \|\beta\|_1 = \lambda \sum_{j=1}^p |\beta_j|$$

- It is non-differentiable. It forces some β_j 's to be exactly zero
- It can be used for feature selection (model selection). It selects important features (removing non-informative or redundant features)
- When λ is larger, fewer features will be selected

Lasso regularized problems

- Logistic regression + lasso regularization (Gradient methods are no longer applicable)

$$\min_{\beta_0, \beta_1, \dots, \beta_p} \sum_{i=1}^n \log(1 + e^{\beta_0 + \beta^T x_i}) - y_i(\beta_0 + \beta^T x_i) + \lambda \sum_{j=1}^p |\beta_j|$$

- Linear regression + lasso regularization (Gradient methods are no longer applicable)

$$\min_{\beta_0, \beta_1, \dots, \beta_p} \frac{1}{2} \sum_{i=1}^n (\beta^T x_i + \beta_0 - y_i)^2 + \lambda \sum_{j=1}^p |\beta_j|$$

- In the following, we always assume $\beta_0 = 0$. Note that the intercept should be zero $\beta_0 = 0$ if the data is standardized. Given feature matrix $X \in \mathbb{R}^{n \times p}$ and response vector $Y \in \mathbb{R}^p$, the famous lasso problem:

$$\min_{\beta \in \mathbb{R}^p} \frac{1}{2} \|X\beta - Y\|^2 + \lambda \|\beta\|_1$$

Lecture 4

Norms

Definition: A vector norm on \mathbb{R}^n is a function $\|\cdot\| : \mathbb{R}^n \rightarrow \mathbb{R}$ satisfying:

1. $\|x\| \geq 0 \quad \forall x \in \mathbb{R}^n$, and $\|x\| = 0 \iff x = 0$
2. $\|\alpha x\| = |\alpha| \|x\| \quad \forall \alpha \in \mathbb{R}, x \in \mathbb{R}^n$
3. $\|x + y\| \leq \|x\| + \|y\| \quad \forall x, y \in \mathbb{R}^n$

Examples:

- $\|x\|_1 = \sum_{i=1}^n |x_i|$ (ℓ_1 norm)
- $\|x\|_2 = \left(\sum_{i=1}^n |x_i|^2\right)^{1/2}$ (ℓ_2 norm)
- $\|x\|_p = \left(\sum_{i=1}^n |x_i|^p\right)^{1/p}$, $1 \leq p < \infty$ (ℓ_p norm)
- $\|x\|_\infty = \max_{1 \leq i \leq n} |x_i|$ (ℓ_∞ norm)
- $\|x\|_{W,p} = \|Wx\|_p$, where W is a fixed nonsingular matrix, $1 \leq p \leq \infty$

Inner product

Definition:

The trace of a square matrix $C \in \mathbb{R}^{n \times n}$ is

$$\text{Tr}(C) = \sum_{i=1}^n C_{ii}.$$

For matrices $A, B \in \mathbb{R}^{m \times n}$, define the standard inner product:

$$\langle A, B \rangle = \text{Tr}(A^T B) = \sum_{i=1}^m \sum_{j=1}^n A_{ij} B_{ij}.$$

Properties:

- $\|A\|_F^2 := \langle A, A \rangle \geq 0$
- $\langle A, A \rangle = 0$ if and only if $A = 0$
- $\langle C, aA + bB \rangle = a\langle C, A \rangle + b\langle C, B \rangle$
- $\langle A + B, A + B \rangle = \langle A, A \rangle + 2\langle A, B \rangle + \langle B, B \rangle$

(i.e., $\|A + B\|_F^2 = \|A\|_F^2 + 2\langle A, B \rangle + \|B\|_F^2$)

Projection onto a closed convex set

Theorem (Projection theorem): Let C be a closed convex set of \mathbb{R}^n .

- (1) For every z , there exists a unique minimizer of

$$\min_{x \in C} \frac{1}{2} \|x - z\|^2,$$

denoted as $\Pi_C(z)$, the projection of z onto C .

- (2) $x^* := \Pi_C(z)$ is the projection of z onto C if and only if

$$\langle z - x^*, x - x^* \rangle \leq 0 \quad \forall x \in C.$$

Proof:

- (1) By definition, there exists $x_k \in C$ such that

$$\min_{x \in C} \|z - x\| \leq \|z - x_k\| < \min_{x \in C} \|z - x\| + \frac{1}{k}.$$

It follows that $\{x_k\}$ is bounded. Since C is closed, there exists a convergent subsequence $\{x_{k_l}\}$ such that $x_{k_l} \rightarrow x^* \in C$. Therefore,

$$\|z - x^*\| = \min_{x \in C} \|z - x\|.$$

For uniqueness, suppose $x^* \neq \tilde{x}$ both satisfy $x^*, \tilde{x} \in C$ and

$$\|z - x^*\| = \|z - \tilde{x}\| = \min_{x \in C} \|z - x\|.$$

Then,

$$2\|z - x^*\|^2 = \|z - x^*\|^2 + \|z - \tilde{x}\|^2 = 2\left\|z - \frac{x^* + \tilde{x}}{2}\right\|^2 + \frac{1}{2}\|x^* - \tilde{x}\|^2.$$

Since C is convex, $\frac{x^* + \tilde{x}}{2} \in C$. Thus,

$$\left\|z - \frac{x^* + \tilde{x}}{2}\right\|^2 \geq \min_{x \in C} \|z - x\|^2 = \|z - x^*\|^2,$$

which implies

$$2\|z - x^*\|^2 \geq 2\|z - x^*\|^2 + \frac{1}{2}\|x^* - \tilde{x}\|^2.$$

Thus, $\|x^* - \tilde{x}\|^2 \leq 0$ and hence $x^* = \tilde{x}$.

- (2) Now, let $x^* = \Pi_C(z)$, and for any $x \in C$, since C is convex,

$$\lambda x + (1 - \lambda)x^* \in C, \quad \forall \lambda \in (0, 1).$$

By minimality,

$$\|z - x^*\|^2 \leq \|z - (\lambda x + (1 - \lambda)x^*)\|^2.$$

Expanding the right-hand side,

$$\|z - (\lambda x + (1 - \lambda)x^*)\|^2 = \|z - x^*\|^2 - 2\lambda \langle z - x^*, x - x^* \rangle + \lambda^2 \|x - x^*\|^2.$$

Thus,

$$0 \leq -2\lambda \langle z - x^*, x - x^* \rangle + \lambda^2 \|x - x^*\|^2.$$

Dividing by $\lambda > 0$ and taking $\lambda \rightarrow 0^+$, we get

$$\langle z - x^*, x - x^* \rangle \leq 0.$$

Conversely, if $\langle z - x^*, x - x^* \rangle \leq 0 \quad \forall x \in C$, then

$$\|z - x\|^2 = \|z - x^*\|^2 + 2\langle z - x^*, x^* - x \rangle + \|x - x^*\|^2 \geq \|z - x^*\|^2,$$

thus x^* minimizes $\|z - x\|$ over C .

arg min

The notation

$$\arg \min_x f(x)$$

denotes the solution set of x for which $f(x)$ attains its minimum (argument of the minimum).

Extended real-valued function

Definition.

Let \mathcal{X} be a Euclidean space (e.g., $\mathcal{X} = \mathbb{R}^n$ or $\mathbb{R}^m \times \mathbb{R}^n$). Let $f : \mathcal{X} \rightarrow (-\infty, +\infty]$ be an extended real-valued function.

1. The **effective domain** of f is defined as

$$\text{dom}(f) := \{x \in \mathcal{X} \mid f(x) < +\infty\}.$$

2. f is said to be **proper** if $\text{dom}(f) \neq \emptyset$.
3. f is said to be **closed** if its epi-graph

$$\text{epi}(f) := \{(x, \alpha) \in \mathcal{X} \times \mathbb{R} \mid f(x) \leq \alpha\}$$

is closed.

4. f is said to be **convex** if its epi-graph is convex.

Extended real-valued function (continued)

- For a real-valued function $f : \mathcal{X} \rightarrow \mathbb{R}$, the convexity of $\text{epi}(f)$ coincides with the following condition:

$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y), \quad \forall x, y \in \text{dom}(f), \quad \lambda \in [0, 1].$$

(Exercise: (1) \iff (2))

- A convex function $f : D \subseteq \mathbb{R}^n \rightarrow \mathbb{R}$ can be extended to a convex function on all of \mathbb{R}^n by setting $f(x) = +\infty$ for $x \notin D$.

Proof of Exercise (1) \iff (2)

(1) \Rightarrow (2):

Assume $\text{epi}(f)$ is convex. Take any $(x, f(x)) \in \text{epi}(f)$ and $(y, f(y)) \in \text{epi}(f)$. For any $\lambda \in [0, 1]$, by convexity of $\text{epi}(f)$,

$$\lambda(x, f(x)) + (1 - \lambda)(y, f(y)) = (\lambda x + (1 - \lambda)y, \lambda f(x) + (1 - \lambda)f(y)) \in \text{epi}(f).$$

Hence,

$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y).$$

Thus, (2) holds.

(2) \Rightarrow (1):

Assume

$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y), \quad \forall x, y \in \text{dom}(f), \quad \lambda \in [0, 1].$$

Let $(x, \alpha), (y, \beta) \in \text{epi}(f)$. Then $f(x) \leq \alpha$ and $f(y) \leq \beta$.

For any $\lambda \in [0, 1]$,

$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y) \leq \lambda \alpha + (1 - \lambda)\beta.$$

Therefore,

$$(\lambda x + (1 - \lambda)y, \lambda \alpha + (1 - \lambda)\beta) \in \text{epi}(f),$$

which shows $\text{epi}(f)$ is convex. Thus, (1) holds.

Indicator function

Let C be a nonempty set in \mathcal{X} .

The **indicator function** of C is

$$\delta_C(x) = \begin{cases} 0, & \text{if } x \in C, \\ +\infty, & \text{if } x \notin C. \end{cases}$$

Let $f(x) = \delta_C(x)$.

Then:

- $\text{dom}(f) = C$, and f is proper.
- $\text{epi}(f) = C \times [0, +\infty)$ is closed if C is closed.
- $\text{epi}(f)$ is convex if C is convex.

Dual/polar cone

Definition (Cone):

A set $C \subseteq \mathcal{X}$ is called a **cone** if $\lambda x \in C$ when $x \in C$ and $\lambda \geq 0$.

Definition (Dual and polar cone):

The **dual cone** of a set $C \subseteq \mathcal{X}$ is

$$C^* := \{y \in \mathcal{X} \mid \langle x, y \rangle \geq 0 \quad \forall x \in C\}.$$

The **polar cone** of C is $C^\circ = -C^*$.

If $C^* = C$, then C is said to be **self-dual**.

- C^* is always a convex cone, even if C is neither convex nor a cone.

Normal cone

Definition (Normal cone):

Let C be a convex set in \mathcal{X} and $\bar{x} \in C$. The normal cone of C at $\bar{x} \in C$ is defined by

$$N_C(\bar{x}) := \{z \in \mathcal{X} \mid \langle z, x - \bar{x} \rangle \leq 0 \quad \forall x \in C\}.$$

By convention, we let $N_C(\bar{x}) = \emptyset$ if $\bar{x} \notin C$.

Proposition:

Let $C \subseteq \mathcal{X}$ be a nonempty convex set and $\bar{x} \in C$. Then:

1. $N_C(\bar{x})$ is a closed convex cone.
2. If $\bar{x} \in \text{int}(C)$ (interior point), then $N_C(\bar{x}) = \{0\}$.
3. If C is a cone, then $N_C(\bar{x}) \subseteq C^\circ$.

Proof:

(1) Proof of closeness:

Let $\{z_k\}$ be a sequence in $N_C(\bar{x})$ such that

$$z_k \rightarrow z \quad (\text{in norm}).$$

We want to show $z \in N_C(\bar{x})$. Since $z_k \in N_C(\bar{x})$, we have

$$\langle z_k, x - \bar{x} \rangle \leq 0 \quad \forall x \in C, \quad \forall k.$$

Fix $x \in C$. By continuity of the inner product and $z_k \rightarrow z$,

$$\langle z_k, x - \bar{x} \rangle \rightarrow \langle z, x - \bar{x} \rangle.$$

Since each $\langle z_k, x - \bar{x} \rangle \leq 0$, taking limits gives

$$\langle z, x - \bar{x} \rangle \leq 0.$$

Thus $z \in N_C(\bar{x})$. Hence, $N_C(\bar{x})$ is closed.

First, we prove that $N_C(\bar{x})$ is a cone. Let $z \in N_C(\bar{x})$ and $\lambda \geq 0$. By definition,

$$\langle z, x - \bar{x} \rangle \leq 0 \quad \forall x \in C$$

implies

$$\langle \lambda z, x - \bar{x} \rangle = \lambda \langle z, x - \bar{x} \rangle \leq 0 \quad \forall x \in C.$$

Thus $\lambda z \in N_C(\bar{x})$.

Next, if $z_1, z_2 \in N_C(\bar{x})$, then

$$\langle z_1, x - \bar{x} \rangle \leq 0 \quad \text{and} \quad \langle z_2, x - \bar{x} \rangle \leq 0 \quad \forall x \in C.$$

Adding, we get

$$\langle z_1 + z_2, x - \bar{x} \rangle \leq 0 \quad \forall x \in C,$$

so $z_1 + z_2 \in N_C(\bar{x})$. Therefore, $N_C(\bar{x})$ is convex.

(2) Let $z \in N_C(\bar{x})$. Since $\bar{x} \in \text{int}(C)$, there exists $\epsilon > 0$ such that $\bar{x} + tz \in C$ for all $|t| < \epsilon$. By definition of the normal cone,

$$0 \geq \langle z, (\bar{x} + tz) - \bar{x} \rangle = t\|z\|^2.$$

For small positive $t > 0$, $t\|z\|^2 \geq 0$, so $t\|z\|^2 \leq 0$. Hence, $\|z\|^2 = 0$, thus $z = 0$. Therefore, $N_C(\bar{x}) = \{0\}$.

(3) Suppose C is a cone. Then for any $x \in C$, $\bar{x} + x \in C$. For any $z \in N_C(\bar{x})$, we have

$$\langle z, x \rangle = \langle z, (\bar{x} + x) - \bar{x} \rangle \leq 0 \quad \forall x \in C.$$

Hence, $z \in C^\circ$. Thus, $N_C(\bar{x}) \subseteq C^\circ$.

Proposition:

Let $C \subseteq \mathcal{X}$ be a nonempty closed convex set. Then for any $y \in C$,

$$u \in N_C(y) \iff y = \Pi_C(y + u),$$

where $\Pi_C(\cdot)$ is the projection onto C .

Proof:

" \Rightarrow " Suppose $u \in N_C(y)$. Then

$$\langle u, x - y \rangle \leq 0 \quad \forall x \in C.$$

Thus,

$$\langle (y + u) - y, x - y \rangle \leq 0 \quad \forall x \in C,$$

which implies that $y = \Pi_C(y + u)$.

" \Leftarrow " Suppose $y = \Pi_C(y + u)$. Then for all $x \in C$,

$$\langle u, x - y \rangle = \langle (y + u) - y, x - y \rangle \leq 0,$$

which implies $u \in N_C(y)$.

Subdifferential

Definition:

Let $f : \mathcal{X} \rightarrow (-\infty, +\infty]$ be a convex function.

We call v a **subgradient** of f at $x \in \text{dom}(f)$ if

$$f(z) \geq f(x) + \langle v, z - x \rangle \quad \forall z \in \mathcal{X}.$$

The set of all subgradients at x is called the **subdifferential** of f at x , denoted by

$$\partial f(x) = \{v \mid f(z) \geq f(x) + \langle v, z - x \rangle \quad \forall z \in \mathcal{X}\}.$$

By convention, $\partial f(x) = \emptyset$ for any $x \notin \text{dom}(f)$.

Subdifferential and optimization

Subgradient is an extension of the gradient:

- If f is differentiable at x , then

$$\partial f(x) = \{\nabla f(x)\}.$$

Proof:

Suppose $v \in \partial f(x)$. Then

$$f(x+h) \geq f(x) + \langle v, h \rangle \quad \forall h.$$

Take $h = t(v - \nabla f(x))$ with $t > 0$, and use Taylor's expansion:

$$f(x+h) = f(x) + \langle \nabla f(x), h \rangle + o(\|h\|).$$

So

$$\langle \nabla f(x), h \rangle + o(\|h\|) \geq \langle v, h \rangle.$$

Dividing by t and letting $t \rightarrow 0^+$, we conclude $v = \nabla f(x)$.

Theorem:

Let $f : \mathcal{X} \rightarrow (-\infty, +\infty]$ be a proper convex function.

Then $\bar{x} \in \mathcal{X}$ is a global minimizer of $\min_{x \in \mathcal{X}} f(x)$ if and only if

$$0 \in \partial f(\bar{x}).$$

Proof:

By the subgradient inequality,

$$f(z) \geq f(\bar{x}) + \langle v, z - \bar{x} \rangle \quad \forall z \in \mathcal{X}.$$

Take $v = 0 \in \partial f(\bar{x})$, then

$$f(z) \geq f(\bar{x}),$$

so \bar{x} is a global minimizer.

Lipschitz continuous

Definition (Lipschitz continuous):

A function $F : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is said to be **locally Lipschitz continuous** if for any open set $\mathcal{O} \subseteq \mathbb{R}^n$, there exists a constant L such that

$$\|F(x) - F(y)\| \leq L\|x - y\| \quad \forall x, y \in \mathcal{O}.$$

If $\mathcal{O} = \mathbb{R}^n$, then F is called **globally Lipschitz continuous**.

Examples:

- $f(x) = |x|$, $x \in \mathbb{R}$ is globally Lipschitz continuous with Lipschitz constant $L = 1$.
- $f(x) = x^2$, $x \in \mathbb{R}$ is locally Lipschitz continuous but not globally Lipschitz continuous.

Fenchel conjugate

Definition:

Let $f : \mathcal{X} \rightarrow [-\infty, +\infty]$.

The (Fenchel) conjugate of f is defined as

$$f^*(y) = \sup\{\langle y, x \rangle - f(x) \mid x \in \mathcal{X}\}, \quad y \in \mathcal{X}^*.$$

Remark:

- f^* is always closed and convex, even if f is neither convex nor closed.
- If $f : \mathcal{X} \rightarrow (-\infty, +\infty]$ is closed and proper convex, then

$$(f^*)^* = f.$$

Proposition:

Let $f : \mathcal{X} \rightarrow (-\infty, +\infty]$ be a closed proper convex function. The following are equivalent:

1. $f(x) + f^*(y) = \langle x, y \rangle$
2. $y \in \partial f(x)$
3. $x \in \partial f^*(y)$

The equivalence $y \in \partial f(x) \iff x \in \partial f^*(y)$ means that ∂f^* is the **inverse** of ∂f (in the sense of multi-valued mappings).

Moreau envelope and proximal mapping

Let $f : \mathcal{X} \rightarrow (-\infty, +\infty]$ be a closed proper convex function. We define:

- **Moreau envelope** (Moreau-Yosida regularization) of f at x :

$$M_f(x) = \min_y \left\{ f(y) + \frac{1}{2} \|y - x\|^2 \right\}.$$

- **Proximal mapping** of f at x :

$$P_f(x) = \arg \min_y \left\{ f(y) + \frac{1}{2} \|y - x\|^2 \right\}.$$

Properties:

- $M_f(x)$ is differentiable, and its gradient is

$$\nabla M_f(x) = x - P_f(x).$$

- $P_f(x)$ exists and is unique.
- $M_f(x) \leq f(x)$.
- $\arg \min_x f(x) = \arg \min_x M_f(x)$.

(The Moreau envelope is a way to **smooth a possibly non-differentiable convex function**.)

Example 1: Projection onto a closed convex set

Let $C \subseteq \mathcal{X}$ be a nonempty closed convex set, and $f(x) = \delta_C(x)$ (indicator function). Its proximal mapping:

$$P_f(x) = \arg \min_{y \in \mathcal{X}} \left\{ \delta_C(y) + \frac{1}{2} \|y - x\|^2 \right\} = \arg \min_{y \in C} \frac{1}{2} \|y - x\|^2 = \Pi_C(x),$$

where $\Pi_C(x)$ denotes projection onto C . Its Moreau envelope:

$$M_f(x) = \frac{1}{2} \|x - \Pi_C(x)\|^2.$$

Example 2: Huber function and soft thresholding

Let $f(x) = \lambda|x|$, $x \in \mathbb{R}$. Then its Moreau envelope (Huber function) is

$$M_f(x) = \begin{cases} \frac{1}{2}x^2, & |x| \leq \lambda, \\ \lambda|x| - \frac{\lambda^2}{2}, & |x| > \lambda. \end{cases}$$

Its proximal mapping (soft thresholding):

$$P_f(x) = \text{sign}(x) \max\{|x| - \lambda, 0\}.$$

Soft thresholding operator

The soft thresholding operator $S_\lambda : \mathbb{R}^n \rightarrow \mathbb{R}^n$ is defined as:

$$S_\lambda(x) = \begin{bmatrix} \text{sign}(x_1) \max\{|x_1| - \lambda, 0\} \\ \text{sign}(x_2) \max\{|x_2| - \lambda, 0\} \\ \vdots \\ \text{sign}(x_n) \max\{|x_n| - \lambda, 0\} \end{bmatrix},$$

for any $x = [x_1, \dots, x_n] \in \mathbb{R}^n$ and $\lambda > 0$.

Moreau decomposition

Theorem (Moreau decomposition):

Let $f : \mathcal{X} \rightarrow (-\infty, +\infty]$ be a closed proper convex function and f^* its Fenchel conjugate. Then, for any $x \in \mathcal{X}$,

$$x = P_f(x) + P_{f^*}(x),$$

and

$$\frac{1}{2} \|x\|^2 = M_f(x) + M_{f^*}(x).$$

Example: Let $C \subseteq \mathcal{X}$ be a nonempty closed convex cone. Take $f(x) = \delta_C(x)$, so $f^*(x) = \delta_{C^\circ}(x)$. Therefore,

$$x = \Pi_C(x) + \Pi_{C^\circ}(x).$$

Remarks:

- $M_f(\cdot)$ is always differentiable, even if f is non-differentiable.
- $P_f(\cdot)$ is important for optimization algorithms (e.g., accelerated proximal gradient methods).
- For many regularizers, $P_f(\cdot)$ and $M_f(\cdot)$ have explicit expressions.

A proximal point view of gradient methods

To minimize a differentiable function $\min_{\beta} f(\beta)$, the gradient update is

$$\beta^{(k+1)} = \beta^{(k)} - \alpha_k \nabla f(\beta^{(k)}).$$

The gradient step can be equivalently written as ((linear approximation + proximal term):

$$\beta^{(k+1)} = \arg \min_{\beta} \left\{ f(\beta^{(k)}) + \langle \nabla f(\beta^{(k)}), \beta - \beta^{(k)} \rangle + \frac{1}{2\alpha_k} \|\beta - \beta^{(k)}\|^2 \right\}.$$

Optimizing composite functions

Minimize

$$\min_{\beta \in \mathbb{R}^p} f(\beta) + g(\beta),$$

where:

- $f: \mathbb{R}^p \rightarrow \mathbb{R}$ is convex, differentiable, ∇f is L -Lipschitz continuous,
- $g: \mathbb{R}^p \rightarrow (-\infty, +\infty]$ is closed, proper convex, possibly non-differentiable.

Example (Lasso):

$$\min_{\beta \in \mathbb{R}^p} \frac{1}{2} \|X\beta - Y\|^2 + \lambda \|\beta\|_1,$$

where the first term is $f(\beta)$ and the second is $g(\beta)$. Since g is non-differentiable, gradient methods alone cannot be applied.

Proximal gradient step

Gradient step (if $g(\beta)$ disappears):

$$\beta^{(k+1)} = \beta^{(k)} - \alpha_k \nabla f(\beta^{(k)}),$$

or equivalently,

$$\beta^{(k+1)} = \arg \min_{\beta} \left\{ f(\beta^{(k)}) + \langle \nabla f(\beta^{(k)}), \beta - \beta^{(k)} \rangle + \frac{1}{2\alpha_k} \|\beta - \beta^{(k)}\|^2 \right\}.$$

Proximal gradient step for $f + g$:

$$\beta^{(k+1)} = \arg \min_{\beta} \left\{ f(\beta^{(k)}) + \langle \nabla f(\beta^{(k)}), \beta - \beta^{(k)} \rangle + g(\beta) + \frac{1}{2\alpha_k} \|\beta - \beta^{(k)}\|^2 \right\}.$$

After ignoring constant terms and completing the square:

$$\beta^{(k+1)} = \arg \min_{\beta} \left\{ \frac{1}{2\alpha_k} \left\| \beta - \left(\beta^{(k)} - \alpha_k \nabla f(\beta^{(k)}) \right) \right\|^2 + g(\beta) \right\}.$$

That is,

$$\beta^{(k+1)} = P_{\alpha_k g} \left(\beta^{(k)} - \alpha_k \nabla f(\beta^{(k)}) \right),$$

where $P_{\alpha_k g}$ denotes the proximal operator associated to $\alpha_k g$.

Derivation (completing the square):

$$\begin{aligned} & \langle \nabla f(\beta^{(k)}), \beta \rangle + \frac{1}{2\alpha_k} \|\beta - \beta^{(k)}\|^2 \\ &= \langle \nabla f(\beta^{(k)}), \beta \rangle - \frac{1}{\alpha_k} \langle \beta^{(k)}, \beta \rangle + \frac{1}{2\alpha_k} \|\beta\|^2 + \text{constant}. \end{aligned}$$

Convergence Rate of PG:

In convex problems, PG method satisfies

$$f(\beta^{(k)}) + g(\beta^{(k)}) - \min_{\beta \in \mathbb{R}^p} (f(\beta) + g(\beta)) \leq O\left(\frac{1}{k}\right).$$

If stopping condition

$$f(\beta^{(k)}) + g(\beta^{(k)}) - \text{optimal value} \leq 10^{-4},$$

then around $O(10^4)$ iterations needed.

Algorithm 6 Proximal Gradient (PG) Method

1: **Initialization:** Choose initial point $\beta^{(0)}$, step size $\alpha > 0$, set $k \leftarrow 0$.
2: **while** not converged **do**
3: Update the iterate:

$$\beta^{(k+1)} = P_{\alpha g} \left(\beta^{(k)} - \alpha \nabla f(\beta^{(k)}) \right)$$

4: Increment $k \leftarrow k + 1$.

5: **end while**

6: **Output:** $\beta^{(k)}$ (approximate solution)

Algorithm 7 Accelerated Proximal Gradient (APG) Method

1: **Initialization:** Choose initial point $\beta^{(0)}$, set $t_0 = t_1 = 1$, step size $\alpha > 0$, set $k \leftarrow 0$.
2: **while** not converged **do**
3: Compute extrapolated point:

$$\bar{\beta}^{(k)} = \beta^{(k)} + \frac{t_k - 1}{t_{k+1}} \left(\beta^{(k)} - \beta^{(k-1)} \right)$$

4: Proximal step:

$$\beta^{(k+1)} = P_{\alpha g} \left(\bar{\beta}^{(k)} - \alpha \nabla f(\bar{\beta}^{(k)}) \right)$$

5: Update:

$$t_{k+1} = \frac{1 + \sqrt{1 + 4t_k^2}}{2}$$

6: Increment $k \leftarrow k + 1$.

7: **end while**

8: **Output:** $\beta^{(k)}$ (approximate solution)

Convergence Rate of APG:

In convex problems, APG method satisfies

$$f(\beta^{(k)}) + g(\beta^{(k)}) - \min_{\beta \in \mathbb{R}^p} (f(\beta) + g(\beta)) \leq O\left(\frac{1}{k^2}\right).$$

If stopping condition

$$f(\beta^{(k)}) + g(\beta^{(k)}) - \text{optimal value} \leq 10^{-4},$$

then around $O(10^2)$ iterations needed.

Accelerated Proximal Gradient (APG) Methods

- Backtracking line search can also be used for finding step length α_k .
- For simplicity, we often take a constant step length. It should satisfy

$$\alpha \in \left(0, \frac{1}{L} \right),$$

where L is the Lipschitz constant of $\nabla f(\cdot)$ (typically unknown).

- APG methods enjoy the same computational cost per iteration as PG methods.
- Iteration complexity:

$$\text{APG: } O\left(\frac{1}{k^2}\right), \quad \text{PG: } O\left(\frac{1}{k}\right).$$

APG for Lasso

Solve:

$$\min_{\beta \in \mathbb{R}^p} \frac{1}{2} \|X\beta - Y\|^2 + \lambda \|\beta\|_1 \quad (\text{lasso problem}),$$

where

$$f(\beta) = \frac{1}{2} \|X\beta - Y\|^2, \quad g(\beta) = \lambda \|\beta\|_1.$$

Then

$$\nabla f(\beta) = X^T (X\beta - Y) \quad \text{with Lipschitz constant } L = \lambda_{\max}(X^T X).$$

Choose step length $\alpha = 1/L$.

APG iterations:

$$\tilde{\beta}^{(k)} = \beta^{(k)} + \frac{t_k - 1}{t_{k+1}} (\beta^{(k)} - \beta^{(k-1)}),$$

$$\beta^{(k+1)} = S_{\lambda/L} \left(\tilde{\beta}^{(k)} - \frac{1}{L} X^T (X \tilde{\beta}^{(k)} - Y) \right),$$

where $S_{\lambda/L}$ is the soft-thresholding operator.

APG is also applicable to "logistic regression + lasso regularization".

Restart Strategy

- To speed up APG, restart the algorithm after a fixed number of iterations.
- Use the latest iterate as the starting point for a new APG round.
- A reasonable choice is to restart every 100 or 200 iterations.