

DSA5105 Principles of Machine Learning

AY2024/25 Sem1 By Zhao Peiduo

Supervised Learning

Lecture 1

Empirical Risk Minimization (ERM) The learning process aims to find a function $f \in \mathcal{H}$ that minimizes the empirical risk, $R_{\text{emp}}(f) = \frac{1}{N} \sum_{i=1}^N L(y_i, f(x_i))$, where $y_i = f^*(x_i)$.

Common Loss Function Mean squared error (MSE) for regression tasks: $L(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2$, cross-entropy loss for classification task: $L(y, p) = -\sum_i y_i \log(p_i)$, and huber loss for robust regression: $L_{\delta}(y, \hat{y}) = \begin{cases} \frac{1}{2}(y - \hat{y})^2 & \text{if } |y - \hat{y}| \leq \delta \\ \delta|y - \hat{y}| - \frac{1}{2}\delta^2 & \text{otherwise} \end{cases}$

Softmax Function For a multi-class classification problem with K classes: $\text{softmax}(z_i) = \frac{\exp(z_i)}{\sum_{j=1}^K \exp(z_j)}$

General Ordinary Least Squares Formula Consider $x \in \mathbb{R}^d$ and the new hypothesis space $\mathcal{H}_M = \{f : f(x) = \sum_{j=0}^{M-1} w_j \varphi_j(x)\}$ Each $\varphi_j : \mathbb{R}^d \rightarrow \mathbb{R}$ is called a **basis function** or **feature map**.

We can rewrite the ERM $\min_{w \in \mathbb{R}^M} \frac{1}{2N} \sum_{i=1}^N \left(\sum_{j=0}^{M-1} w_j \varphi_j(x_i) - y_i \right)^2$ into $\min_{w \in \mathbb{R}^M} \frac{1}{2N} \|\Phi w - y\|^2$. Solving by setting $\nabla R_{\text{emp}}(\hat{w}) = 0$, we have $\hat{w} = (\Phi^T \Phi)^{-1} \Phi^T y$, given invertible $\Phi^T \Phi$.

For cases where $\Phi^T \Phi$ is not invertible, the formula using the Moore-Penrose pseudoinverse is: $\hat{w}(u) = \Phi^\dagger y + (I - \Phi^\dagger \Phi) u \quad u \in \mathbb{R}^M$

Regularization To prevent overfitting, regularization techniques add a penalty to the loss function: $\min_{w \in \mathbb{R}^M} \frac{1}{2N} \|\Phi w - y\|^2$. For ridge regression, minimizing the ERM we get $\hat{w} = (\Phi^T \Phi + \lambda I_M)^{-1} \Phi^T y$ which is always invertible for positive λ . Common regularization terms include Ridge(L2): $\lambda \sum_{j=1}^p w_j^2$ and Lasso(L1): $\lambda \sum_{j=1}^p |w_j|$

Lecture 2

Reformulation of Ridge Regression We rewrite the regularized least squares solution in another way:

$$\hat{w} = (\Phi^T \Phi + \lambda I_M)^{-1} \Phi^T y = \Phi^\top (\Phi \Phi^\top + \lambda I_N)^{-1} y$$

Reformulation of Ridge Regression

$\hat{f}(x) = \sum_{i=1}^N \alpha_i \varphi(x_i)^\top \varphi(x) \alpha = (G + \lambda I_N)^{-1} y$ where $G_{ij} = \varphi(x_i)^\top \varphi(x_j)$ is the gram matrix
Mercer's Theorem and SPD Kernels Suppose k is a SPD kernel. Then, there exists a feature space \mathcal{H} and a feature map $\varphi : \mathbb{R}^d \rightarrow \mathcal{H}$ such that $k(x, x') = \varphi(x)^\top \varphi(x')$

SPD kernels properties: $K(x, x') = K(x', x)$ (**Symmetry**) For any n and $\{x_1, \dots, x_n\}$, the Gram matrix $G_{ij} = k(x_i, x_j)$ is positive semi-definite. (Recall: a matrix G is positive semi-definite if $c^T G c \geq 0$ for any vector c) (**Positive Semi-definiteness**)

Examples of SPD Kernels: **Linear Kernel:** $K(x, x') = x^\top x'$ **Polynomial Kernel:** $K(x, x') = (1 + x^\top x')^d$

Gaussian RBF Kernel: $K(x, x') = \exp\left(-\frac{\|x - x'\|^2}{2\sigma^2}\right)$

Constructing kernels Given valid kernels $k_1(x, x')$ and $k_2(x, x')$, the following new kernels will also be valid:
 $ck_1(x, x')$ $f(x)k_1(x, x')f(x')$ $q(k_1(x, x'))$ $\exp(k_1(x, x'))$ $k_1(x, x') + k_2(x, x')$ $k_1(x, x')k_2(x, x')$
 $k_3(\varphi(x), \varphi(x'))$ $x^\top A x'$ $k_a(x_a, x'_a) + k_b(x_b, x'_b)$ $k_a(x_a, x'_a)k_b(x_b, x'_b)$

Lecture 3

SVM Max Margin Formulation $\max_{\mathbf{w}, b} \frac{1}{\|\mathbf{w}\|} \min_{i=1, \dots, N} |\mathbf{w}^\top \mathbf{x}_i + b|$ subject to $y_i(\mathbf{w}^\top \mathbf{x}_i + b) > 0 \quad \forall i$

Optimization Problem $\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$ subject to: $y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 \quad \forall i$. Introducing Lagrange multipliers $\alpha_i \geq 0$, the Lagrangian is: $\mathcal{L}(\mathbf{w}, b, \alpha) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^N \alpha_i [y_i(\mathbf{w}^\top \mathbf{x}_i + b) - 1]$

KKT conditions for SVM

- From Stationarity** $\hat{\mathbf{w}} = \sum_{i=1}^N \hat{\mu}_i y_i x_i, \quad 0 = \sum_{i=1}^N \hat{\mu}_i y_i$
- From Dual Feasibility** $\hat{\mu}_i \geq 0 \quad \text{for } i = 1, \dots, N$
- From Complementary Slackness** $\hat{\mu}_i = 0 \quad \text{or } y_i(\hat{\mathbf{w}}^\top x_i + \hat{b}) = 1$
- The multipliers** $\hat{\mu}$ can be found by the dual problem $\max_{\mu \in \mathbb{R}^N} \sum_{i=1}^N \mu_i - \frac{1}{2} \sum_{i,j} \mu_i \mu_j y_i y_j (x_i^\top x_j)$

Dual formulation of SVM

$\hat{\mu} = \arg \max_{\mu \in \mathbb{R}^N} \sum_{i=1}^N \mu_i - \frac{1}{2} \sum_{i,j} \mu_i \mu_j y_i y_j (x_i^\top x_j)$ Subject to: $\hat{\mu} \geq 0$ and $\sum_{i=1}^N \hat{\mu}_i y_i = 0$

Decision function: $\hat{f}(x) = \text{sgn} \left(\sum_{i=1}^N \hat{\mu}_i y_i x_i^\top x + \hat{b} \right)$ Complementary slackness: $\hat{\mu}_i = 0 \quad \text{or } 1 = y_i(\hat{\mathbf{w}}^\top x_i + \hat{b})$

Kernel SVMs

$\hat{\mu} = \arg \max_{\mu \in \mathbb{R}^N} \sum_{i=1}^N \mu_i - \frac{1}{2} \sum_{i,j} \mu_i \mu_j y_i y_j k(x_i, x_j)$ Subject to: $\hat{\mu} \geq 0$ and $\sum_{i=1}^N \hat{\mu}_i y_i = 0$

Decision function: $\hat{f}(x) = \text{sgn} \left(\sum_{i=1}^N \hat{\mu}_i y_i k(x_i, x) + \hat{b} \right)$

Only support vectors satisfying $1 = y_i(\hat{\mathbf{w}}^\top \varphi_i(x) + \hat{b})$ matter for predictions. This is a **sparse kernel method**.

Lecture 4

Classification and Regression Trees Suppose that the input space is \mathcal{X} . A **partition** of \mathcal{X} is a collection of subsets $\mathcal{R}_1, \mathcal{R}_2, \dots, \mathcal{R}_J$ such that $\mathcal{R}_i \cap \mathcal{R}_j = \emptyset$ for $i \neq j$ and $\bigcup_{j=1}^J \mathcal{R}_j = \mathcal{X}$

The general decision tree hypothesis space is: $\mathcal{H} = \left\{ f : f(x) = \sum_{j=1}^J a_j \mathbb{I}_{x \in \mathcal{R}_j}, \{ \mathcal{R}_j \} \text{ is a partition of } \mathcal{X}, a_j \in \mathcal{Y} \right\}$ where $\mathbb{I}_{x \in \mathcal{R}_j}$ is an indicator variable returning 1 if x is in \mathcal{R}_j .

A decision tree model $f(x) = \sum_{j=1}^J a_j \mathbb{I}_{x \in \mathcal{R}_j}$ depends on both a_j and \mathcal{R}_j . For regression we take the **average** label

values $a_j = y_j = \frac{\sum_i y_i \mathbb{I}_{x \in \mathcal{R}_j}}{\sum_i \mathbb{I}_{x \in \mathcal{R}_j}}$; For classification we take the **mode** label values $a_j = \text{mode}\{y_i : x_i \in \mathcal{R}_j\}$

Loss function for Decision Trees Classification Entropy: $-\sum_{k=1}^K \sum_{j=1}^J p_{jk} \log p_{jk}$

Gini Impurity: $\sum_{k=1}^K \sum_{j=1}^J p_{jk}(1 - p_{jk})$ where p_{jk} is the proportion of samples in \mathcal{R}_j belonging to class k .

Bagging reduces the variance by aggregating predictions from multiple models trained on different random subsamples of the data: **Regression** $\hat{f}(x) = \frac{1}{m} \sum_{j=1}^m \hat{f}_j(x)$ **Classification** $\hat{f}(x) = \text{Mode}\{\hat{f}_j(x) : j = 1, \dots, m\}$

Boosting works by training weak learners sequentially, where each learner focuses on correcting the mistakes of the previous ones. The combined model is a weighted sum of the weak learners: $\hat{f}(x) = \sum_{t=1}^T \alpha_t \hat{f}_t(x)$ where α_t are coefficients based on each learner's performance. Boosting helps reduce bias.

Key Ideas of AdaBoost

- Initialize with uniform weight across all training samples
- Train a classifier/regressor f_1
- Identify the samples that f_1 got wrong (classification) or has large errors (regression)
- Weight these samples more heavily and train f_2 on this reweighted dataset
- Repeat steps 3-5

AdaBoost Implementation

Data: $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$, Initialize $w_i^{(1)} = \frac{1}{N}$ for all $i = 1, \dots, N$;
For $j = 1, \dots, m$ do

- Obtain f_j from: $f_j = \arg \min_{f \in \mathcal{H}} \sum_{i=1}^N w_i^{(j)} \mathbb{I}_{y_i \neq f(x_i)}$
- Compute combination coefficients: $\delta_j = \frac{\sum_{i=1}^N w_i^{(j)} \mathbb{I}_{y_i \neq f_j(x_i)}}{\sum_{i=1}^N w_i^{(j)}} \quad \alpha_j = \log \left(\frac{1 - \delta_j}{\delta_j} \right)$
- Update weights: $w_i^{(j+1)} = w_i^{(j)} \exp \left(\alpha_j \mathbb{I}_{y_i \neq f_j(x_i)} \right)$

Return: $\hat{f}(x) = \text{Sign} \left(\sum_{j=1}^m \alpha_j f_j(x) \right)$

Cross-Validation In k -fold cross-validation, the data is split into k subsets. The model is trained on $k - 1$ subsets and validated on the remaining one. This process is repeated k -times.

Lecture 5

Activation Functions

Sigmoid: $\sigma(z) = \frac{1}{1 + e^{-z}}, [0, 1]$ **ReLU** $\sigma(z) = \max(0, z), [0, \text{inf}]$, leaky ReLU: $\delta \text{ if } z < 0$ instead of 0

Gradient Descent $w^{(t+1)} = w^{(t)} - \eta \frac{\partial \mathcal{L}}{\partial w}$

Stochastic Gradient Descent (SGD): $\theta_{k+1} = \theta_k - \eta \frac{1}{J} \sum_{i \in I_B} \nabla \Phi_i(\theta_k)$, less likely to be locally optimal.

Deep Neural Networks (DNNs) Deep neural networks are an extension of shallow networks. The idea is to stack n hidden layers together and forward pass the x sequentially.

Back-propagation Algorithm

Initialize $x_0 = x \in \mathbb{R}^d$.

For $t = 0, 1, \dots, T$: $x_{t+1} = g_t(x_t, W_t) = \sigma(W_t x_t)$

Set $p_{T+1} = \nabla_x L(x_{T+1}, y)$.

For $t = T, T-1, \dots, 1$:

$$-\nabla W_t \Phi = p_{t+1}^\top \nabla W g_t(x_t, W_t)$$

$$-p_t = [\nabla_x g_t(x_t, W_t)]^\top p_{t+1}$$

Return $\{\nabla_{W_t} \Phi : t = 0, \dots, T\}$.

Back-propagation Example Computation

The model is defined as: $y(x) = v \delta(w_1 \delta(w_0 x))$, $x, w_0, w_1, v \in \mathbb{R}$, where δ is the identity function, i.e., $\delta(z) = z$.

The loss function is: $L = (y(x) - y)^2$ We want to compute the gradients: $\frac{\partial L}{\partial v}, \frac{\partial L}{\partial w_1}, \frac{\partial L}{\partial w_0}$

Forward pass: $x_1 = w_0 x \quad x_2 = v w_1 x_1 \quad L = (x_2 - y)^2$

Backward pass:

$$p_2 = \frac{\partial L}{\partial x_2} = 2(x_2 - y) \quad \frac{\partial L}{\partial v} = \frac{\partial L}{\partial x_2} \cdot \frac{\partial x_2}{\partial v} = 2(x_2 - y) w_1 x_1 \quad \frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial x_2} \cdot \frac{\partial x_2}{\partial w_1} = 2(x_2 - y) v x_1$$

$$p_1 = \frac{\partial L}{\partial x_1} = \frac{\partial L}{\partial x_2} \cdot \frac{\partial x_2}{\partial x_1} = 2(x_2 - y) v w_1 \quad \frac{\partial L}{\partial w_0} = \frac{\partial L}{\partial x_1} \cdot \frac{\partial x_1}{\partial w_0} = 2(x_2 - y) v w_1 x$$

Lecture 6

PCA Algorithm Simplified Flow Center the data \rightarrow compute sample covariance matrix $S = \frac{1}{N} \sum_{i=1}^N x_i x_i^\top \rightarrow$ compute top k eigenvalues and retrieve their corresponding eigenvectors. In feature space we replace all x with ϕ , the feature maps, and compute design matrix as $\Phi_{ij} \leftarrow \Phi_{ij} - \frac{1}{N} \sum_{i=1}^N \Phi_{ij}$, where $\Phi_{ij} = \phi_j(x_i)$ is the raw design matrix.
PCA Whitening Transform Principal component scores are given by $Z = XU$ where X is the original features and U is the matrix of eigenvectors. The transformation $X' = XU \Lambda^{-\frac{1}{2}}$, where Λ is the matrix of eigenvalues makes $\text{cov}(X') = I$.

Neural Network AutoEncoders Choose encoder T_{enc} and decoder T_{dec} as:

$T_{\text{enc}}(x; \theta) = A \sigma(Wx + b) \quad T_{\text{dec}}(x; \phi) = B \sigma(Vx + c)$

$\theta = (A, W, b) \in \mathbb{R}^{m \times q} \times \mathbb{R}^{q \times d} \times \mathbb{R}^q \quad \phi = (B, V, c) \in \mathbb{R}^{d \times q'} \times \mathbb{R}^{q' \times m} \times \mathbb{R}^{q'}$

Given a dataset $\mathcal{D} = \{x_i\}_{i=1}^N$, we solve the empirical risk minimization to minimize the distance between x_i and x'_i : $\min_{\theta, \phi} \frac{1}{2N} \sum_i \|x_i - T_{\text{dec}}(T_{\text{enc}}(x_i; \theta); \phi)\|^2$, where the input is used as labels.

Unsupervised Learning

Lecture 7

Clustering partition a dataset \mathcal{D} into disjoint groups $\mathcal{D} = \mathcal{D}_1 \cup \mathcal{D}_2 \cup \dots \cup \mathcal{D}_K$, such that data grouped together are similar and dissimilar if they are in different groups.

K-means Algorithm To minimize the distortion $\min_{Z,R} J(R, Z) = \frac{1}{2N} \sum_{i=1}^N \sum_{k=1}^K r_{ik} \|x_i - z_k\|^2$:

Data: $\mathcal{D} = \{x_i\}_{i=1}^N, x_i \in \mathbb{R}^d$ for all i

Hyperparameters: K (number of clusters); stopping criterion (convergence of loss function j for instance)

Initialize $Z \in \mathbb{R}^{K \times d}$, that is, starting with some centroid set (which could determine algo performance)

While stopping criterion not reached do

Assign point to the nearest centroid: $r_{ik} = \begin{cases} 1 & \text{if } k = \arg \min_j \|x_i - z_j\|^2, i = 1, \dots, N \\ 0 & \text{otherwise} \end{cases}$

Recompute centroid by taking average of members of cluster: $z_k = \frac{\sum_{i=1}^N r_{ik} x_i}{\sum_{i=1}^N r_{ik}}, k = 1, \dots, K$

Return: centroids Z , assignment matrix R

Depending on the initial condition, we can end up with global/local optimum or saddle point.

Maximum Likelihood Estimation Given dataset x_i , model them as i.i.d. samples from $p(x|\theta)$, θ are parameters to be determined: $\theta_{MLE} = \arg \max_{\theta} \log(\prod_i p(x_i|\phi)) = \sum_i \log(p(x_i|\phi))$

Gaussian Mixture Models (GMMs) Soft probabilistic labeling as compared to K-means (hard deterministic labeling): Model the data as samples from a linear convex combination of K Gaussian:

$p(x) = \sum_{k=1}^K \pi_k p_g(x; z_k, \Sigma_k)$, where $p_g(x|z, \Sigma) = (2\pi)^{-\frac{d}{2}} |\Sigma|^{-\frac{1}{2}} \exp\left\{-\frac{1}{2}(x-z)^T \Sigma^{-1}(x-z)\right\}$,

$\pi_k \geq 0, \sum_k \pi_k = 1$ (Mixture weights)

Sampling from GMM

- Sample r (one-hot) according to π_k - find a member Gaussian distribution
- Identify the "hot" coordinate of r - Go to that distribution
- Sample x from $p_g(x; z_\ell, \Sigma_\ell)$ - sampling by its mean and covariance matrix

MLE for GMM

Data: $\mathcal{D} = \{x_i\}_{i=1}^N, x_i \in \mathbb{R}^d$ for all i **Hyperparameters:** K (num_clusters); stopping criterion

Initialize: $\pi_k = 1/K, z_k \in \mathbb{R}^d, \Sigma_k \in \mathbb{R}^{d \times d}$ for $k = 1, \dots, K$

While stopping criterion not reached do

Update γ_{ik} , the responsibility term: $r_{ik} = \frac{\pi_k p_g(x; z_k, \Sigma_k)}{\sum_{\ell=1}^K \pi_\ell p_g(x; z_\ell, \Sigma_\ell)}, i = 1, \dots, N, k = 1, \dots, K$

Compute $\{N_k = \sum_{i=1}^N \gamma_{ik}\}$ and Update $\{\pi_k, z_k, \Sigma_k\}$:

$z_k = \frac{1}{N_k} \sum_{i=1}^N \gamma_{ik} x_i, \quad \Sigma_k = \frac{1}{N_k} \sum_{i=1}^N \gamma_{ik} (x_i - z_k)(x_i - z_k)^T, \quad \pi_k = \frac{N_k}{N}$ for $k = 1, \dots, K$

Return: cluster centers, covariances, and mixture coefficients $\{z_k, \Sigma_k, \pi_k\}$, cluster responsibilities or soft assignments $\{\gamma_{ik}\}$

Relationship between K-means and GMM Set $\Sigma_k = \epsilon I_d$, then the responsibility term becomes $\gamma_{ik} = \frac{\pi_k \exp\left(-\frac{1}{2\epsilon} \|x_i - z_k\|^2\right)}{\sum_{\ell} \pi_{\ell} \exp\left(-\frac{1}{2\epsilon} \|x_i - z_{\ell}\|^2\right)}$, and when we take $\epsilon \rightarrow 0$ we will get the hard label assignment as K-means.

Reinforcement Learning

Lecture 8

Key Elements of RL Action (action space, policy/chain of actions leading to optimal reward), State (match with action to form transition function, terminal / goal state), Reward (reward function)

Markov Property The future does not depend on the past, i.e. for transition probability $\mathbb{P}(S_{t+1} = s' | S_t = s, S_{t-1} = s_{t-1}, \dots, S_0 = s_0)$, it is simply $\mathbb{P}(S_{t+1} = s' | S_t = s)$.

Transition Matrix In time homogeneous Markov Chain, The transition probability is independent of time: $\mathbb{P}(S_{t+1} = s' | S_t = s) = P_{ss'} = p(s'|s)$ is the transition matrix.

Markov Decision Process (MDP) $p(s', r|s, a) = \mathbb{P}(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a)$, which integrate reward R and action a to the transition probability. It is finite if the action space is discrete finite for each state.

Policy $\pi(a|s) = \mathbb{P}(A_t = a | S_t = s)$, and deterministic policy is a constant mapping from the state space to action space $\pi: S \rightarrow \mathcal{A}$

(Discounted) Return $G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$, $\gamma \in (0, 1]$, where γ is the discount factor. The goal is typically to maximize the expected return, and in the context of time homogeneous case, $\mathbb{E}_{\pi} [G_t | S_t = s] = \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s \right]$

Dynamic Programming Memoization / caching the result of overlapping optimal substructures. Trade space for cost of recomputation for naive implementations and space usage can be optimized through iterative approach.

Bellman's Equation $v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$

This can be derived using DP to formulate a recursive solver for value function $v_{\pi}(s) = \mathbb{E}_{\pi} [G_t | S_t = s]$ and action value function $q_{\pi}(s, a) = \mathbb{E}_{\pi} [G_t | S_t = s, A_t = a]$, and substitute the action value function to the value function.

Bellman's Equation for Finite MDP $v_{\pi} = \gamma P(\pi) v_{\pi} + b(\pi)$, where:

- $P(\pi)_{s,s'} = \sum_a \pi(a|s) \sum_r p(s', r|s, a)$ is the state transition matrix for the policy π ($N * N$ matrix with N being number of states, so we only sum over r)
- $b(\pi)(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) \cdot r$ is the expected immediate reward for the policy π ($N * 1$ matrix)

The unique solution would thus be $v_{\pi} = (I - \gamma P(\pi))^{-1} b(\pi)$.

Optimal Policy $\pi_{*} \geq \pi \quad \forall \pi \neq \pi_{*}$ Always exist but may not be unique.

Policy Improvement For any two policies π, π' , if $\sum_a \pi'(a|s) q_{\pi}(s, a) \geq \sum_a \pi(a|s) q_{\pi}(s, a) \forall s$ then $v_{\pi'}(s) \geq v_{\pi}(s) \forall s$, and we assert π' is better than π .

Bellman's Optimality Equations Given an optimal policy $\pi_{*}(s) \in \arg \max_a q_{*}(s, a)$, we can recursively find optimal value and action value functions:

$v_{*}(s) = \max_{a \in \mathcal{A}} \sum_{s', r} p(s', r|s, a) [r + \gamma v_{*}(s')] \quad q_{*}(s, a) = \sum_{s', r} p(s', r|s, a) [r + \gamma \max_{a' \in \mathcal{A}} q_{*}(s', a')]$

$v_{*}(s) = \max_a q_{*}(s, a) \quad \text{Optimal Policy: } \pi_{*}(s) \in \arg \max_a q_{*}(s, a)$

Lecture 9

Model-Based vs Model-Free Model-based assumes that the transition function is known: Value iteration, policy iteration. Model-free improves through sampling/ Monte Carlo simulations: Temporal Difference (TD), Q-learning.

Value Iteration Algorithm Start with some value function, update utility estimate iteratively and keep track of corresponding current optimal actions.

Model Parameters: MDP transition probability $p(s', r|s, a)$, discount rate $\gamma < 1$

Input: Policy π , Stopping criterion

Initialize $v \in \mathbb{R}^n$

While stopping criterion not reached do

Update value function: $v \leftarrow F(v) = \max_{\pi} \{\gamma P(\pi) v + b(\pi)\}$ (Make use of P, model-based)

Return v

Reinforcement Learning

Contraction Mapping Theorem Fix point iteration will lead to a unique $v_{*} \in V$ such that $v_{*} = F(v_{*})$, and v_{*} can be estimated by starting at arbitrary v_0 and iterate for many times (towards infinity).

Value Iteration Convergence Rate Given error ϵ being the difference between v_K and v_{∞} , and number of iterations k , to achieve $\mathcal{O}(\epsilon)$ we need $k \sim \mathcal{O}(\log(1/\epsilon))$

Policy Iteration Algorithm Start with some policy, plug into value function and derive a better policy. Repeat until policy does not change.

Model Parameters: MDP transition probability $p(s', r|s, a)$, discount rate $\gamma < 1$

Initialize: Deterministic policy $\pi \leftarrow \pi_0$

While $\pi \neq \pi'$ do

$\pi \leftarrow \pi'$ (Set the next policy for evaluation)

$v_{\pi} \leftarrow (I - \gamma P(\pi))^{-1} b(\pi)$ (Evaluate policy π using P, model-based)

$\pi' \leftarrow \arg \max_{\pi} \{b(\pi) + \gamma P(\pi) v_{\pi}\}$ (Retrieve subsequent best policy)

Return v

Its convergence can be shown by $v_{\pi_k} \leq v_{\pi_{k+1}} \leq v_{*}$ from the algorithm (argmax, minimally equivalent values) and $\pi_{k+1} \neq \pi_k$ if and only if $v_{\pi_{k+1}} > v_{\pi_k}$ (otherwise the while loop ends)

Monte Carlo Method IID sampling and compute expectation: $\mathbb{E}_{x \sim \mu} f(x) \approx \frac{1}{N} \sum_{i=1}^N f(X_i)$ where $X_i \sim \mu$ is IID

Model-free Policy Iteration Instead of relying on the transition matrix P we sample episodes according current π to get samples of state and reward sequence (S and R), and get the value by averaging discounted rewards. Concretely, update estimation for action-value function by: $q(s, a) \leftarrow \frac{1}{N} \sum_{n=1}^N \left[\sum_{k=0}^{\infty} \gamma^k R_{k+1}^{(n)} | S_0^{(n)} = s, A_0^{(n)} = a \right]$ and perform policy improvement based on this $q(s, a)$.

Temporal Difference (TD) Algorithm

Input: State-reward simulator, initial state sampler, Stopping criterion, Policy π

Initialize: Initial value v_0 , Episode length T , Start with $v \leftarrow v_0$;

While stopping criterion not reached do

$s \leftarrow \text{InitialStateSampler}()$ (The start of the new episode)

For $t = 0$ to $T - 1$ do

$a \leftarrow a \sim \pi(\cdot|s)$

$s', r \leftarrow \text{StateRewardSimulator}(s, a)$

$v(s) \leftarrow (1 - \alpha)v(s) + \alpha (r + \gamma v(s'))$

$s \leftarrow s'$

Return v

Q-learning Algorithm

Input: State-reward simulator, initial state sampler, Stopping criterion, Policy π

Initialize: Initial value q_0 , Episode length T , Start with $q \leftarrow q_0$

While stopping criterion not reached do

$s \leftarrow \text{InitialStateSampler}()$

For $t = 0$ to $T - 1$ do

$a \leftarrow a \sim \pi(\cdot|s)$

$s', r \leftarrow \text{StateRewardSimulator}(s, a)$

$q(s, a) \leftarrow (1 - \alpha)q(s, a) + \alpha [r + \gamma \max_{a'} q(s', a')]$

$s \leftarrow s'$

Return q

Graph-Based Methods

Lecture 10

Page Rank Let $L_k \subset \{1, 2, \dots, n\}$ denote the set of pages with a link to page k ; that is, L_k is the set of page k 's backlinks. For each k we require $x_k = \sum_{j \in L_k} \frac{x_j}{n_j}$ where n_j is the number of outgoing links from page j .

The resultant matrix A representing pairwise values of x_k is a column-stochastic matrix with an eigenvalue equal to

1, and 1 is also its largest eigenvalue. For instance: $A = \begin{bmatrix} 0 & 0 & 1 & \frac{1}{2} \\ \frac{1}{3} & 0 & 0 & 0 \\ \frac{1}{3} & \frac{1}{2} & 0 & 0 \\ \frac{1}{3} & \frac{1}{2} & 0 & 0 \end{bmatrix}$

Power Iteration Let the eigenvalues of A be $\lambda_1 > \lambda_2 > \dots > \lambda_n$ and the associated eigenvectors be v_1, v_2, \dots, v_n . Then any vector x_0 can be written as $x_0 = c_1 v_1 + c_2 v_2 + \dots + c_n v_n$. Then

$A^k x_0 = c_1 \lambda_1^k v_1 + c_2 \lambda_2^k v_2 + \dots + c_n \lambda_n^k v_n = c_1 \lambda_1^k \left(v_1 + \frac{c_2}{c_1} \left(\frac{\lambda_2}{\lambda_1} \right)^k v_2 + \dots + \frac{c_n}{c_1} \left(\frac{\lambda_n}{\lambda_1} \right)^k v_n \right) \rightarrow c_1 \lambda_1^k v_1$ as $k \rightarrow \infty$.

Lecture 11

Graph Laplacian $L_G := D - W = \sum_{i < j} w_{ij} (e_i - e_j)(e_i - e_j)^T$, the diagl entries are the degree of each node and other entries are $-w_{ij}$

Cheeger's cut $h(S) = \frac{\text{cut}(S)}{\min\{\text{vol}(S), \text{vol}(S^c)\}}$ **Normalized cut** (Ncut) $Ncut(S) = \frac{\text{cut}(S)}{\text{vol}(S)} + \frac{\text{cut}(S^c)}{\text{vol}(S^c)}$

$h(S) \leq Ncut(S) \leq 2h(S)$

From Ncut to relaxed balanced cut $Ncut(S) = y^T L_G y \quad y_i = \begin{cases} \left(\frac{\text{vol}(S^c)}{\text{vol}(S) \text{vol}(G)} \right)^{\frac{1}{2}} & \text{if } i \in S, \\ - \left(\frac{\text{vol}(S)}{\text{vol}(S^c) \text{vol}(G)} \right)^{\frac{1}{2}} & \text{if } i \in S^c. \end{cases}$

Spectral Clustering for multi-clusters

Given a graph $G = (V, E, W)$, let v_2, \dots, v_k be the eigenvectors corresponding to the second through $(k-1)$ th eigenvalues of the normalized Laplacian $\mathcal{L}_G = D^{-\frac{1}{2}} L_G D^{-\frac{1}{2}}$;

Let $\phi_m = D^{-\frac{1}{2}} v_m \in \mathbb{R}^n$;

Consider the map $\phi: V \rightarrow \mathbb{R}^{k-1}$ defined as $\phi(v_i) = [\phi_2(i) \quad \dots \quad \phi_k(i)]^T$, where D is the degree matrix with $D_{ii} = \deg(i) = \sum_{j=1}^n w_{ij}$;

Cluster the n points in $k-1$ dimensions into k clusters using k -means.

For example, given $k = 2$, we can compute $\phi_2 = D^{-\frac{1}{2}} v_2$, choose a threshold τ (try and error or k -means for $k = 2$), and set $S = \{i \in V : \phi_2(i) \leq \tau\}$.

Cheeger's inequality $\frac{1}{2} \lambda_2(\mathcal{L}_G) \leq h_G \leq \sqrt{2 \lambda_2(\mathcal{L}_G)}$